

# Reactive Probabilistic Programming for Scalable Bayesian Inference

Dmitry Bagaev



# Reactive Probabilistic Programming for Scalable Bayesian Inference

Dmitry Bagaev



# Probabilistic Programming?

# Probabilistic programming

Why not just programming?



# Probabilistic programming

Why not just programming?

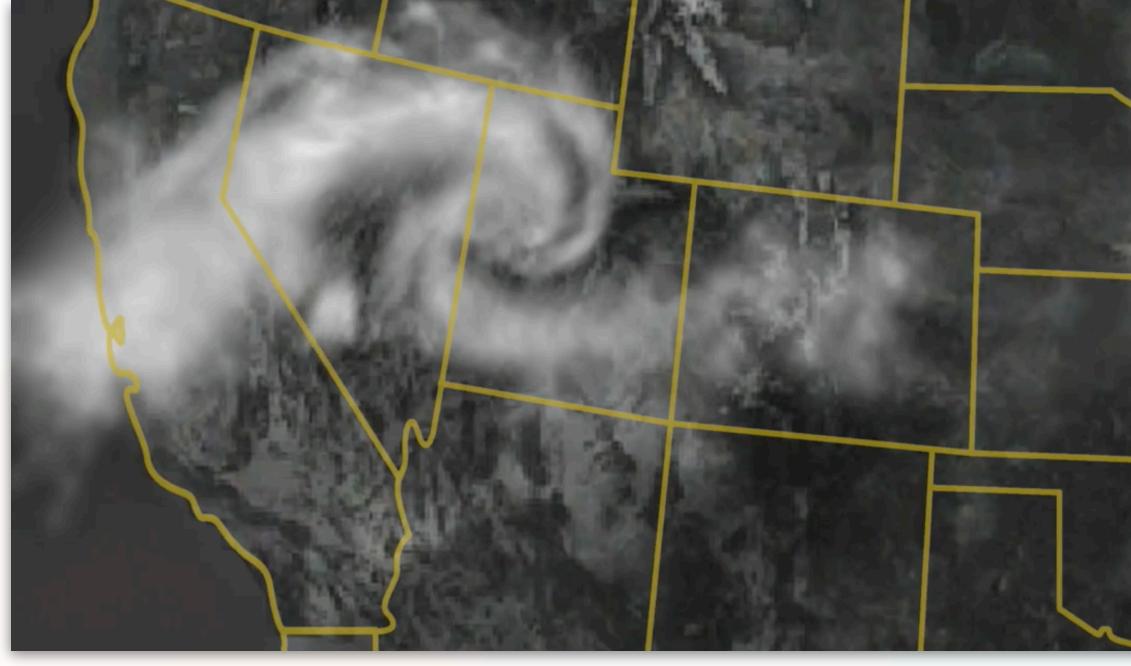


# Probabilistic programming

Do we really need it?



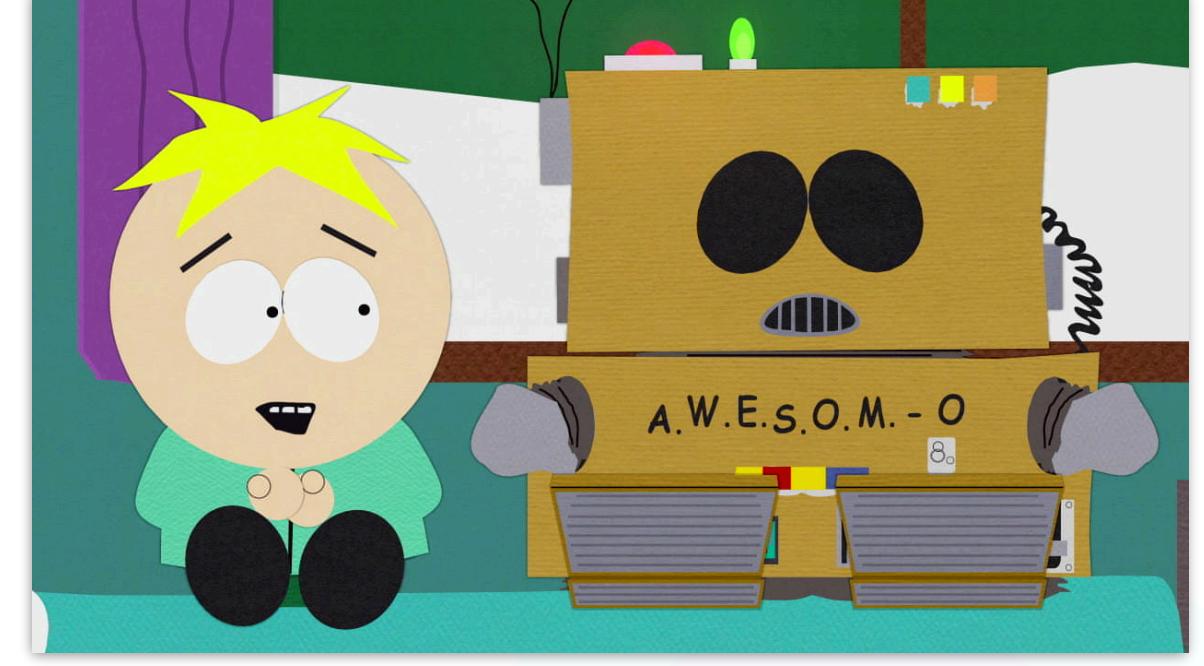
Pharmacology



Climate modelling



Logistics



Robotics



Auto-pilots



Finances



Audio

And so on...

# Probabilistic programming

Bayesian Inference

## Likelihood

how “likely” are the states,  
given the data?

$$p(s|\hat{y}) = \frac{p(\hat{y}|s)p(s)}{p(\hat{y})}$$

## Posterior

updated knowledge: how  
probable are hidden  
states given the observed data

## Posterior

prior knowledge: how probable are  
the hidden states before observing  
any data

## Evidence

how probably are these  
particular observations  
given our assumptions

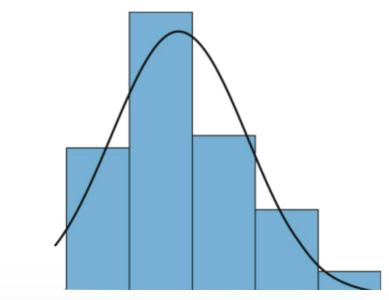
# Why is it so challenging?

# Probabilistic programming

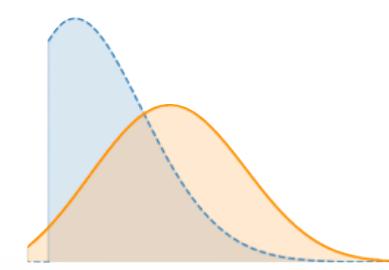
Why is it so challenging?

Account for the uncertainty in data 🎲

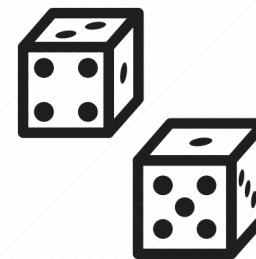
velocity =



position =



windspeed =



$$y = 2x$$

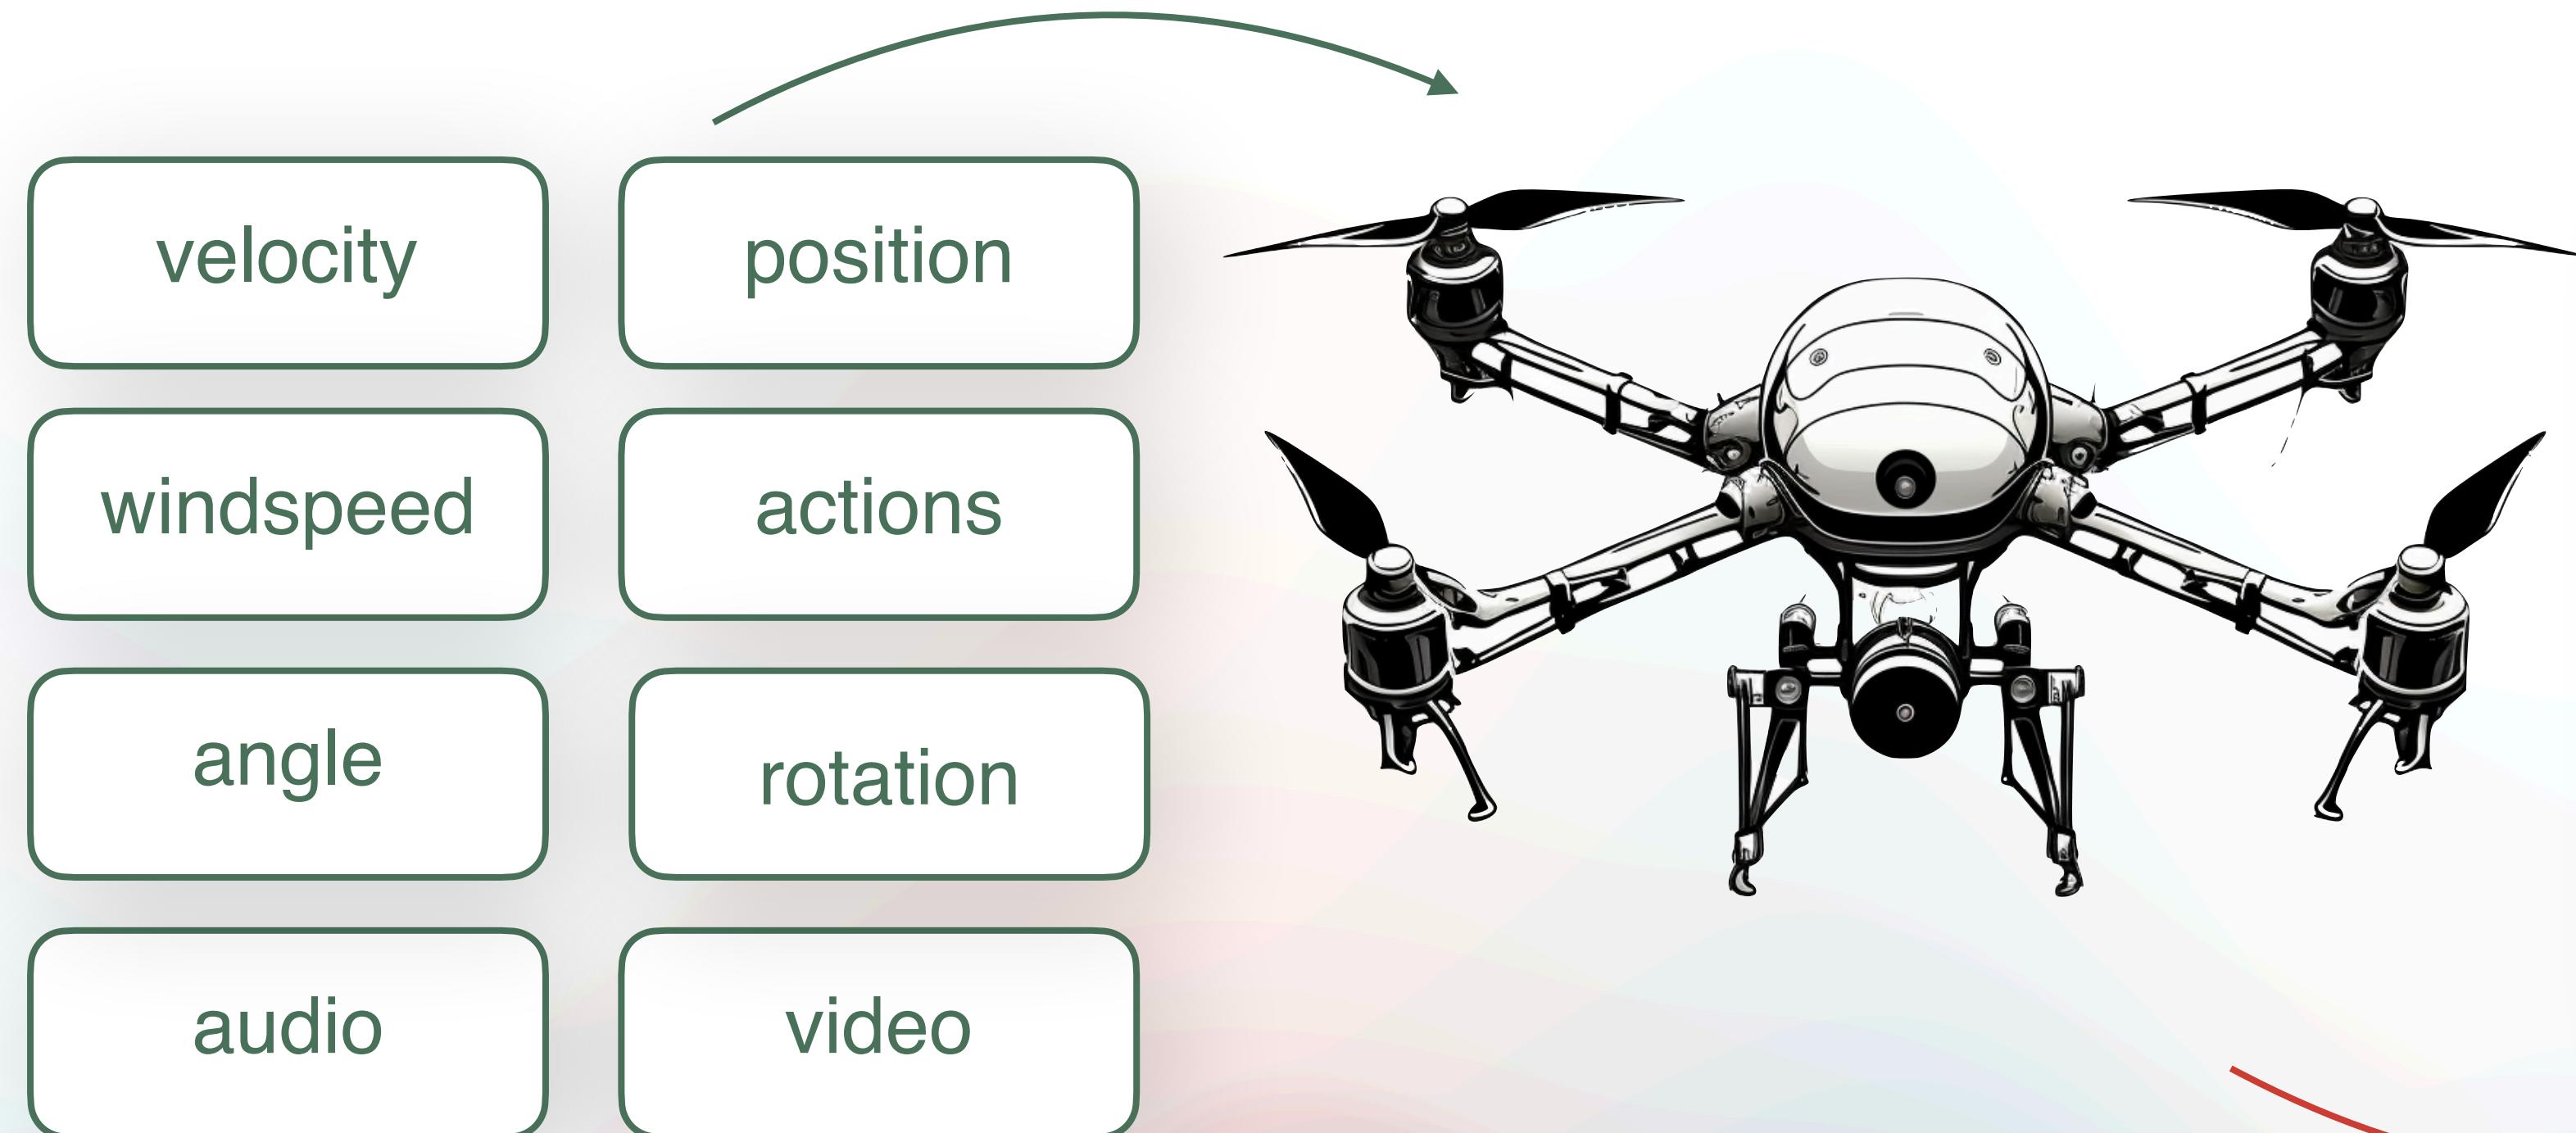
$$x = 1$$

Solve simple math 😎

# Probabilistic programming

Why is it so challenging?

Add more components



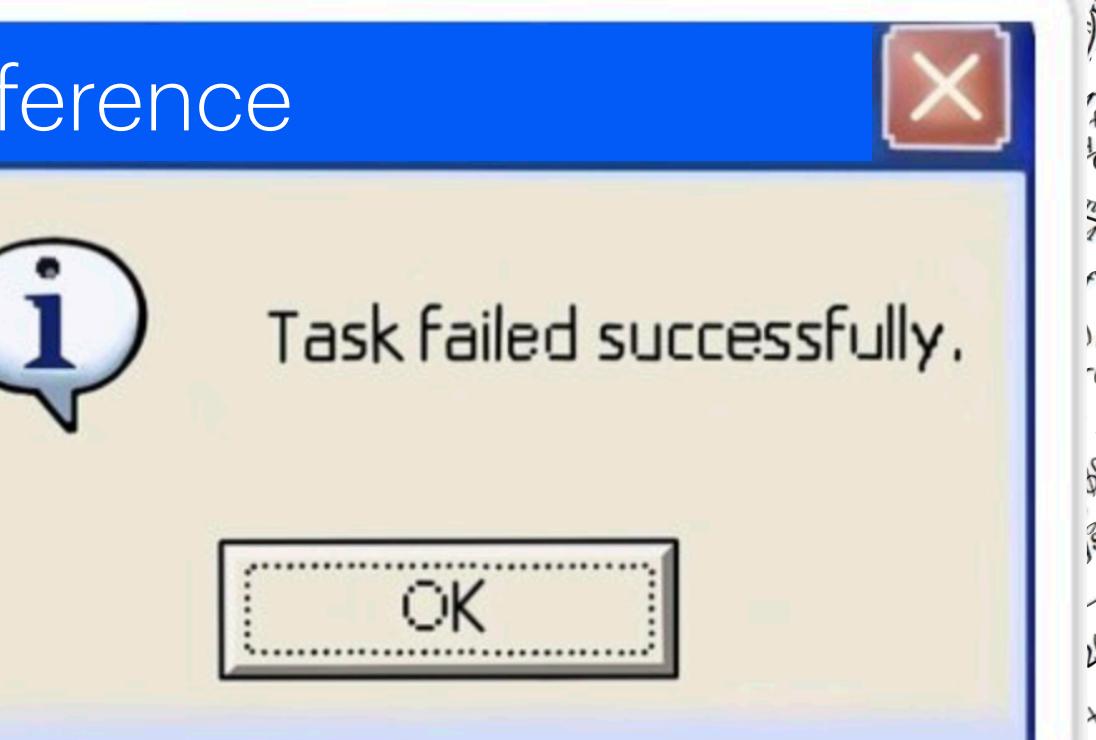
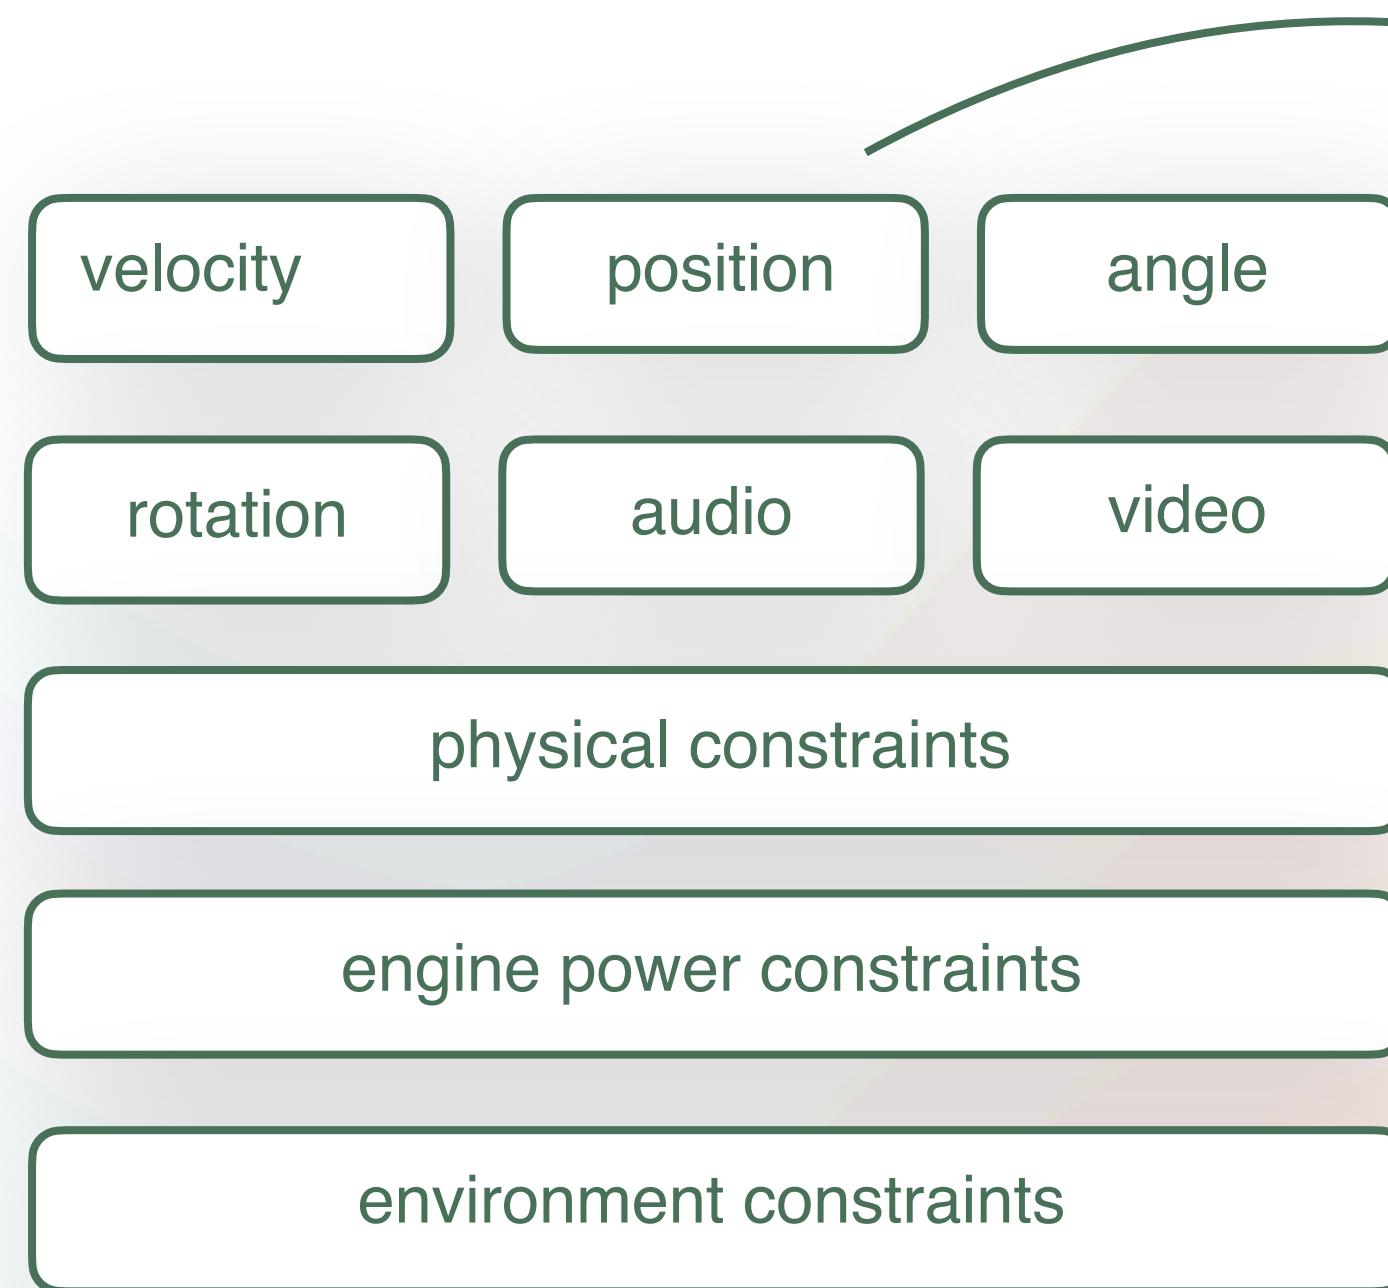
$$\begin{aligned} & Z_\mu^0 Z_\nu^0 W_\mu^+ W_\nu^- + g^2 s_w^2 (A_\mu W_\mu^+ A_\nu W_\nu^- - A_\mu A_\nu W_\mu^+ W_\nu^-) + g^2 s_w c_w (A_\mu Z_\nu^0 (W_\mu^+ W_\nu^- - \\ & W_\nu^+ W_\mu^-) - 2 A_\mu Z_\mu^0 W_\nu^+ W_\nu^-) - \frac{1}{2} \partial_\mu H \partial_\mu H - 2 M^2 \alpha_h H^2 - \partial_\mu \phi^+ \partial_\mu \phi^- - \frac{1}{2} \partial_\mu \phi^0 \partial_\mu \phi^0 - \\ & \beta_h \left( \frac{2M^2}{g^2} + \frac{2M}{g} H + \frac{1}{2} (H^2 + \phi^0 \phi^0 + 2\phi^+ \phi^-) \right) + \frac{2M^4}{g^2} \alpha_h - \\ & g \alpha_h M (H^3 + H \phi^0 \phi^0 + 2H \phi^+ \phi^-) - \\ & \frac{1}{8} g^2 \alpha_h (H^4 + (\phi^0)^4 + 4(\phi^+ \phi^-)^2 + 4(\phi^0)^2 \phi^+ \phi^- + 4H^2 \phi^+ \phi^- + 2(\phi^0)^2 H^2) - \\ & g M W_\mu^+ W_\mu^- H - \frac{1}{2} g \frac{M}{c_w^2} Z_\mu^0 H - \\ & \frac{1}{2} i g (W_\mu^+ (\phi^0 \partial_\mu \phi^- - \phi^- \partial_\mu \phi^0) - W_\mu^- (\phi^0 \partial_\mu \phi^+ - \phi^+ \partial_\mu \phi^0)) + \\ & \frac{1}{2} g (W_\mu^+ (H \partial_\mu \phi^- - \phi^- \partial_\mu H) + W_\mu^- (H \partial_\mu \phi^+ - \phi^+ \partial_\mu H)) + \frac{1}{2} g \frac{1}{c_w} (Z_\mu^0 (H \partial_\mu \phi^0 - \phi^0 \partial_\mu H) + \\ & M (\frac{1}{c_w} Z_\mu^0 \partial_\mu \phi^0 + W_\mu^+ \partial_\mu \phi^- + W_\mu^- \partial_\mu \phi^+) - i g \frac{s_w}{c_w} M Z_\mu^0 (W_\mu^+ \phi^- - W_\mu^- \phi^+) + i g s_w M A_\mu (W_\mu^+ \phi^- - \\ & W_\mu^- \phi^+) - i g \frac{1-2c_w}{2c_w} Z_\mu^0 (\phi^+ \partial_\mu \phi^- - \phi^- \partial_\mu \phi^+) + i g s_w A_\mu (\phi^+ \partial_\mu \phi^- - \phi^- \partial_\mu \phi^+) - \\ & \frac{1}{4} g^2 W_\mu^+ W_\mu^- (H^2 + (\phi^0)^2 + 2\phi^+ \phi^-) - \frac{1}{8} g^2 \frac{1}{c_w^2} Z_\mu^0 (H^2 + (\phi^0)^2 + 2(2s_w^2 - 1)^2 \phi^+ \phi^-) - \\ & \frac{1}{2} g^2 \frac{s_w}{c_w} Z_\mu^0 \phi^0 (W_\mu^+ \phi^- + W_\mu^- \phi^+) - \frac{1}{2} i g^2 \frac{s_w}{c_w} Z_\mu^0 H (W_\mu^+ \phi^- - W_\mu^- \phi^+) + \frac{1}{2} g^2 s_w A_\mu \phi^0 (W_\mu^+ \phi^- + \\ & W_\mu^- \phi^+) + \frac{1}{2} i g^2 s_w A_\mu H (W_\mu^+ \phi^- - W_\mu^- \phi^+) - g^2 \frac{s_w}{c_w} (2c_w^2 - 1) Z_\mu^0 A_\mu \phi^+ \phi^- - \\ & g^2 s_w^2 A_\mu A_\mu \phi^+ \phi^- + \frac{1}{2} i g_s \lambda_{ij}^a (q_i^\sigma \gamma^\mu q_j^\sigma) g_{\mu\nu}^a - \bar{e}^\lambda (\gamma \partial + m_e^\lambda) e^\lambda - \bar{\nu}^\lambda (\gamma \partial + m_v^\lambda) \nu^\lambda - \bar{u}_j^\lambda (\gamma \partial + \\ & m_u^\lambda) u_j^\lambda - \bar{d}_j^\lambda (\gamma \partial + m_d^\lambda) d_j^\lambda + i g s_w A_\mu ((-\bar{e}^\lambda \gamma^\mu e^\lambda) + \frac{2}{3} (\bar{u}_j^\lambda \gamma^\mu u_j^\lambda) - \frac{1}{3} (\bar{d}_j^\lambda \gamma^\mu d_j^\lambda)) + \\ & \frac{i g}{4 c_w} Z_\mu^0 \{ (\bar{\nu}^\lambda \gamma^\mu (1 + \gamma^5) \nu^\lambda) + (\bar{e}^\lambda \gamma^\mu (4s_w^2 - 1 - \gamma^5) e^\lambda) + (d_j^\lambda \gamma^\mu (\frac{4}{3} s_w^2 - 1 - \gamma^5) d_j^\lambda) + \\ & (\bar{u}_j^\lambda \gamma^\mu (1 - \frac{8}{3} s_w^2 + \gamma^5) u_j^\lambda) \} + \frac{i g}{2 \sqrt{2}} W_\mu^+ ((\bar{\nu}^\lambda \gamma^\mu (1 + \gamma^5) U^{lep} \lambda_\kappa e^\kappa) + (\bar{u}_j^\lambda \gamma^\mu (1 + \gamma^5) C_{\lambda\kappa} d_j^\kappa)) + \\ & \frac{i g}{2 \sqrt{2}} W_\mu^- ((\bar{e}^\kappa U^{lep} \lambda_\kappa \gamma^\mu (1 + \gamma^5) \nu^\lambda) + (\bar{d}_j^\kappa C_{\lambda\kappa}^\dagger \gamma^\mu (1 + \gamma^5) u_j^\lambda)) + \\ & \frac{i g}{2 M \sqrt{2}} \phi^+ (-m_e^\kappa (\bar{e}^\lambda U^{lep} \lambda_\kappa (1 - \gamma^5) e^\kappa) + m_v^\kappa (\bar{\nu}^\lambda U^{lep} \lambda_\kappa (1 + \gamma^5) e^\kappa) + \\ & \frac{i g}{2 M \sqrt{2}} \phi^- (m_e^\lambda (\bar{e}^\lambda U^{lep} \lambda_\kappa (1 + \gamma^5) \nu^\kappa) - m_v^\kappa (\bar{e}^\lambda U^{lep} \lambda_\kappa^\dagger (1 - \gamma^5) \nu^\kappa) - \frac{g}{2} \frac{m_\lambda}{M} H (\bar{\nu}^\lambda \nu^\lambda) - \\ & \frac{g}{2} \frac{m_\lambda}{M} H (\bar{e}^\lambda e^\lambda) + \frac{i g}{2} \frac{m_\lambda}{M} \phi^0 (\bar{\nu}^\lambda \gamma^5 \nu^\lambda) - \frac{i g}{2} \frac{m_\lambda}{M} \phi^0 (\bar{e}^\lambda \gamma^5 e^\lambda) - \frac{1}{4} \bar{\nu}_\lambda M_{\lambda\kappa}^R (1 - \gamma_5) \hat{\nu}_\kappa - \end{aligned}$$

Solve simple math 😳

# Probabilistic programming

Why is it so challenging?

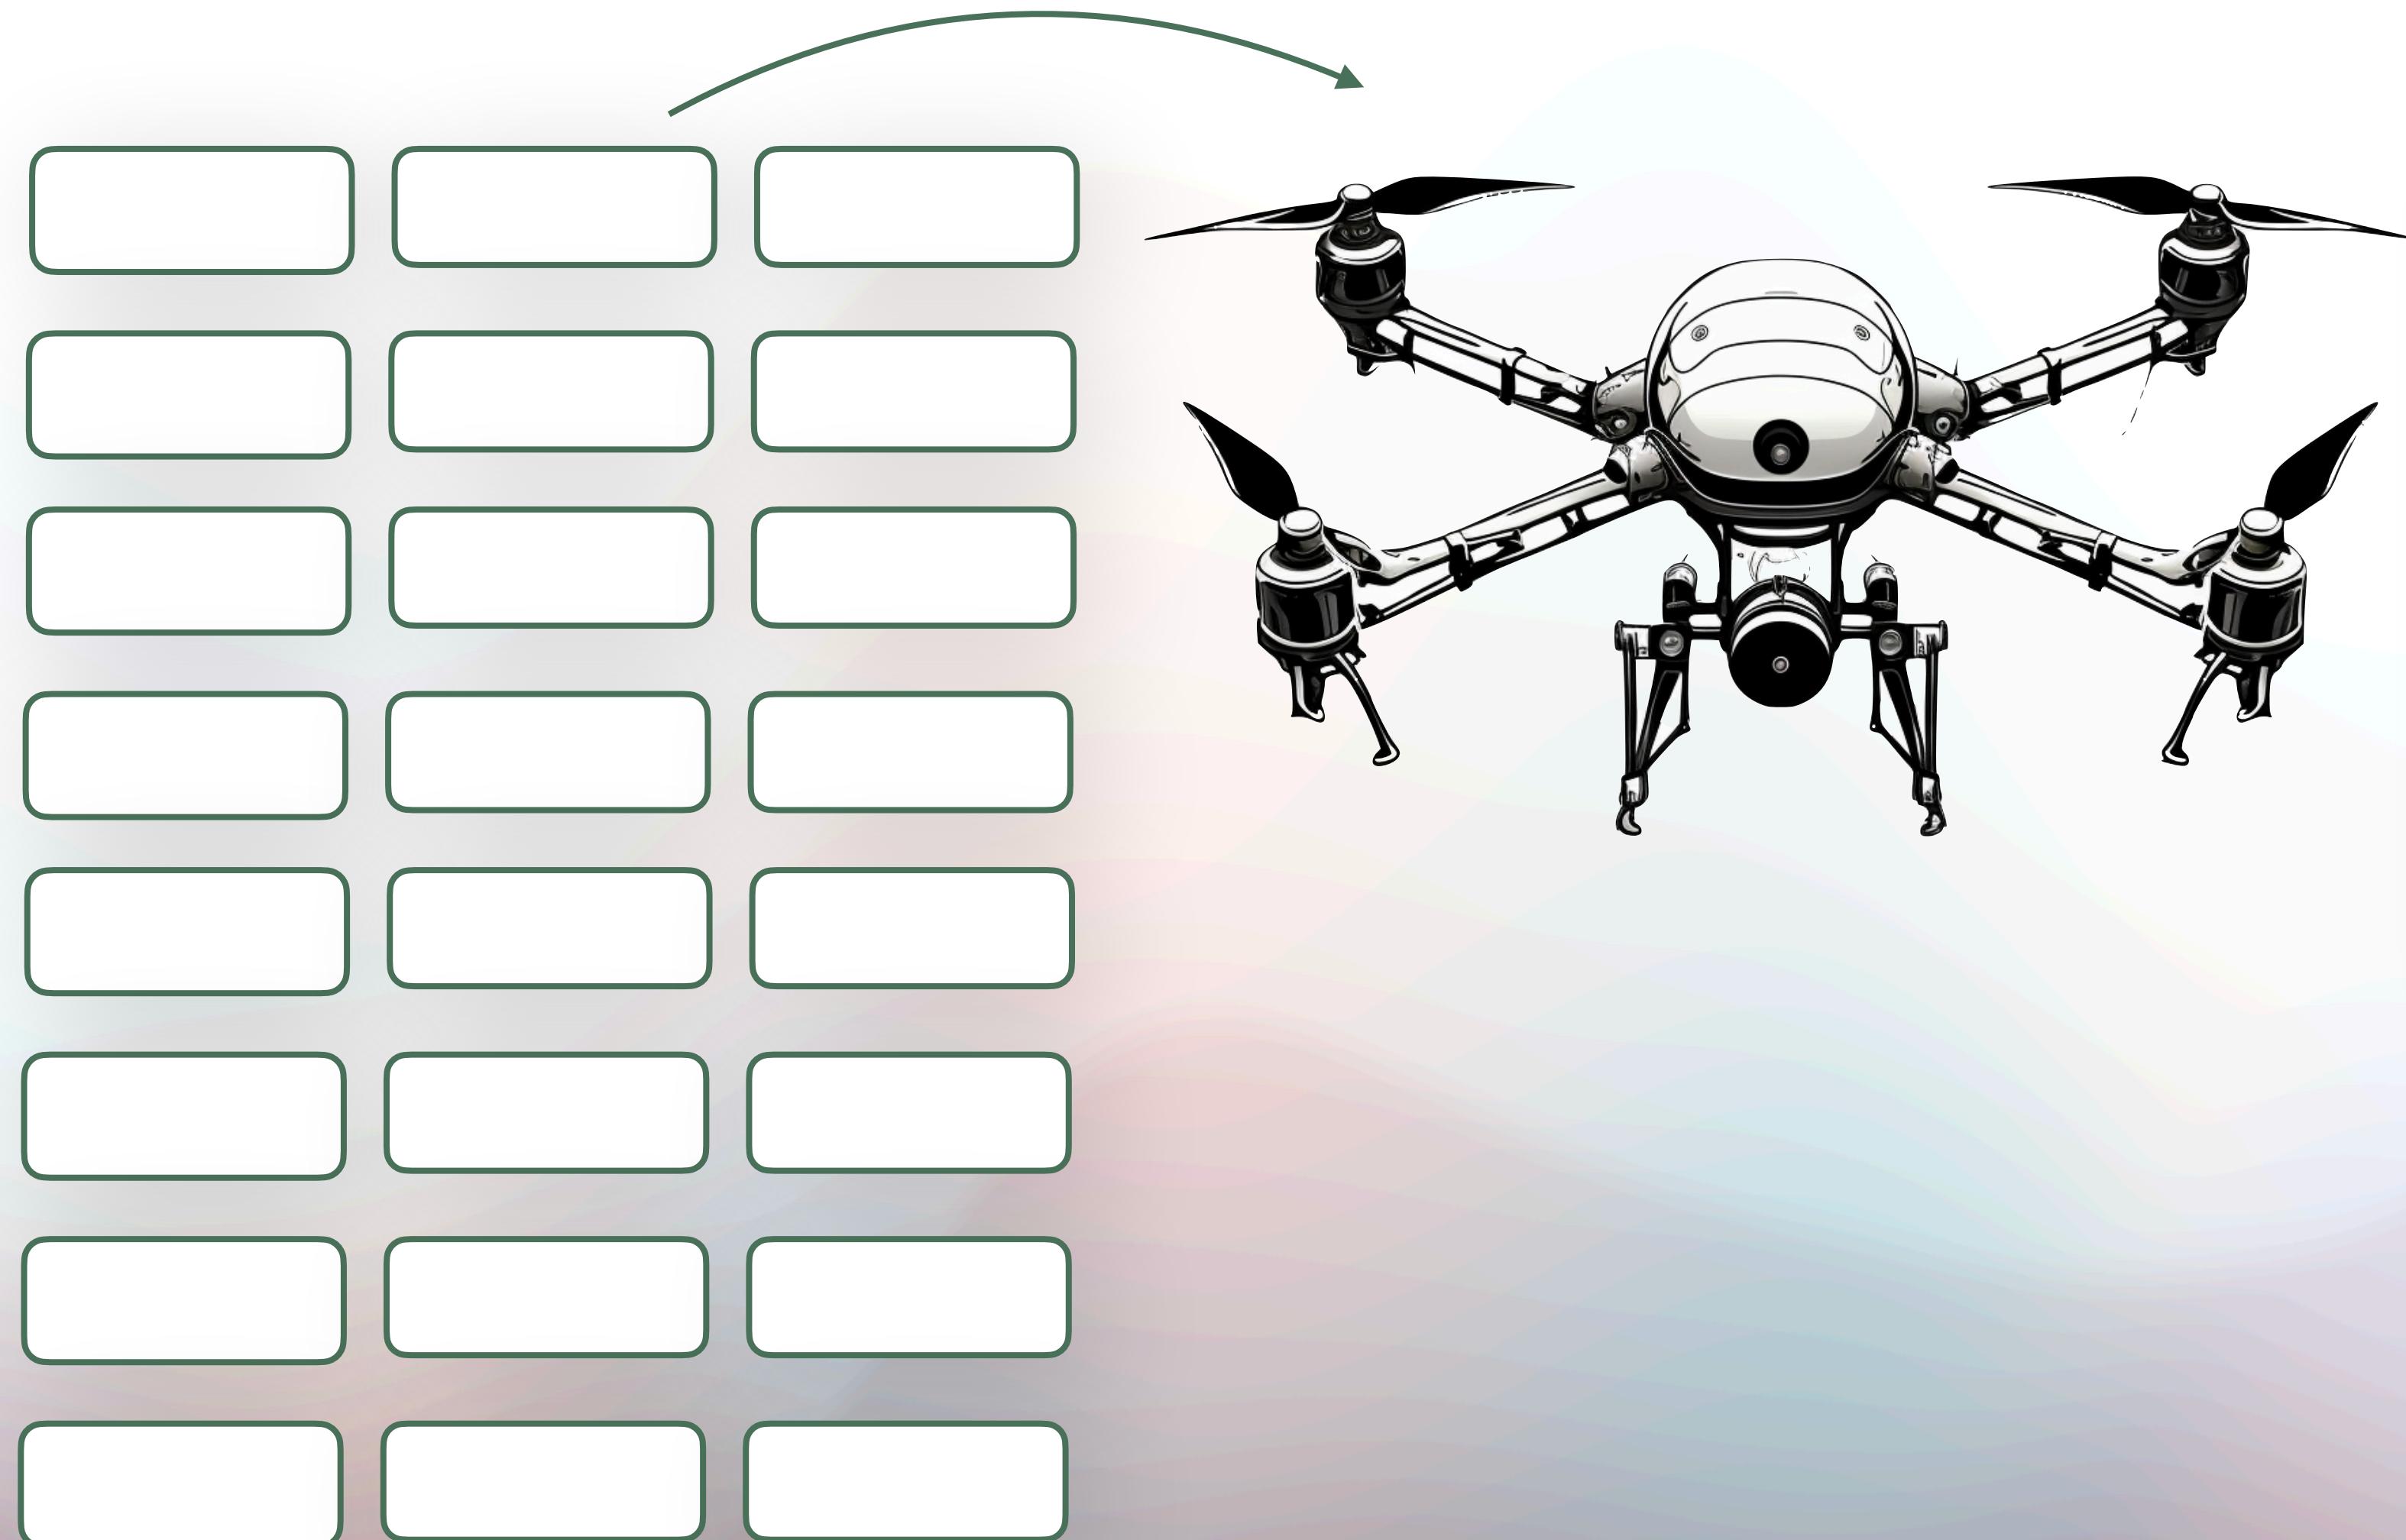
Add even more components



# Probabilistic programming

# Why is it so challenging?

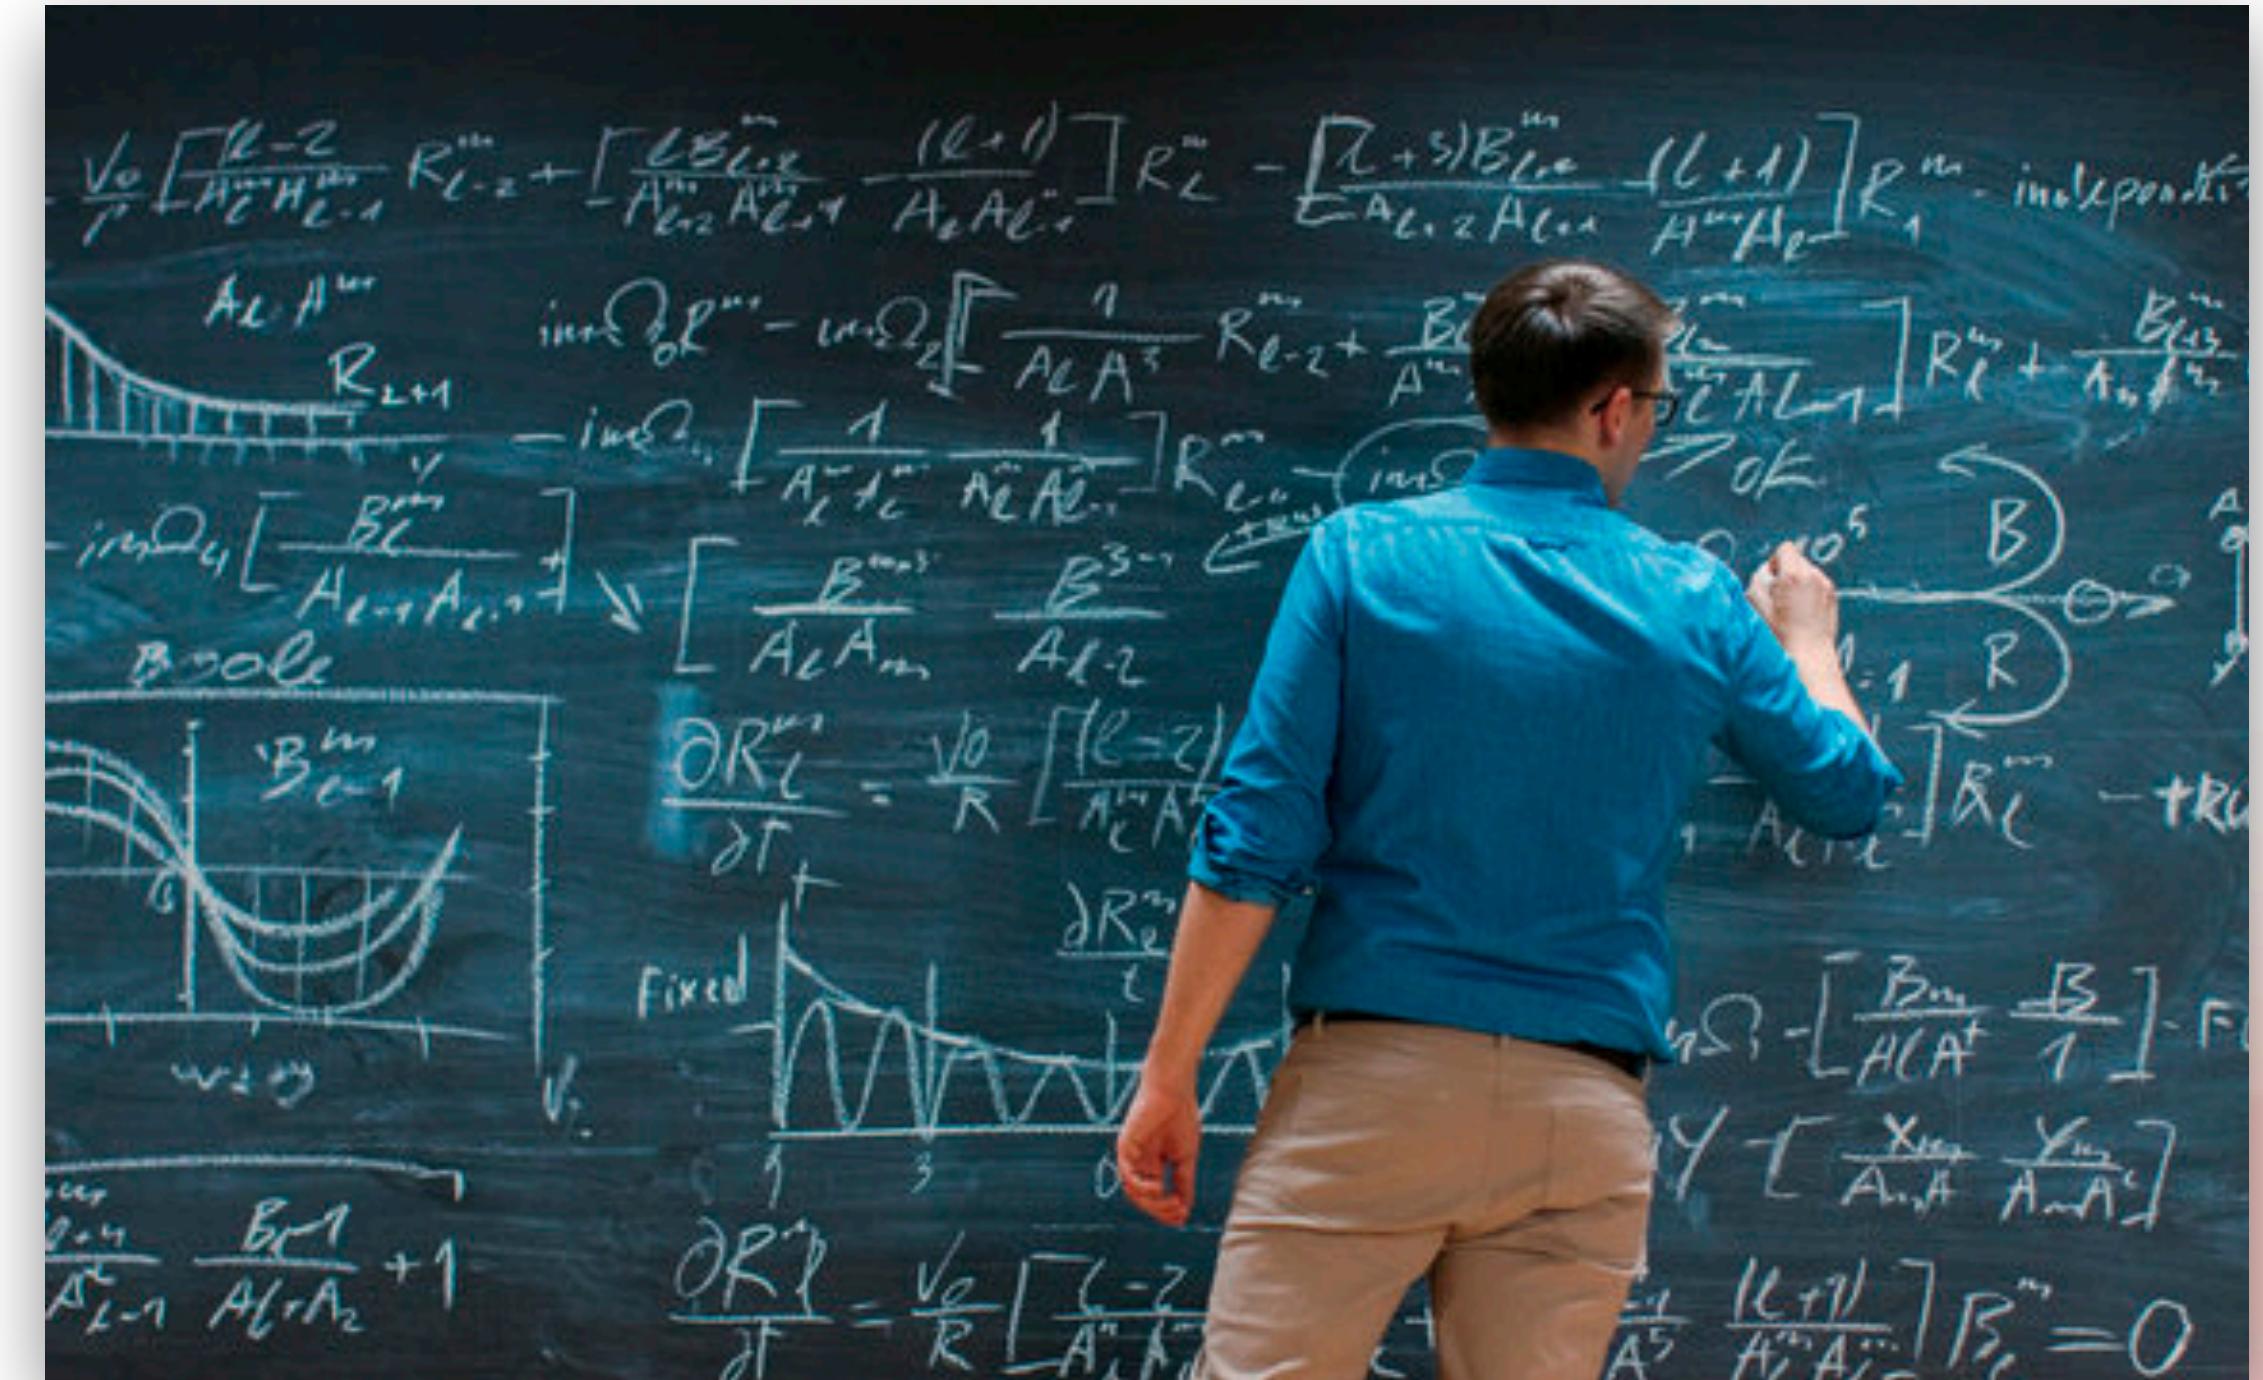
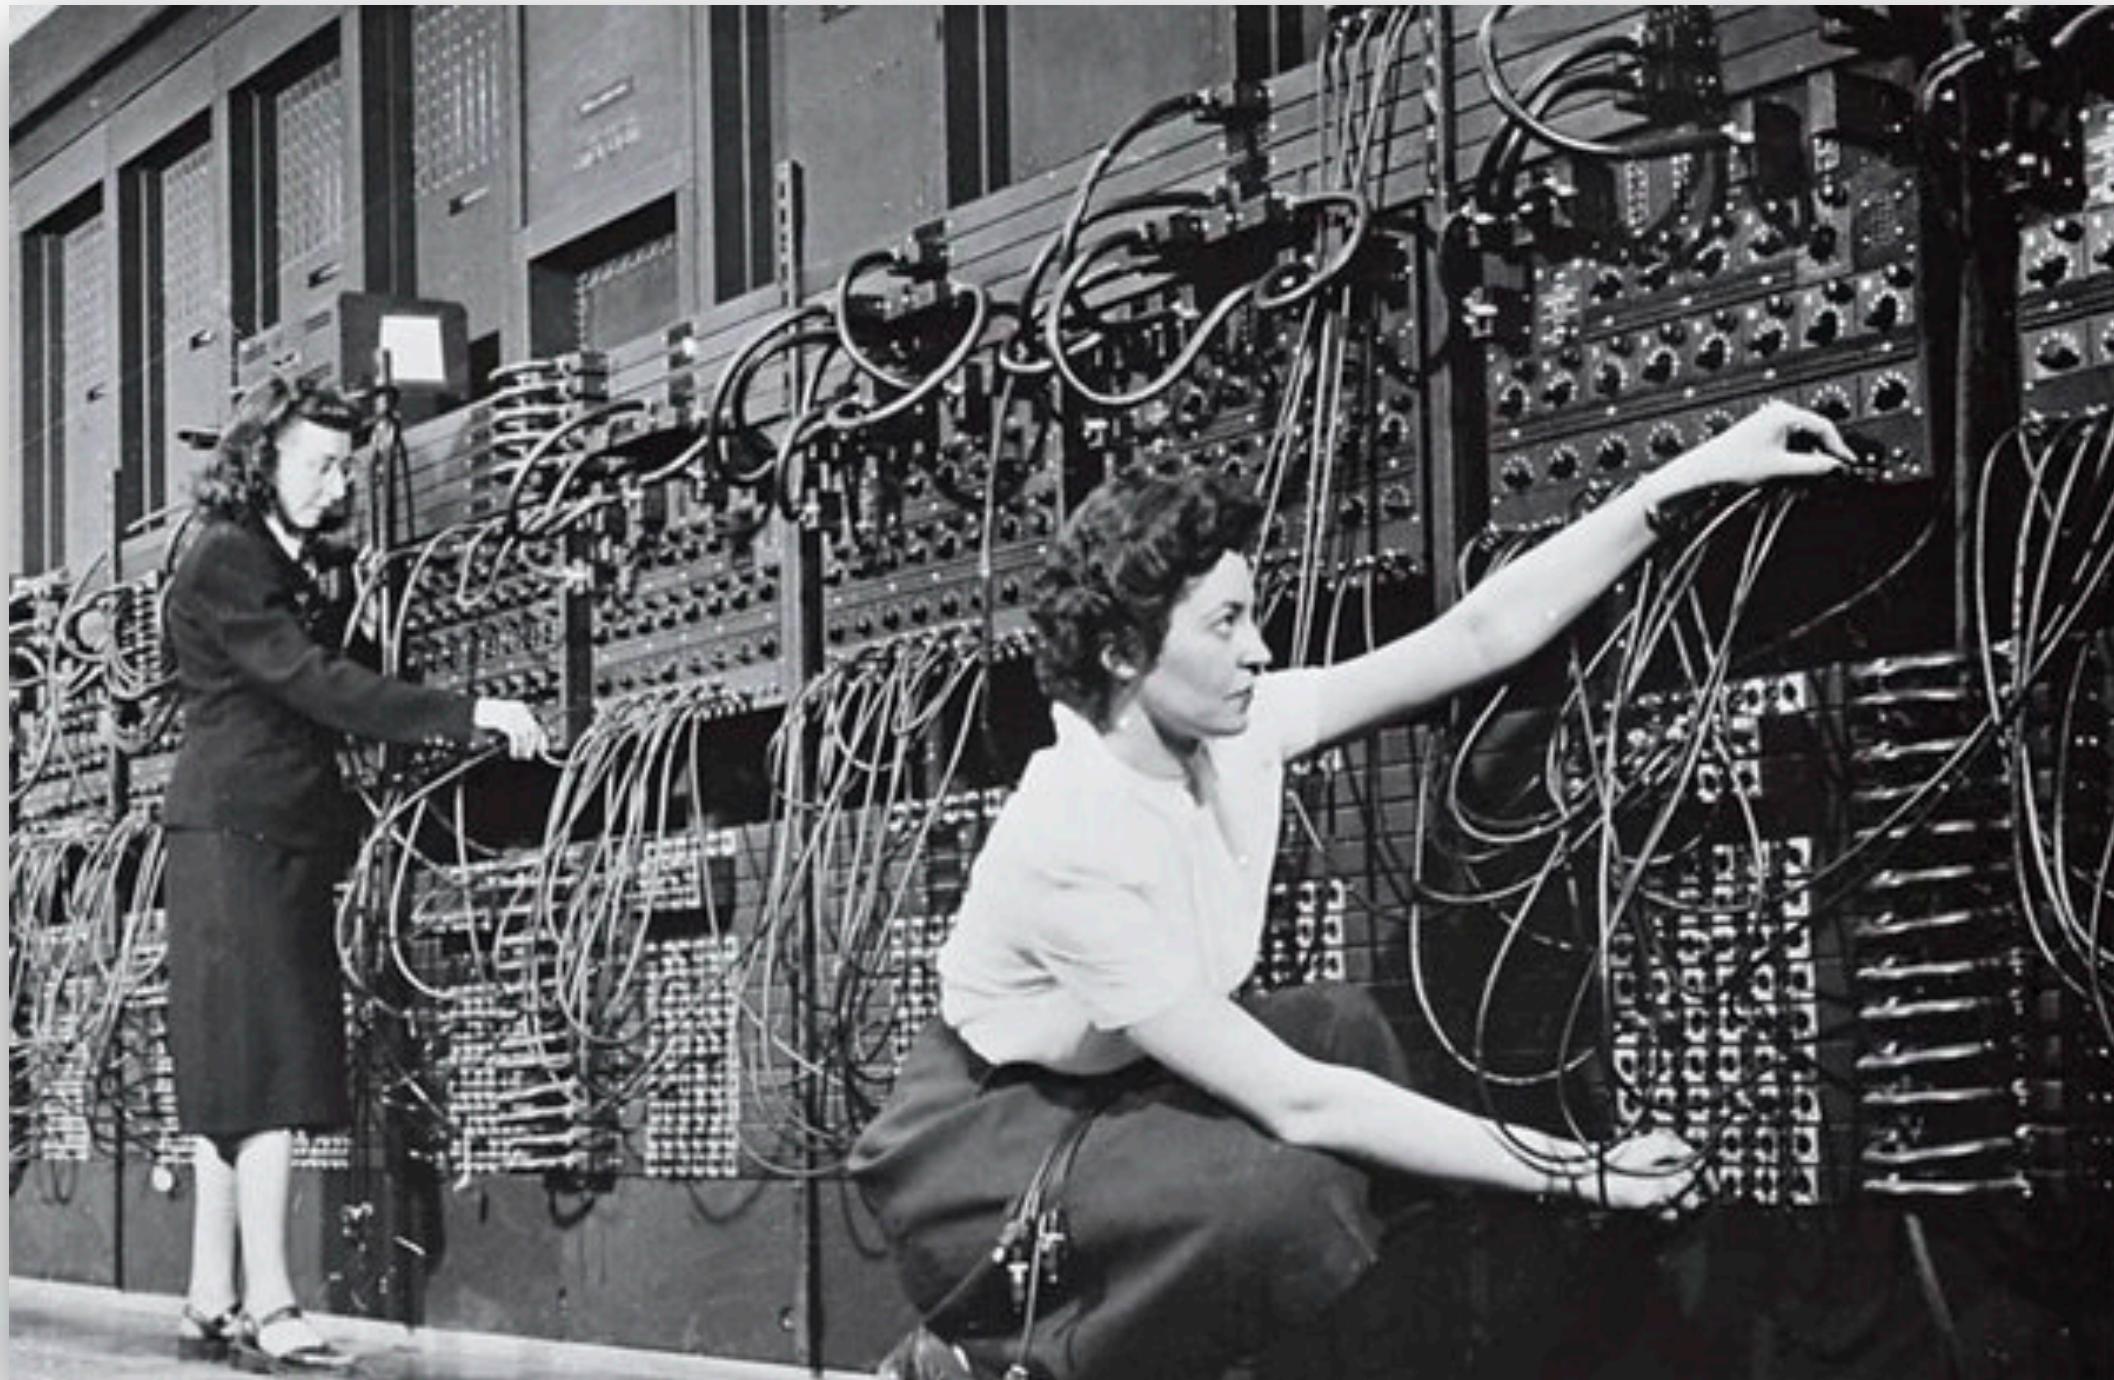
Add as many components as you want



# compile

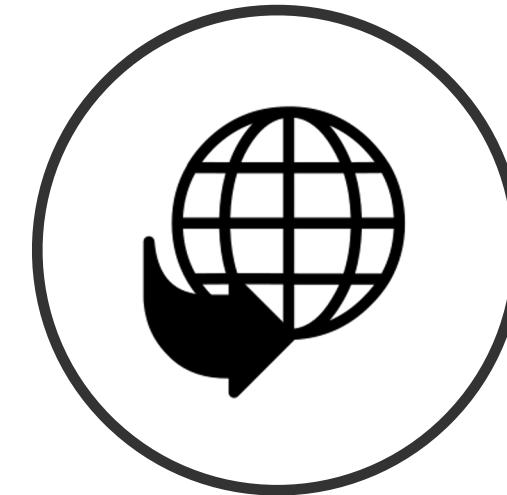
# User-friendliness

Traditional programming and Probabilistic programming

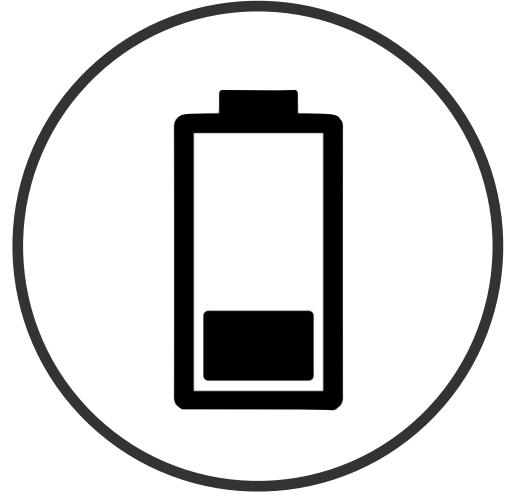


# Probabilistic programming

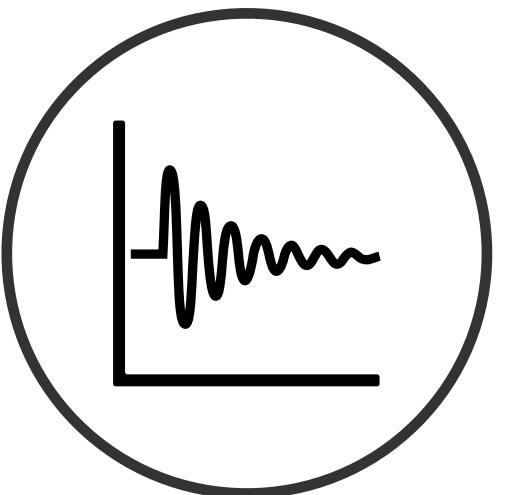
Why is it so challenging?



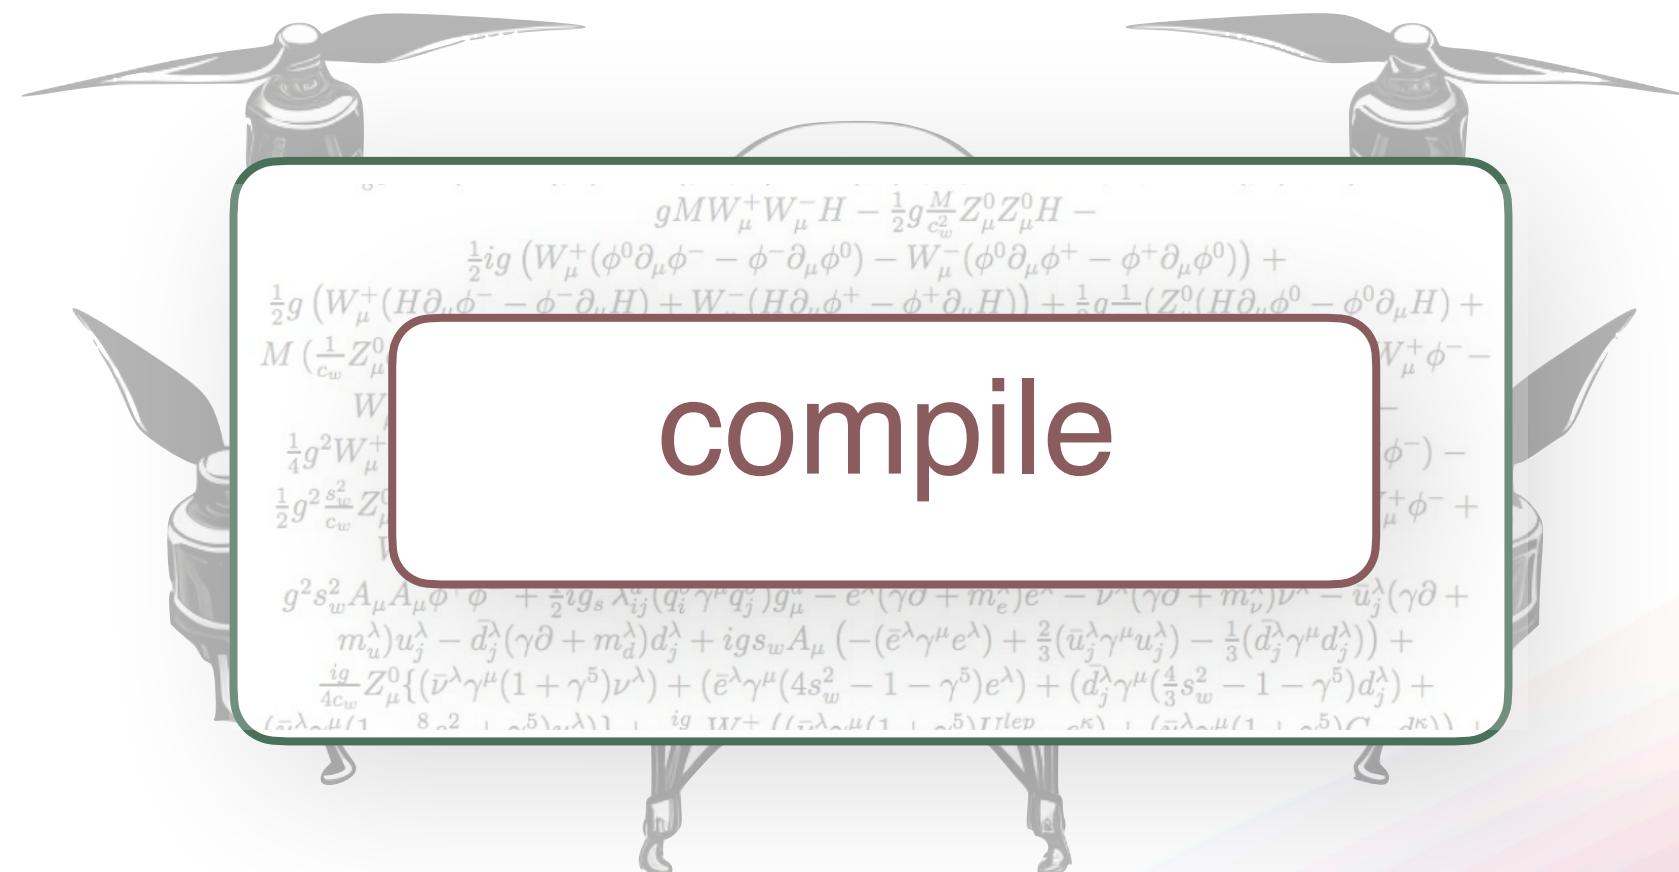
Scalability



Low-power



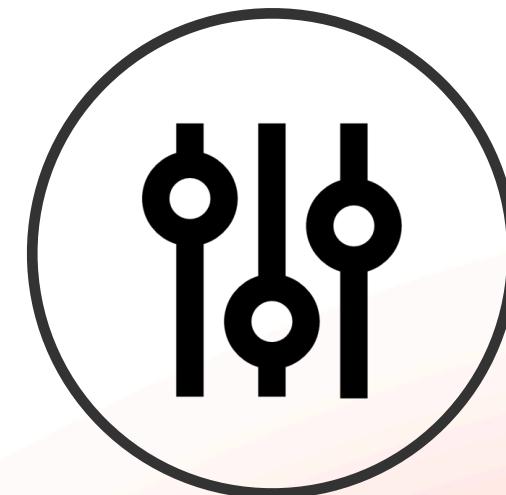
Robustness



Performance



Real-time

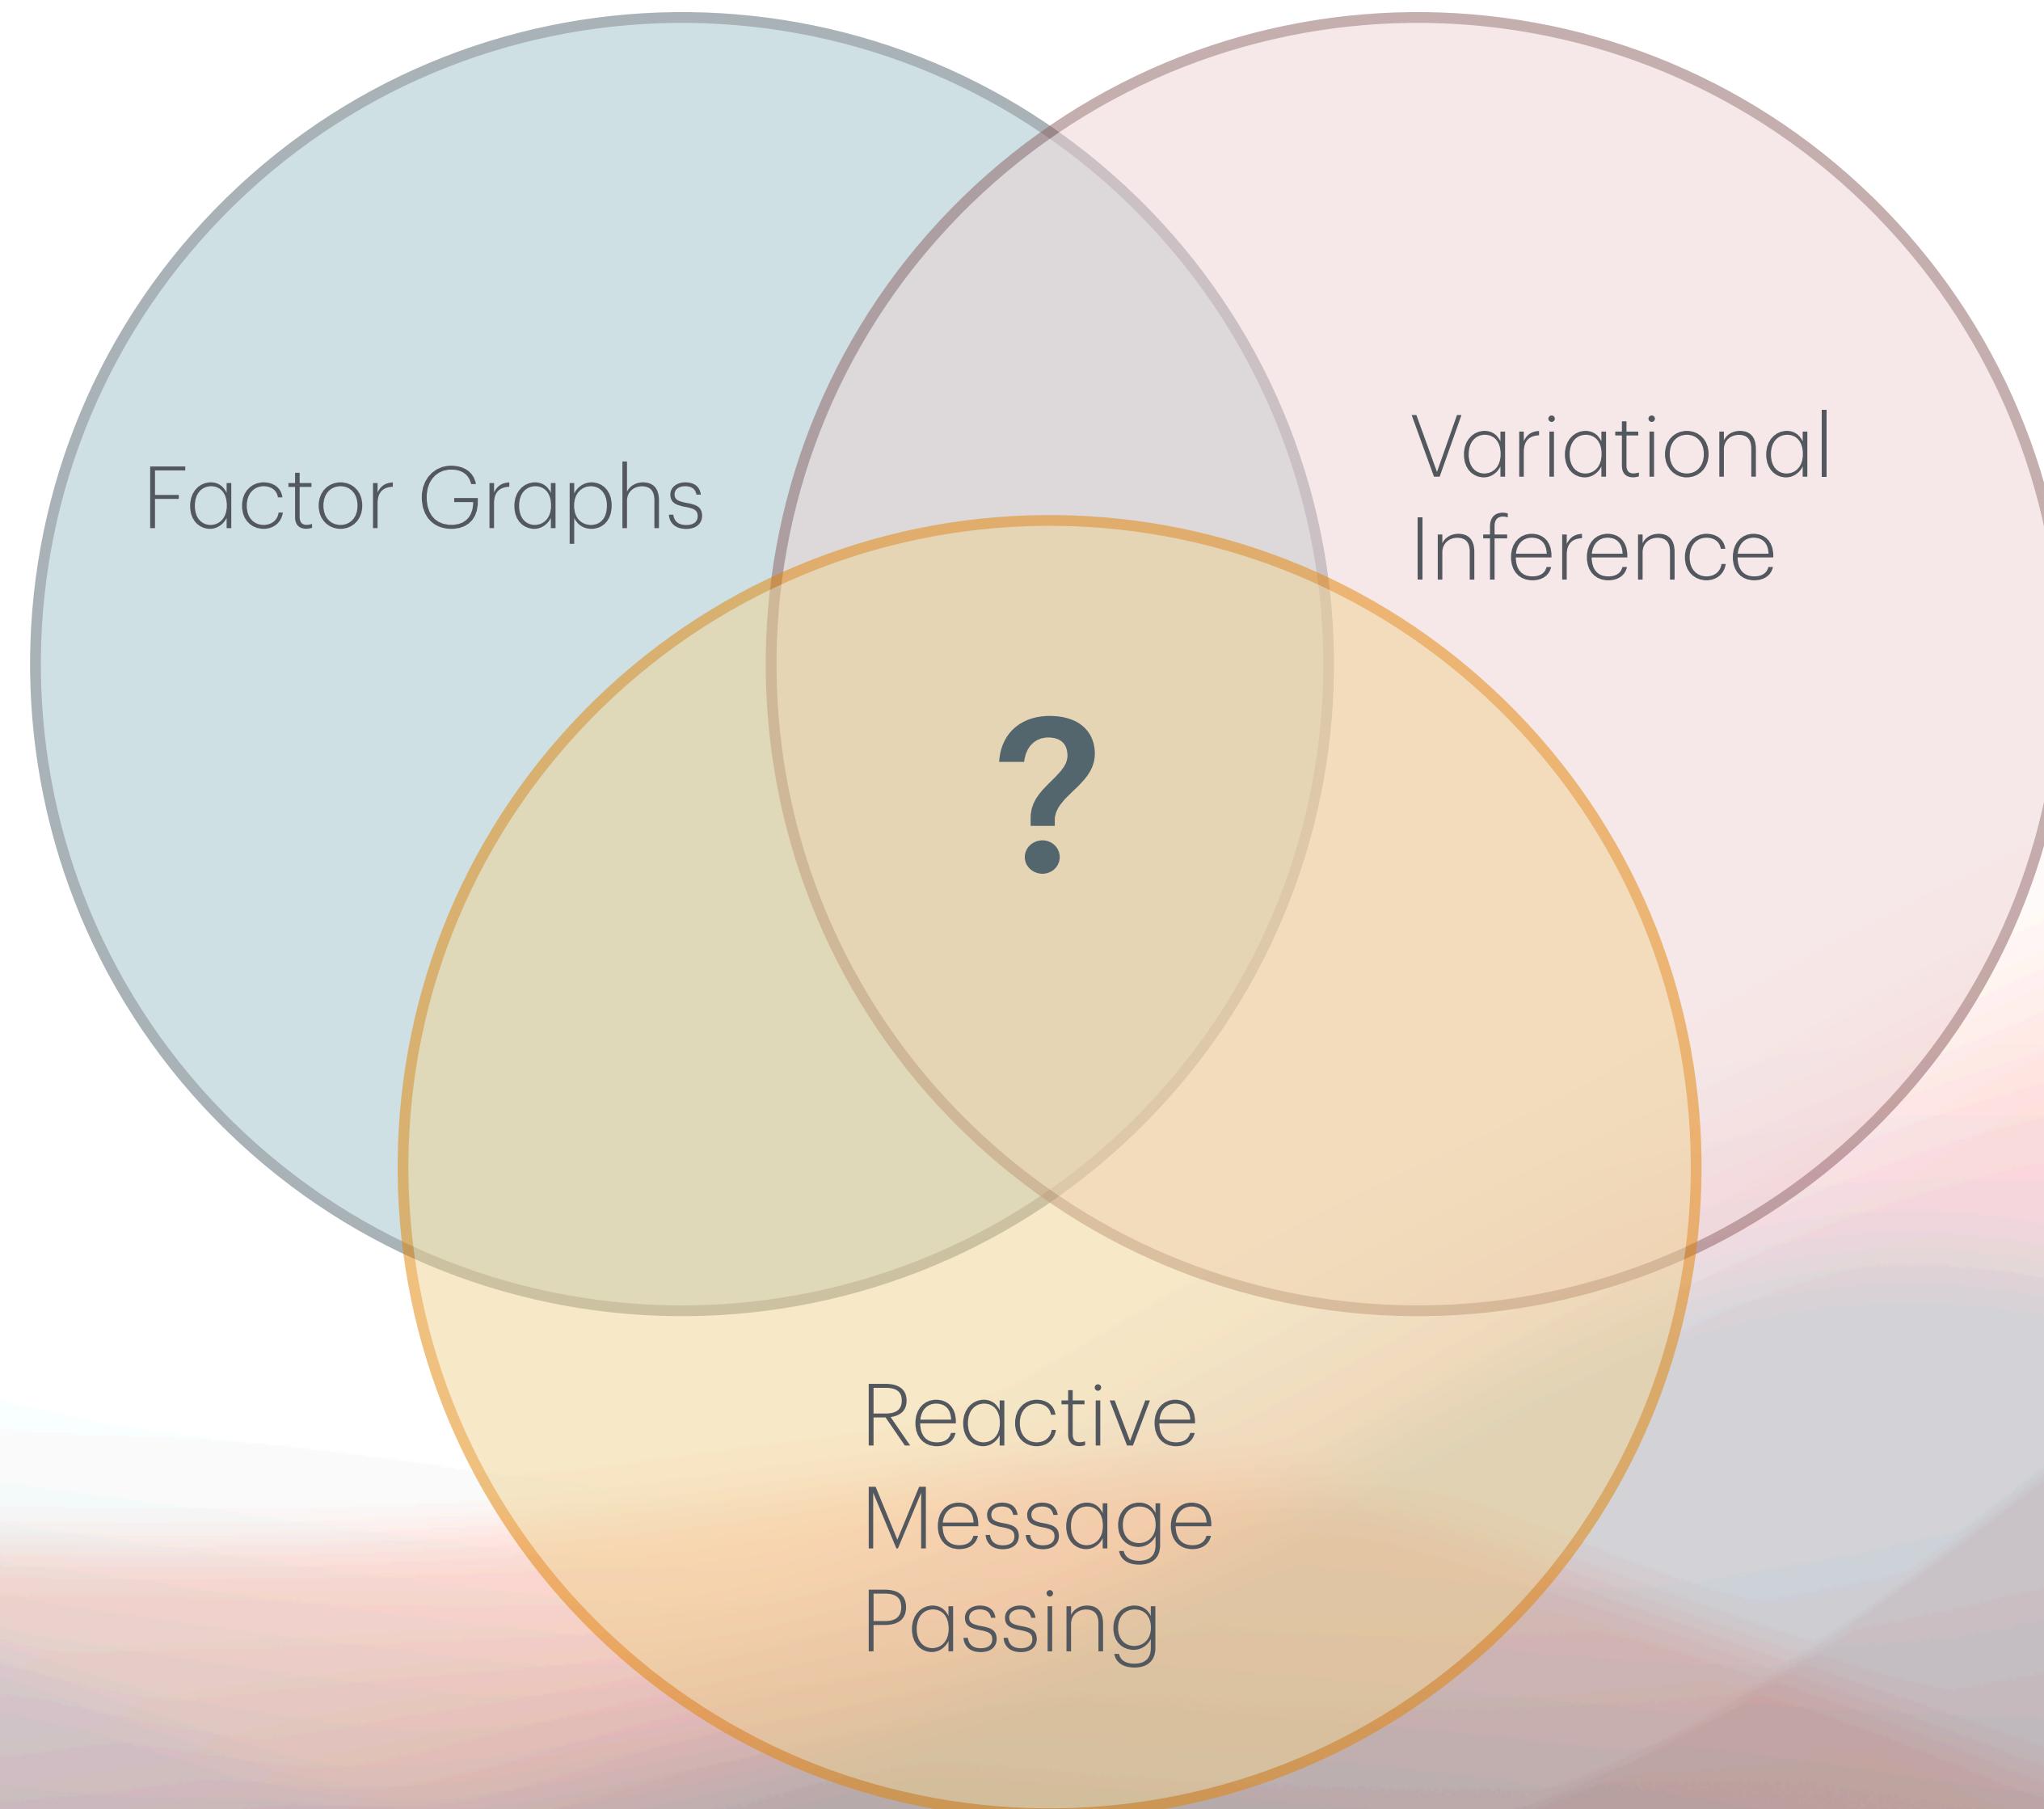


Adaptivity

# How do we approach it?

# How do we approach it?

Scalable Bayesian Inference architecture for autonomous systems



# How do we approach it?

Scalable Bayesian Inference architecture for autonomous systems

## Variational inference

- Scalable and efficient
- Complexity vs accuracy trade-off

$$q^* = \arg \min_{q \in \mathcal{Q}} \int q(s) \log \frac{q(s)}{p(\hat{y}, s)} ds$$

Simpler distributions  
from exponential family

Scalable trade-off between  
complexity and accuracy

# How do we approach it?

Scalable Bayesian Inference architecture for autonomous systems

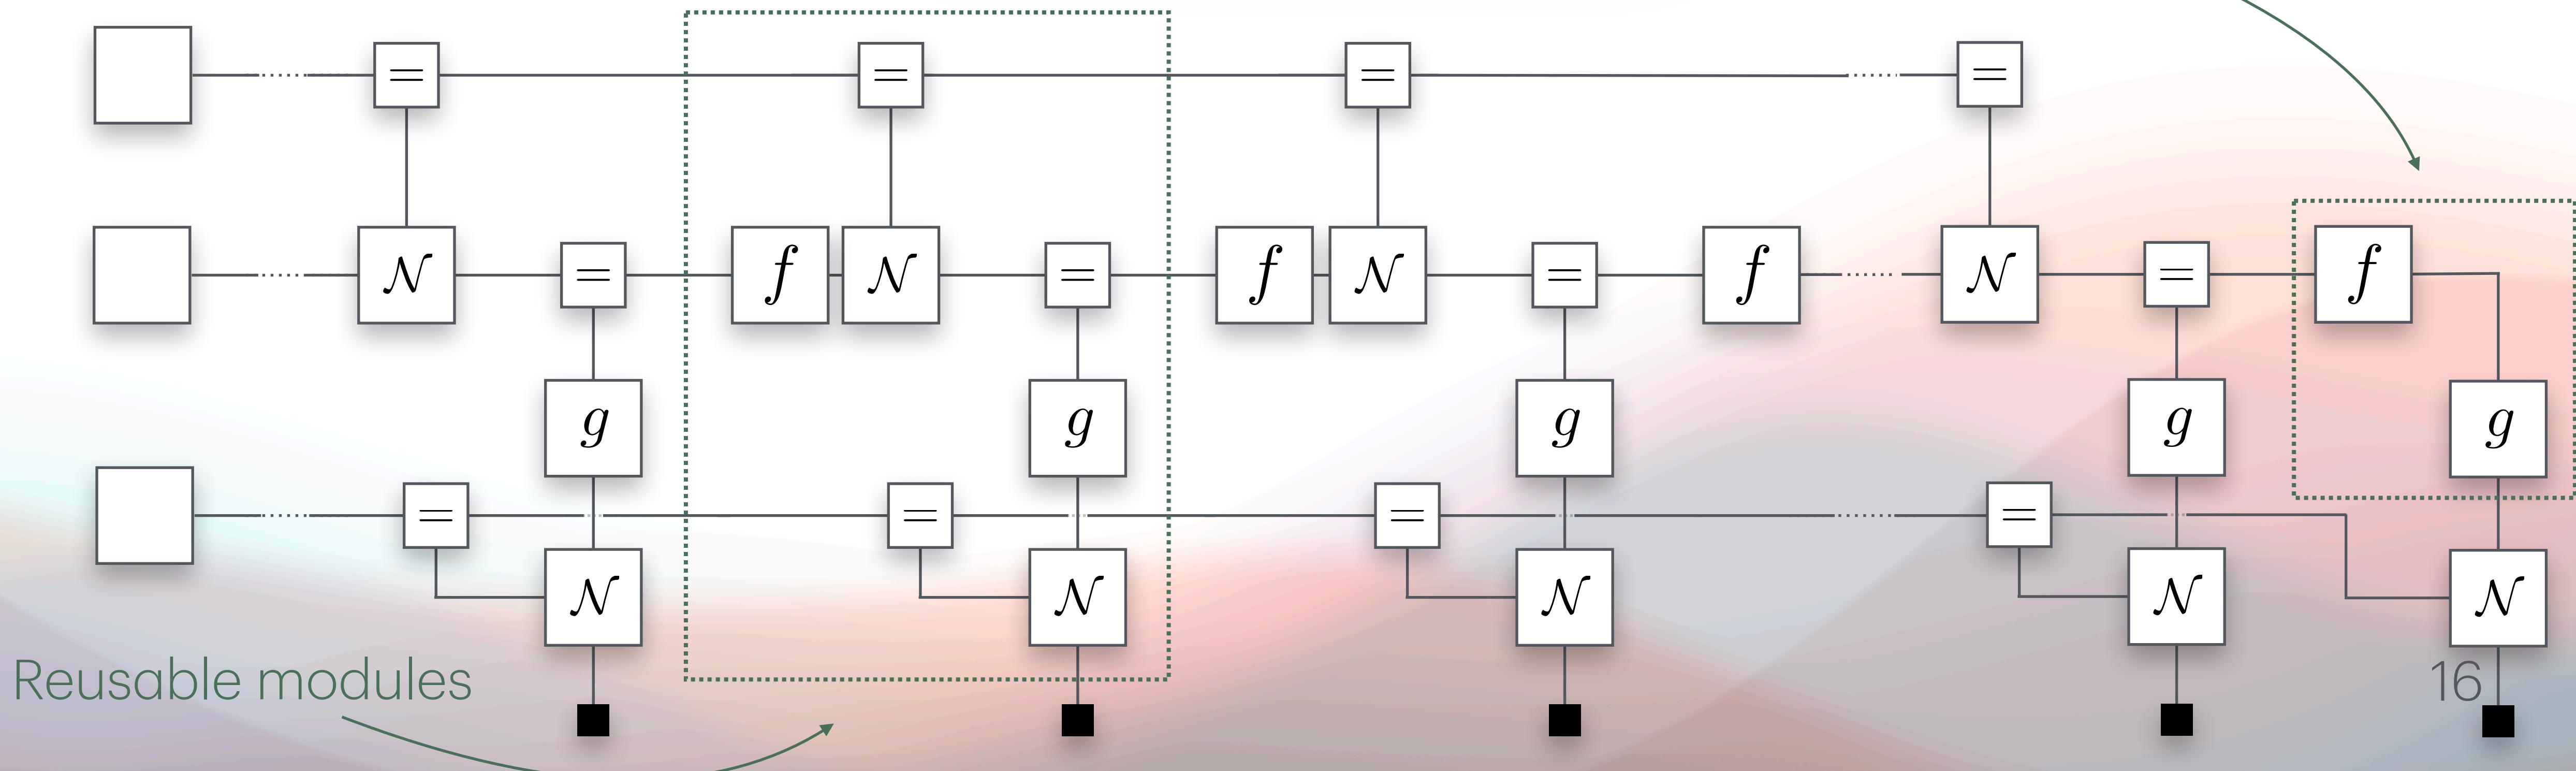
## Variational inference

- Scalable and efficient
- Complexity vs accuracy trade-off

## Factor graphs

- Explainable and not a black-box
- Modular and local

## Local optimisations



# How do we approach it?

Scalable Bayesian Inference architecture for autonomous systems

## Variational inference

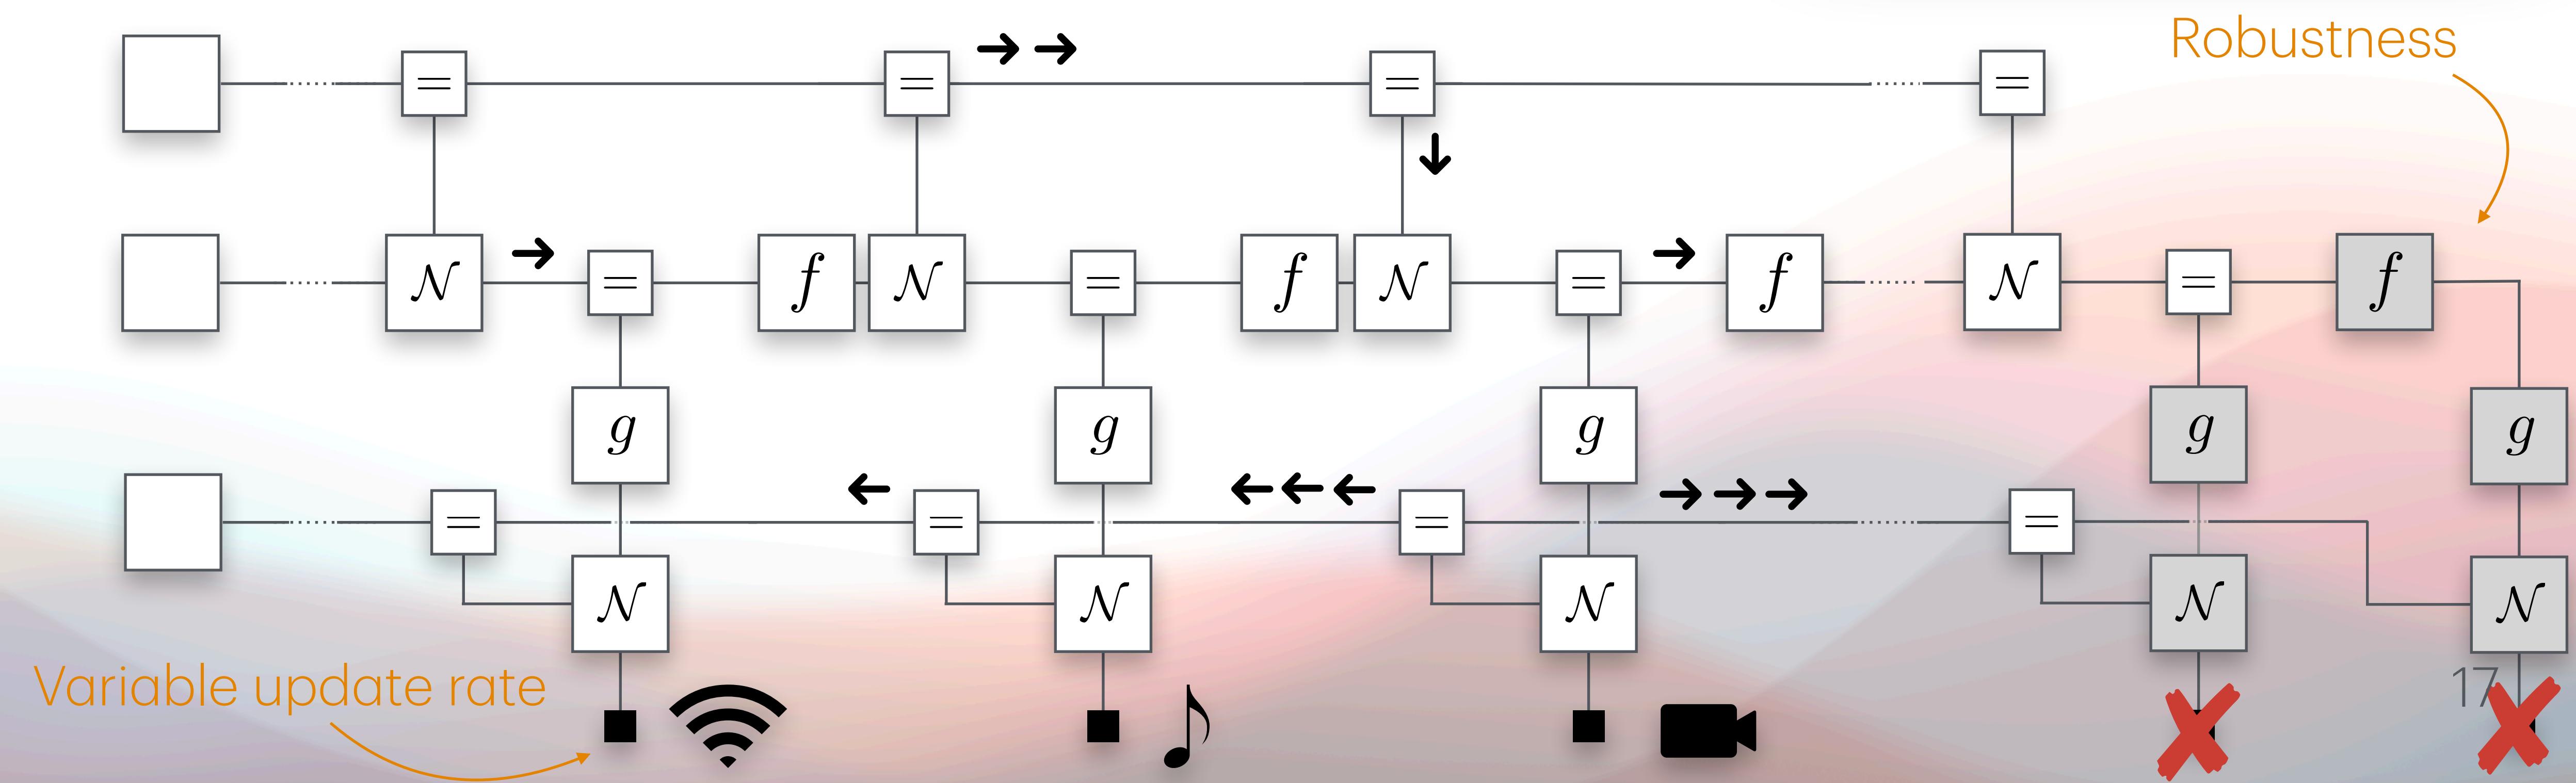
- Scalable and efficient
- Complexity vs accuracy trade-off

## Factor graphs

- Explainable and not a black-box
- Modular and local

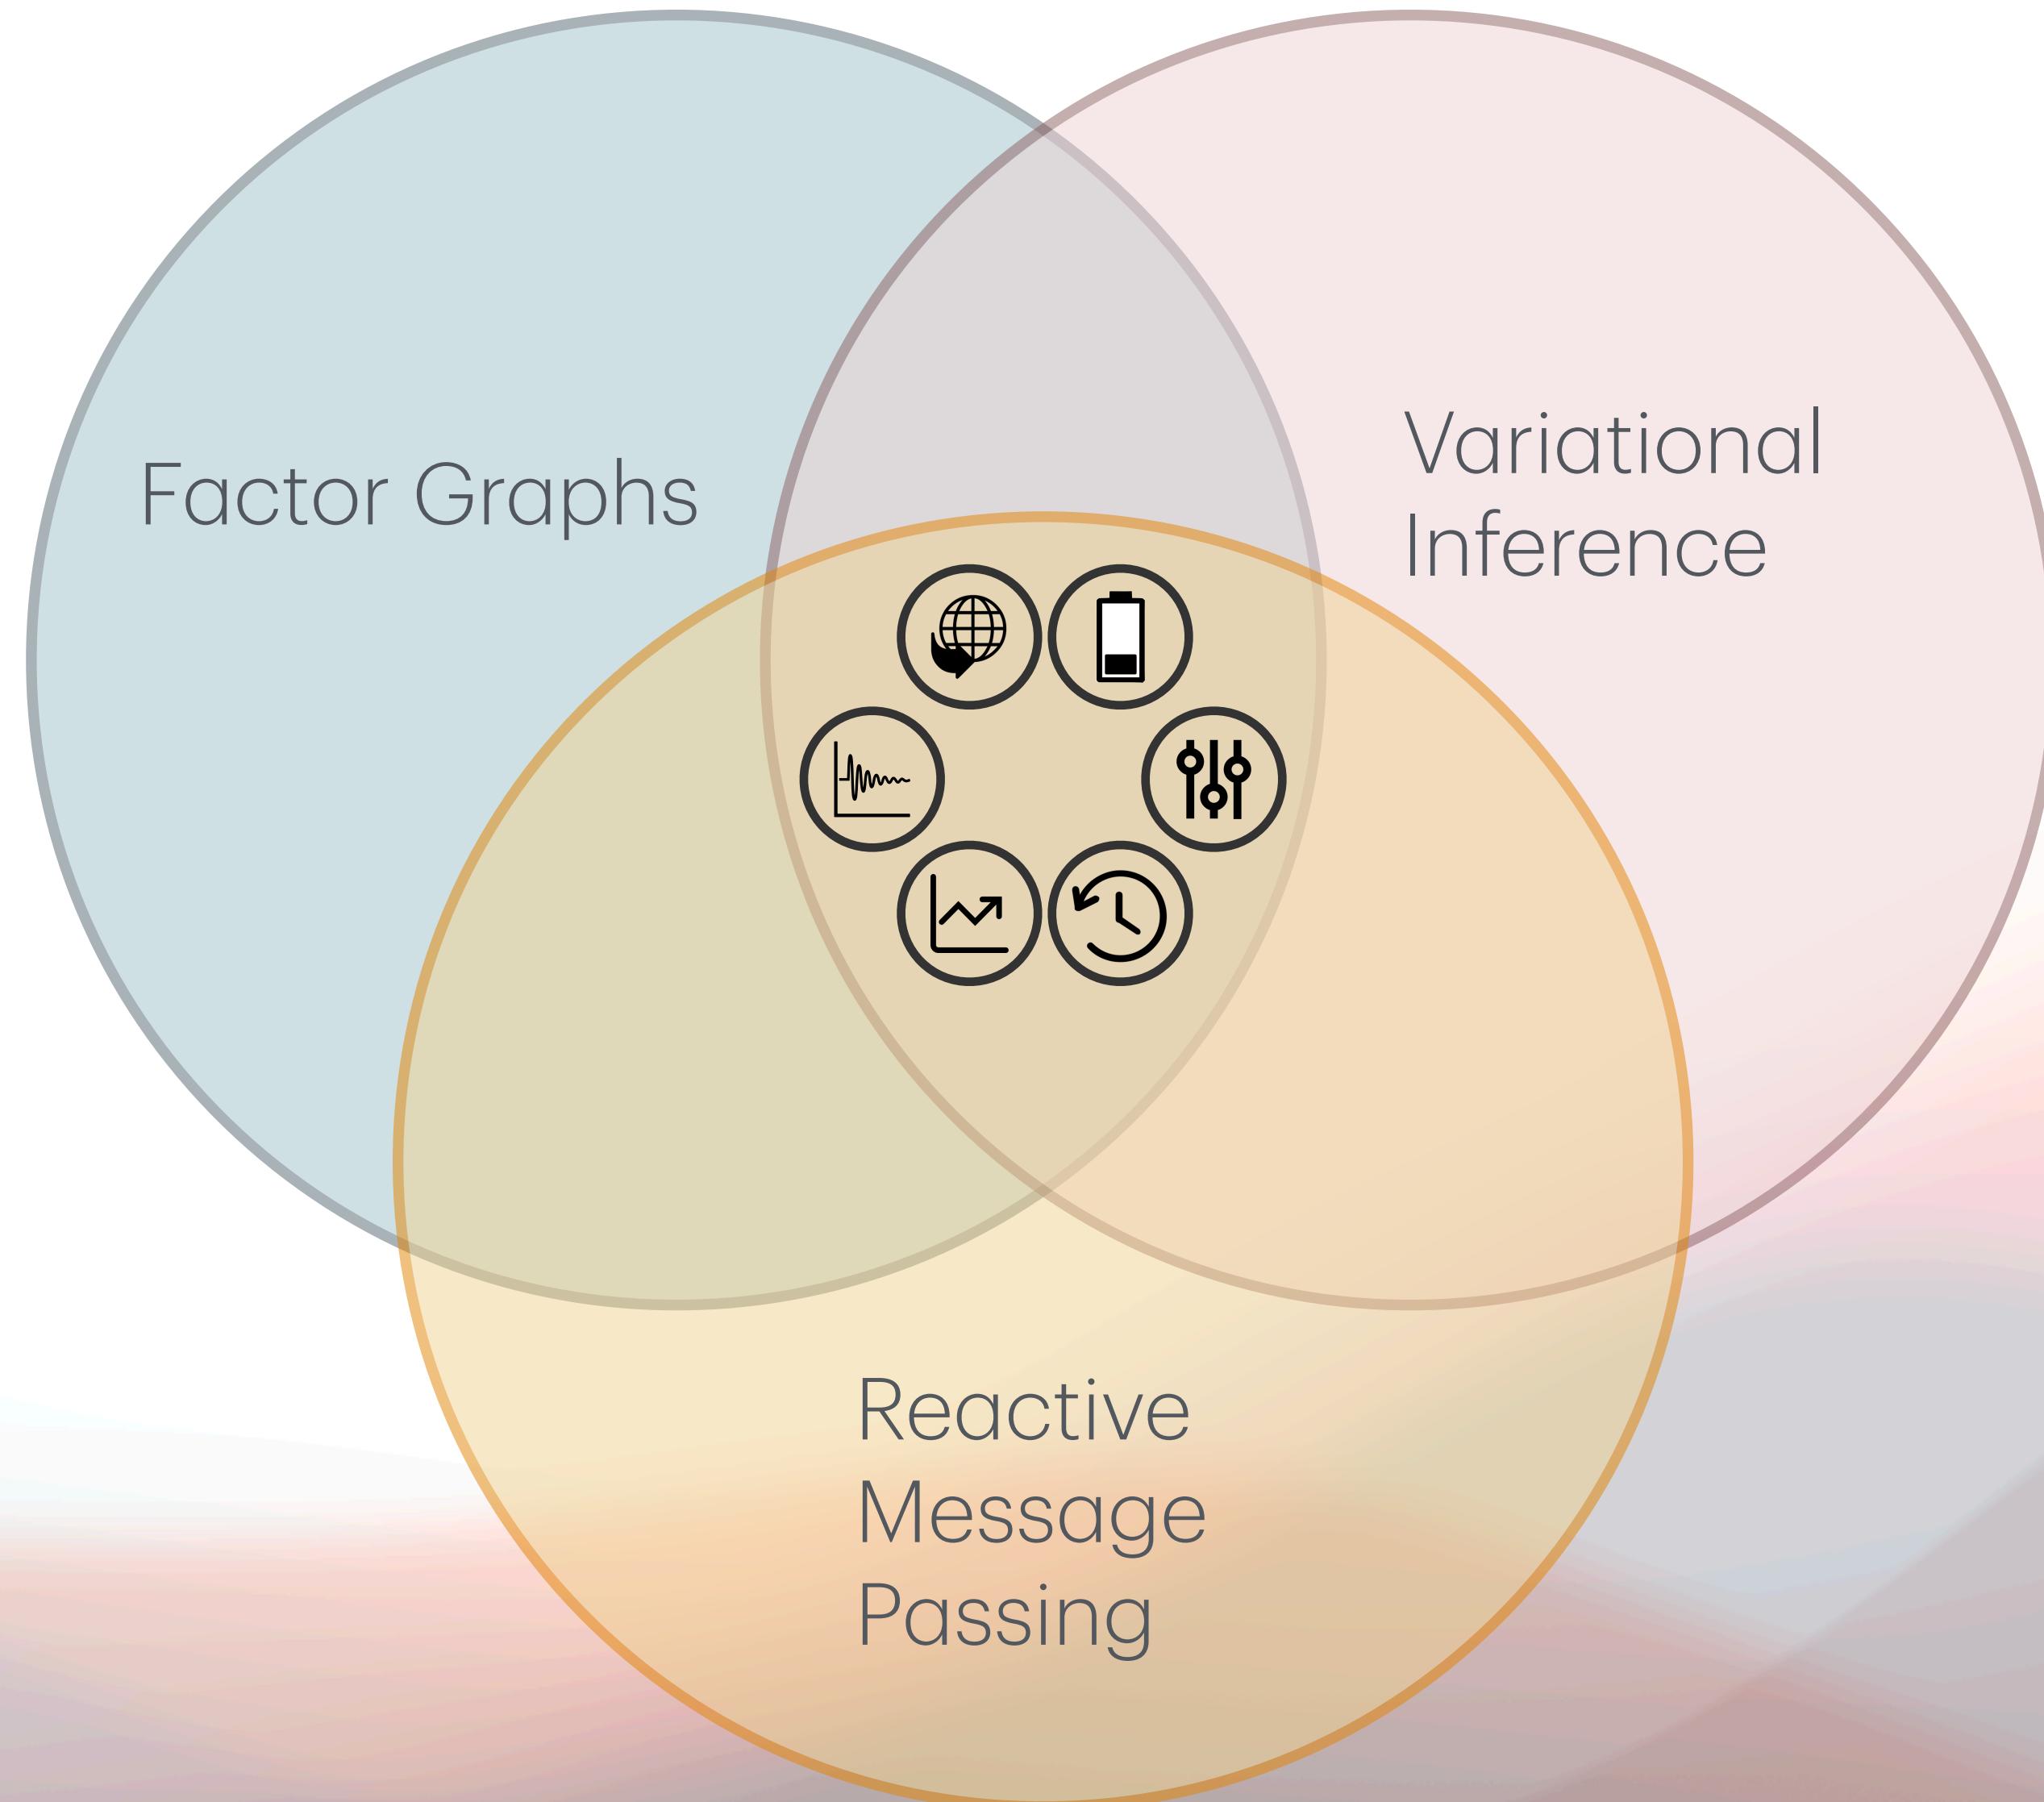
## Reactive Message Passing

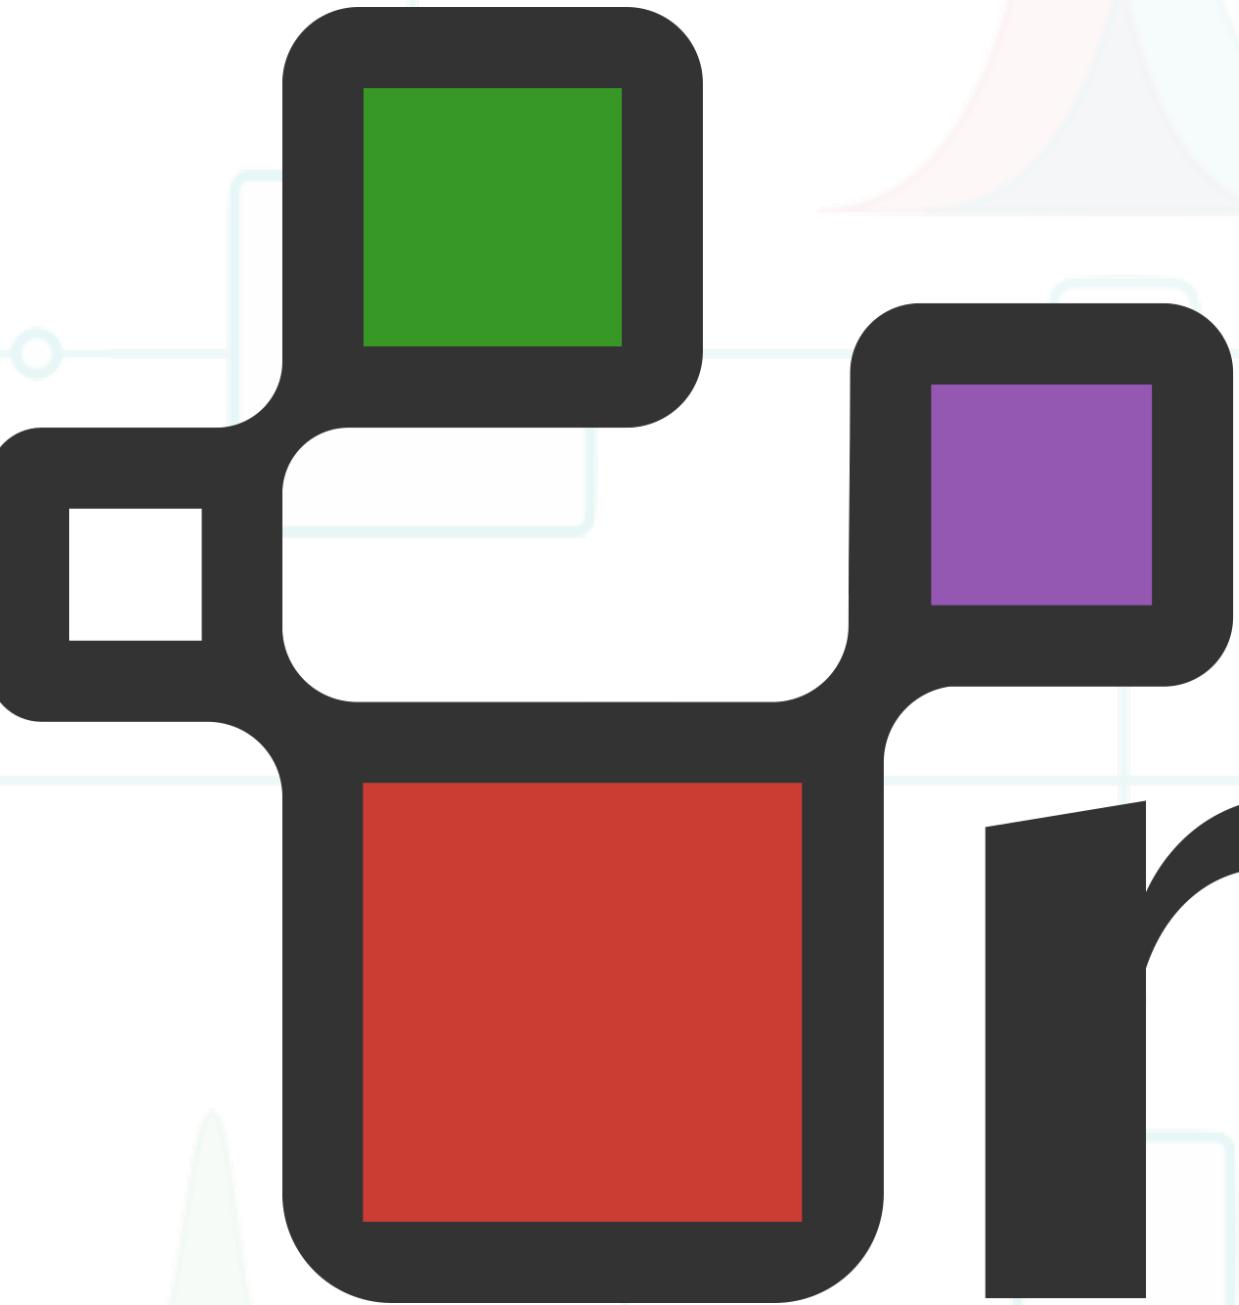
- Adaptive and robust
- Parallel and lazy



# How do we approach it?

Scalable Bayesian Inference architecture for autonomous systems





# rxinfer

Automatic Bayesian Inference through Reactive Message Passing

# Reactive Bayesian Inference

Automatic Bayesian Inference through Reactive Message Passing

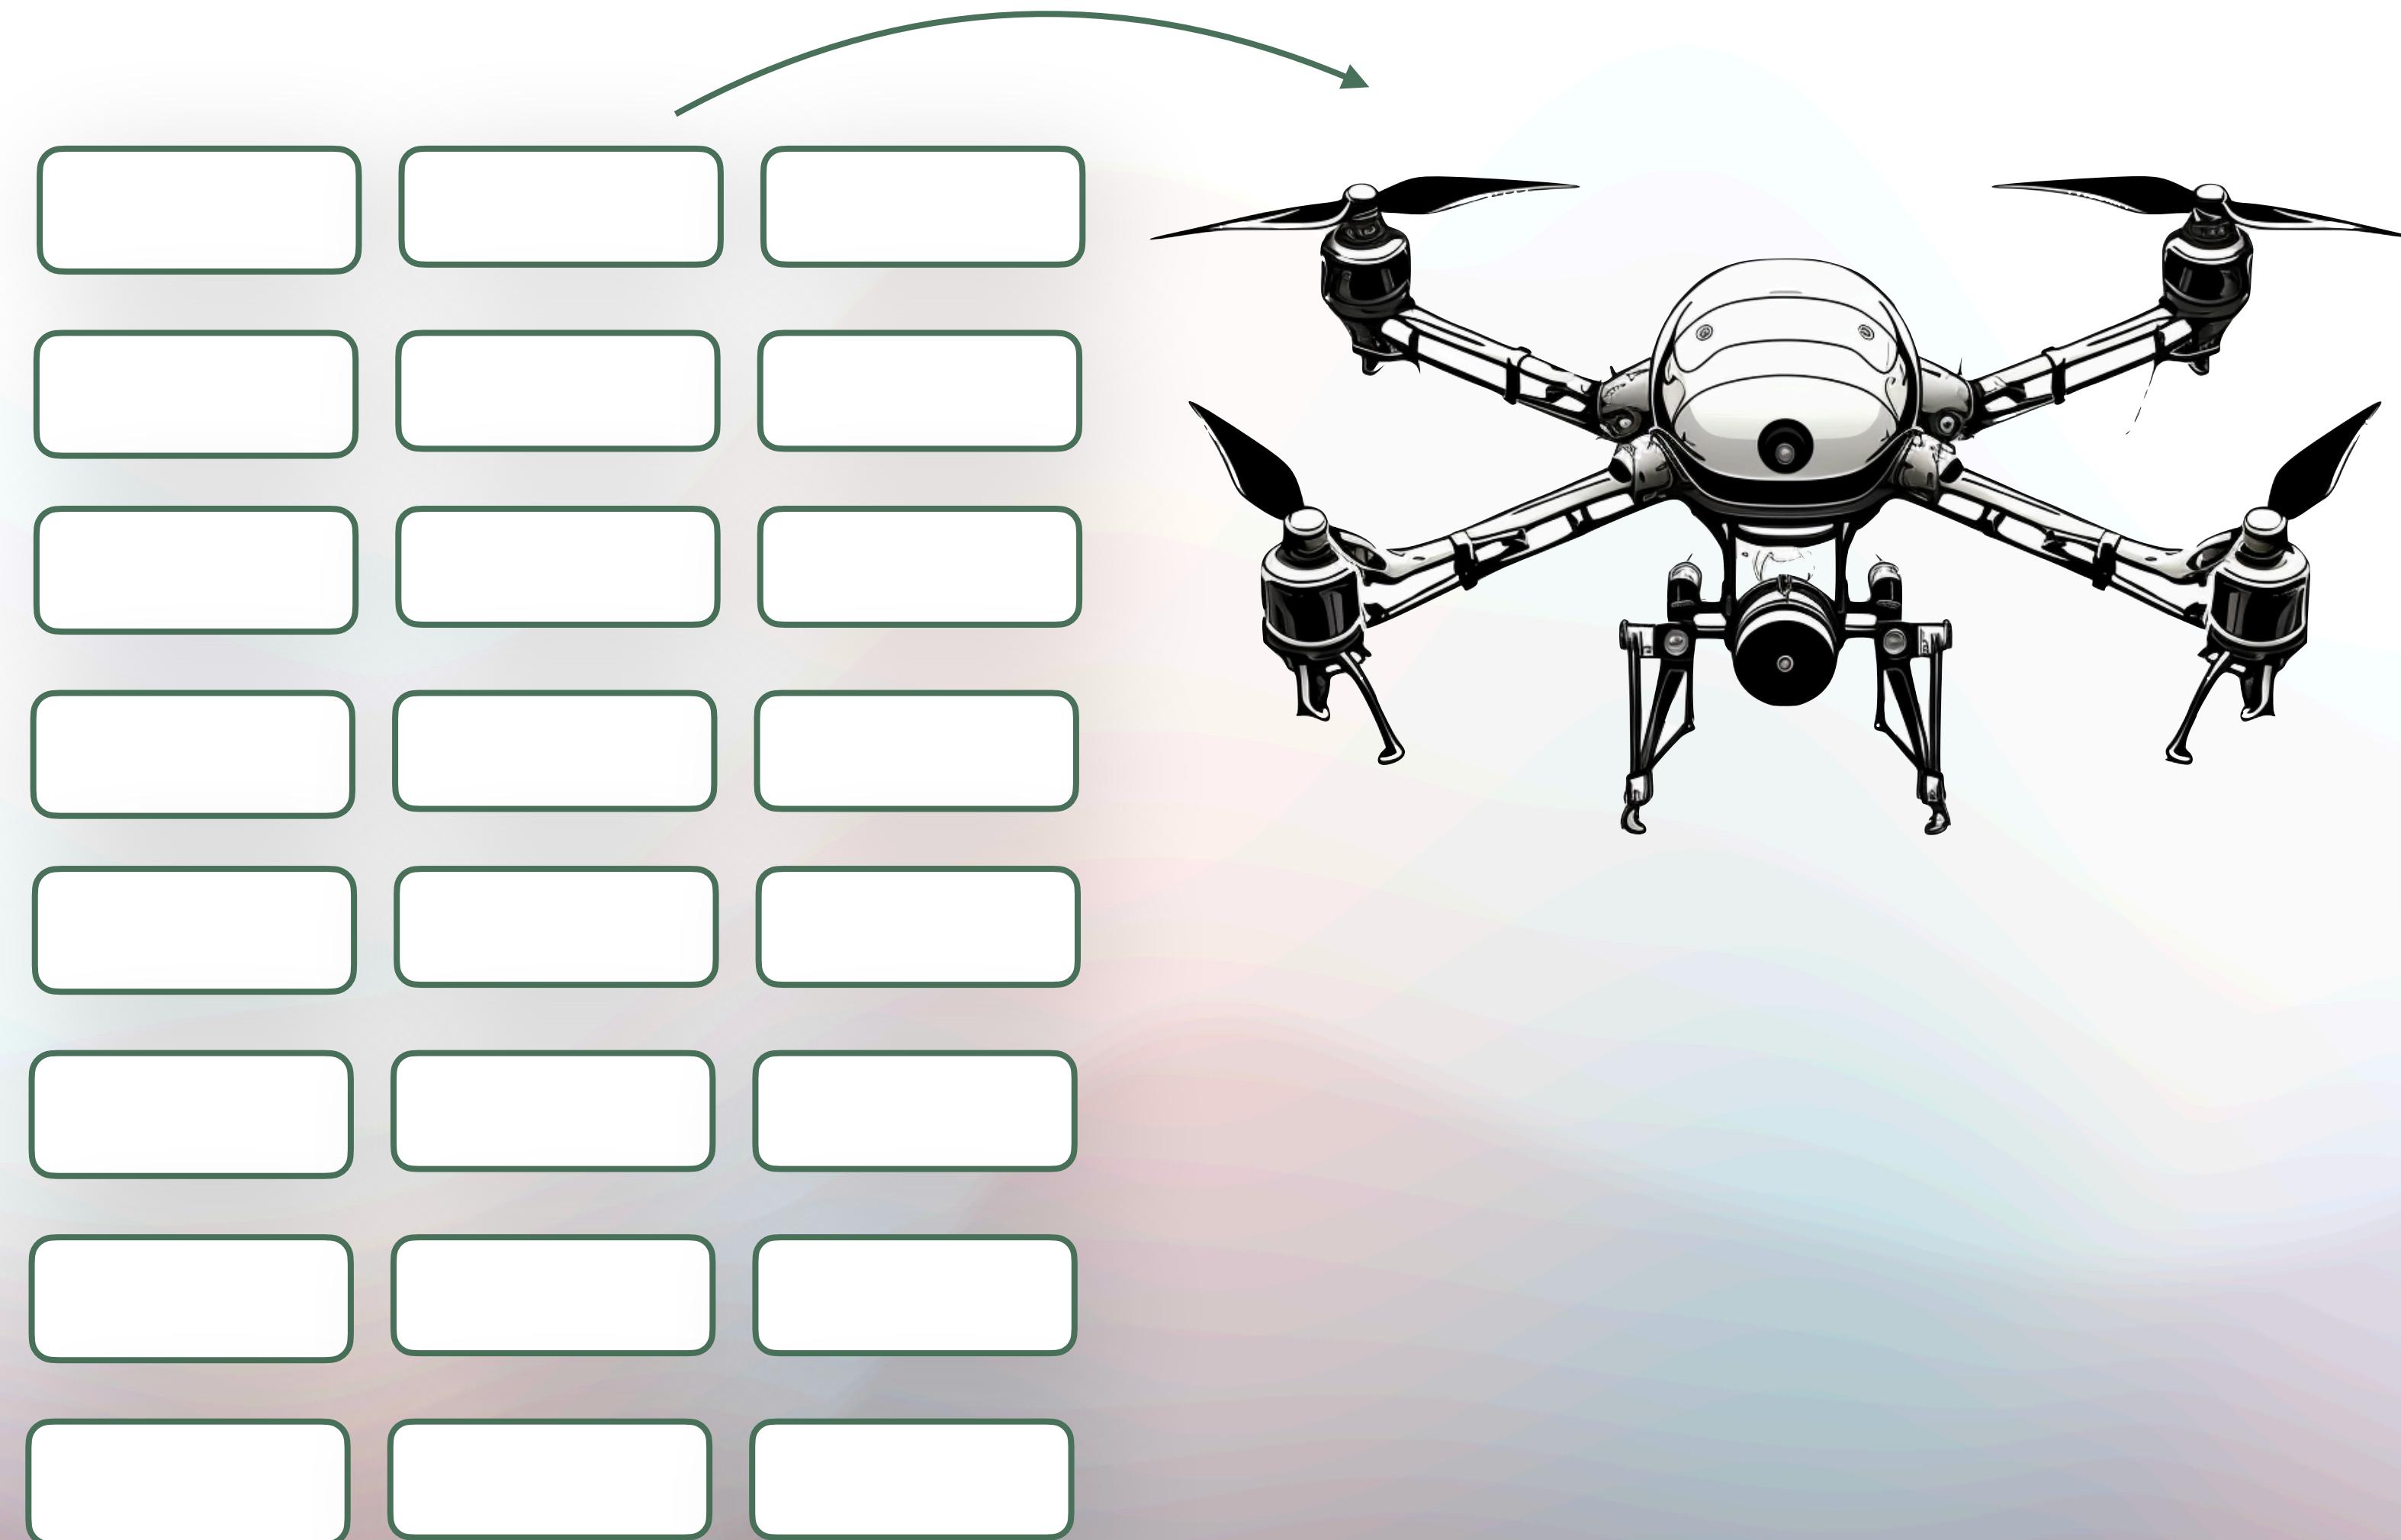
The screenshot shows the RxInfer website ([rxinfer.ml](https://rxinfer.ml)). The header includes a GitHub icon and a "GitHub" link. The main content features the RxInfer logo, a navigation bar with links to "Get Started", "Documentation", "Examples", "Discussions", "Team", and "Contact". Below the navigation is a large "rxinfer" logo with the tagline "Automatic Bayesian Inference through Reactive Message Passing". A diagram at the bottom illustrates a factor graph with nodes labeled with the Greek letter  $\mathcal{N}$  and equality signs (=) representing messages.

The screenshot shows the RxInfer GitHub repository page ([github.com/biaslab/RxInfer.jl](https://github.com/biaslab/RxInfer.jl)). The page displays a list of recent code commits by `bvdmitri`, including updates to version 2.15.0. It also shows sections for "Releases" (with v2.15.0 as the latest), "Packages" (none published), "Contributors" (16), "Deployments" (212 deployments), and "Languages" (Jupyter Notebook 60%, Julia 36.0%, Other 0.4%). The "Overview" section describes RxInfer as a Julia package for automatic Bayesian inference on a factor graph with reactive message passing, mentioning its use of a Forney-style factor graph representation.

# Reactive Bayesian Inference

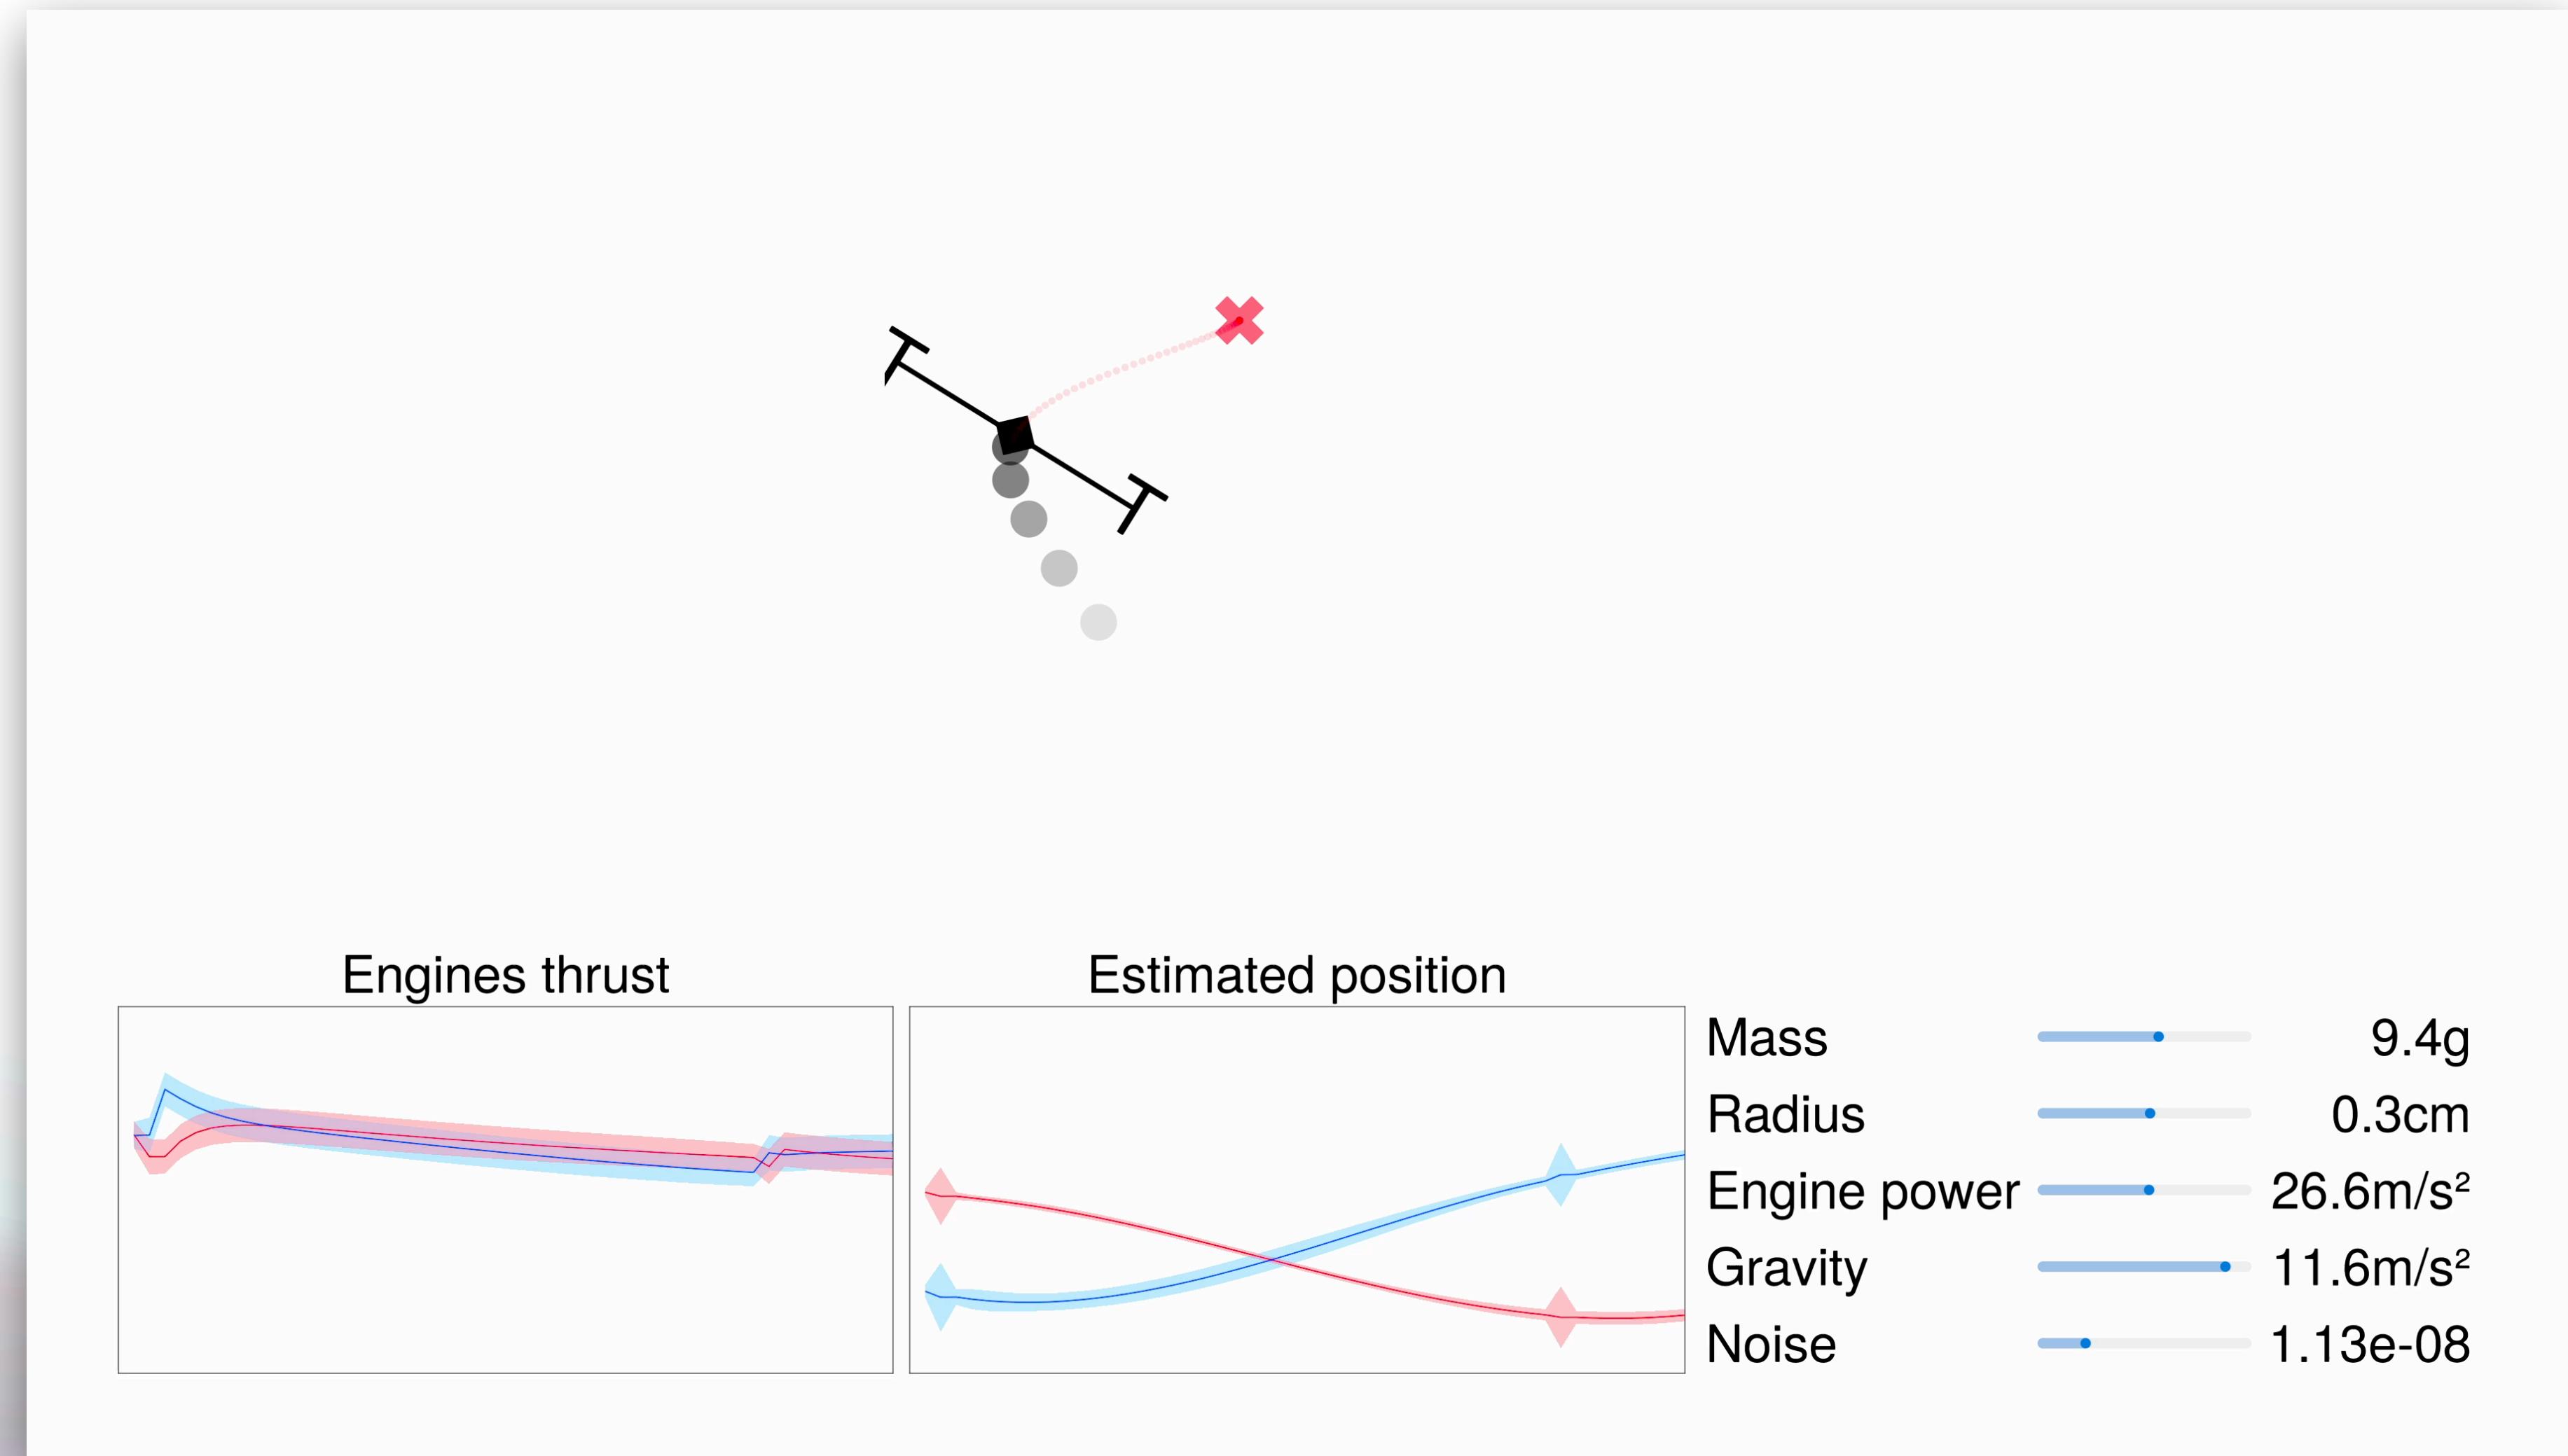
Automatic Bayesian Inference through Reactive Message Passing

Add as many components as you want



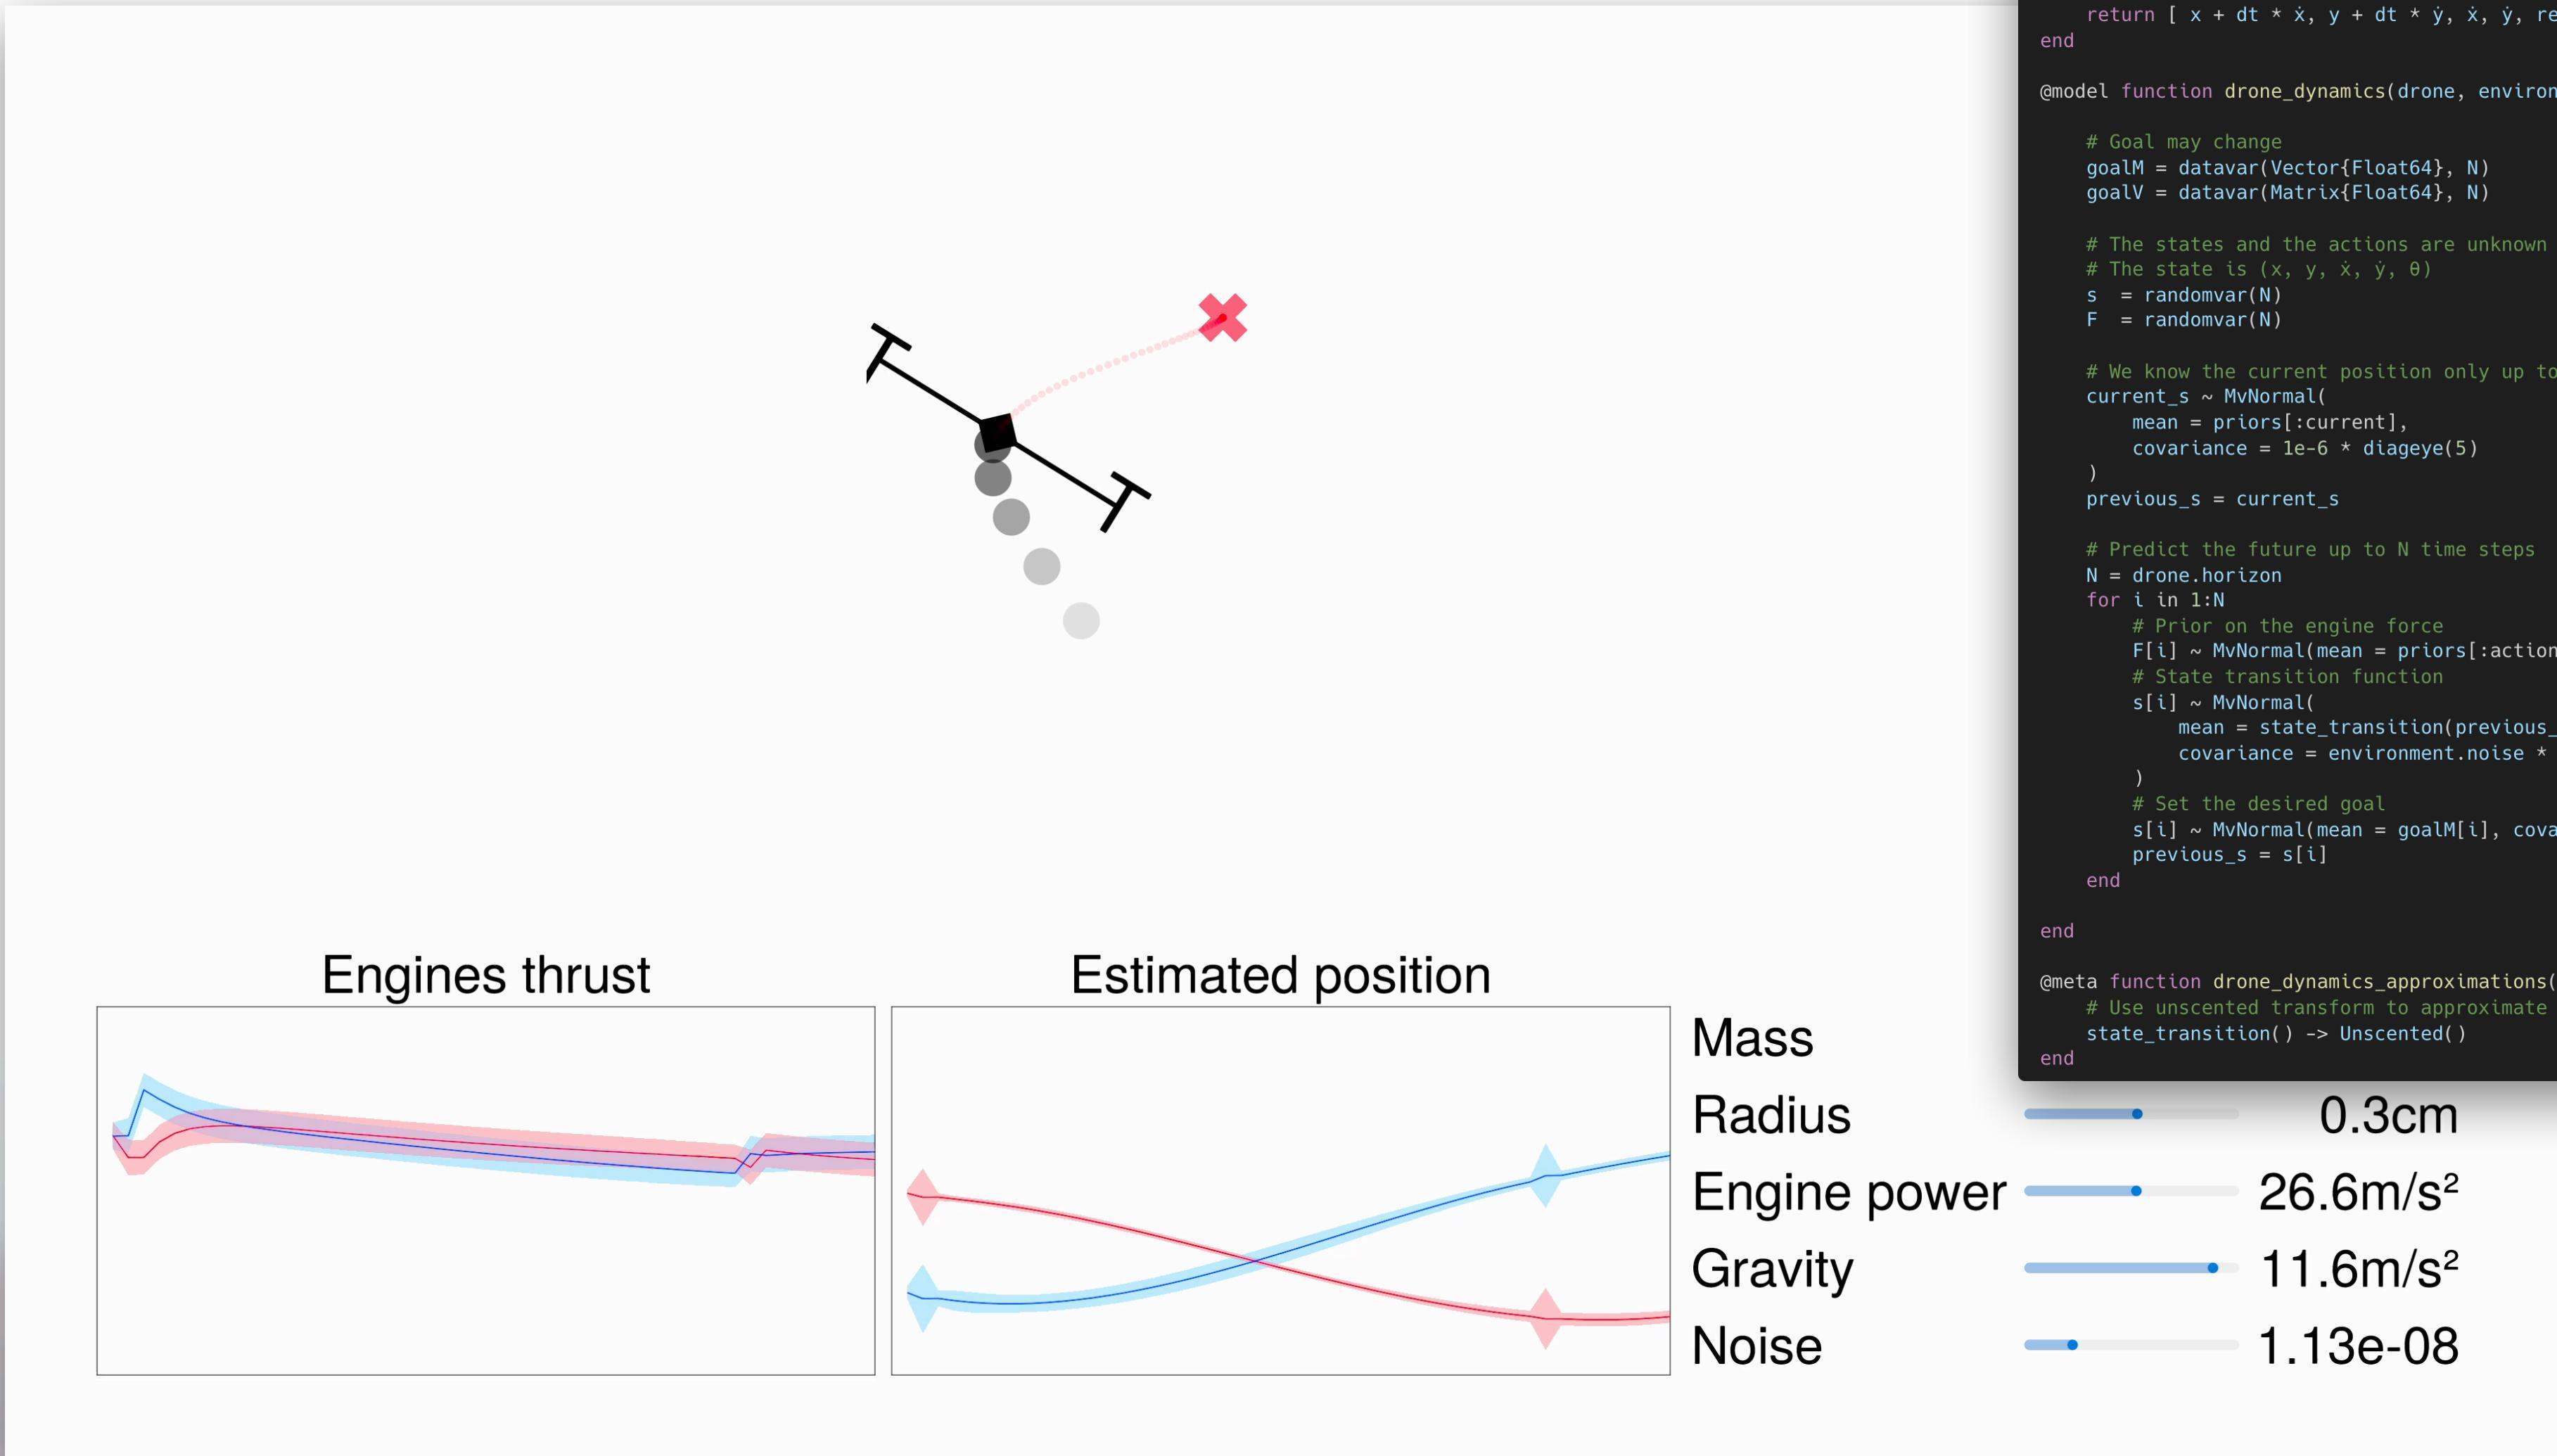
# Reactive Bayesian Inference

Automatic Bayesian Inference through Reactive Message Passing



# Reactive Bayesian Inference

Automatic Bayesian Inference through Reactive Message Passing



```
function state_transition(state, action, drone, environment)
    dt, g, _ = environment
    _, r, m, l = drone
    F_l, F_r = action
    x, y, x̄, ȳ, θ = state

    dx = (F_l + F_r)/m * sin(θ)
    dy = (F_l + F_r)/m * cos(θ) - g/m
    dθ = (F_l - F_r) * r
    x̄ = x + dt * dx
    ȳ = y + dt * dy
    θ̄ = rem(θ + dt * dθ, 2π)

    return [x + dt * x̄, y + dt * ȳ, x̄, ȳ, rem(θ + dt * dθ, 2π)]
end

@model function drone_dynamics(drone, environment, priors)
    # Goal may change
    goalM = datavar(Vector{Float64}, N)
    goalV = datavar(Matrix{Float64}, N)

    # The states and the actions are unknown components
    # The state is (x, y, x̄, ȳ, θ)
    s = randomvar(N)
    F = randomvar(N)

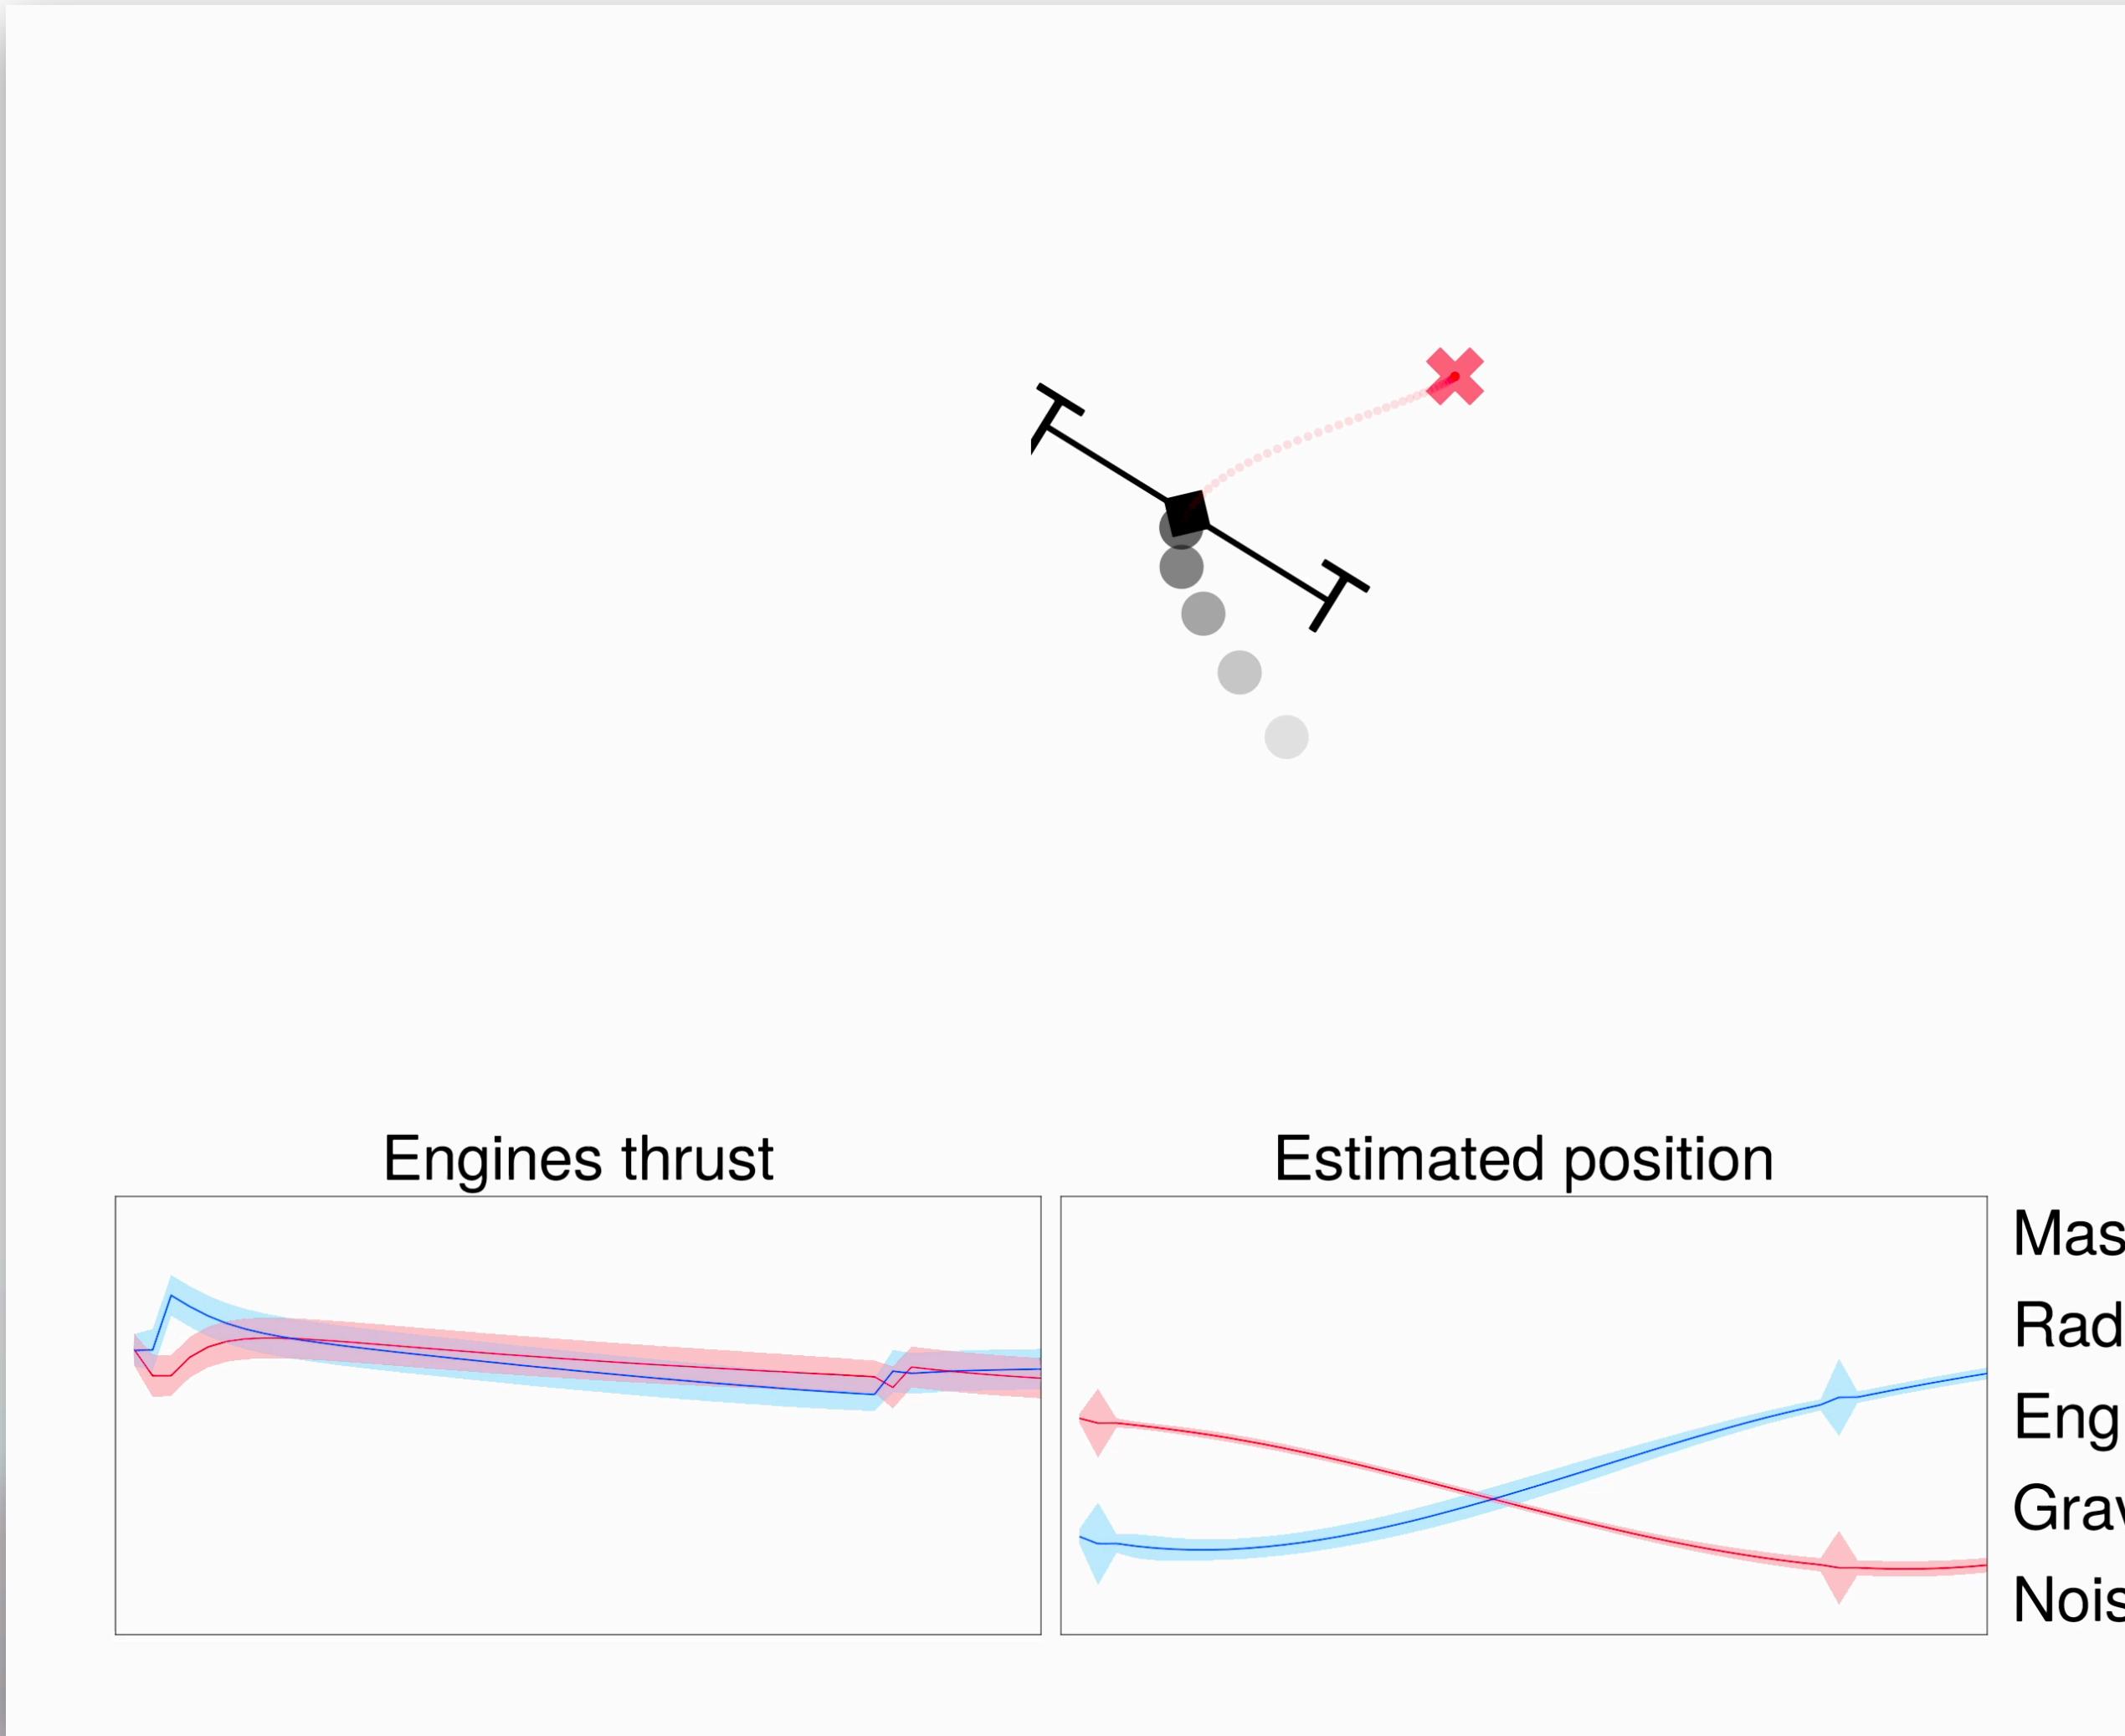
    # We know the current position only up to a certain degree
    current_s ~ MvNormal(
        mean = priors[:current],
        covariance = 1e-6 * diageye(5)
    )
    previous_s = current_s

    # Predict the future up to N time steps
    N = drone.horizon
    for i in 1:N
        # Prior on the engine force
        F[i] ~ MvNormal(mean = priors[:actions][i], covariance = diageye(2))
        # State transition function
        s[i] ~ MvNormal(
            mean = state_transition(previous_s, F[i], drone, environment),
            covariance = environment.noise * diageye(5)
        )
        # Set the desired goal
        s[i] ~ MvNormal(mean = goalM[i], covariance = goalV[i])
        previous_s = s[i]
    end
end

@meta function drone_dynamics_approximations()
    # Use unscented transform to approximate the non-linearity
    state_transition() -> Unscented()
end
```

# Reactive Bayesian Inference

Automatic Bayesian Inference through Reactive Message Passing



```

region
  function state_transition(state, action, drone, environment)
    dt, g, _ = environment
    _, r, m, l = drone
    F_l, F_r = action
    x, y, ḡ, ḙ, θ = state

    dx = (F_l + F_r)/m * sin(θ)
    dy = (F_l + F_r)/m * cos(θ) - g/m
    dθ = (F_l - F_r) * r
    ḡ = ḡ + dt * dx
    ḙ = ḙ + dt * dy

    return [x + dt * ḡ, y + dt * ḙ, ḡ, ḙ, rem(θ + dt)]
end

@model function drone_dynamics(drone, environment, prior)
  # Goal may change
  goalM = datavar(Vector{Float64}, N)
  goalV = datavar(Matrix{Float64}, N)

  # The states and the actions are unknown components
  # The state is (x, y, ḡ, ḙ, θ)
  s = randomvar(N)
  F = randomvar(N)

  # We know the current position only up to a certain uncertainty
  current_s ~ MvNormal()
  previous_s = current_s

  # Predict the future up to N time steps
  N = drone.horizon
  for i in 1:N
    # Prior on the engine force
    F[i] ~ MvNormal(mean = priors[:actions][i], covariance = 1e-6 * diageye(5))
    # State transition function
    s[i] ~ MvNormal(
      mean = state_transition(previous_s, F[i]),
      covariance = environment.noise * diageye(5))
  end

  # Set the desired goal
  s[i] ~ MvNormal(mean = goalM[i], covariance = 1e-6 * diageye(5))
  previous_s = s[i]
end

@meta function drone_dynamics_approximations()
  # Use unscented transform to approximate the non-linear state_transition() -> Unscented()
end

```



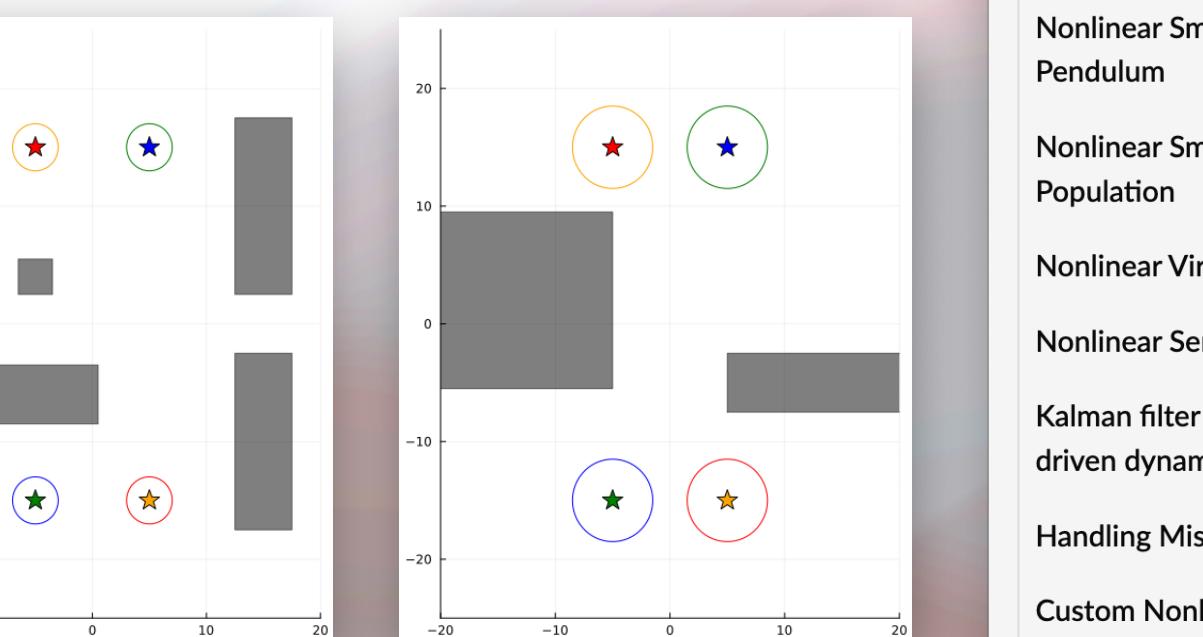
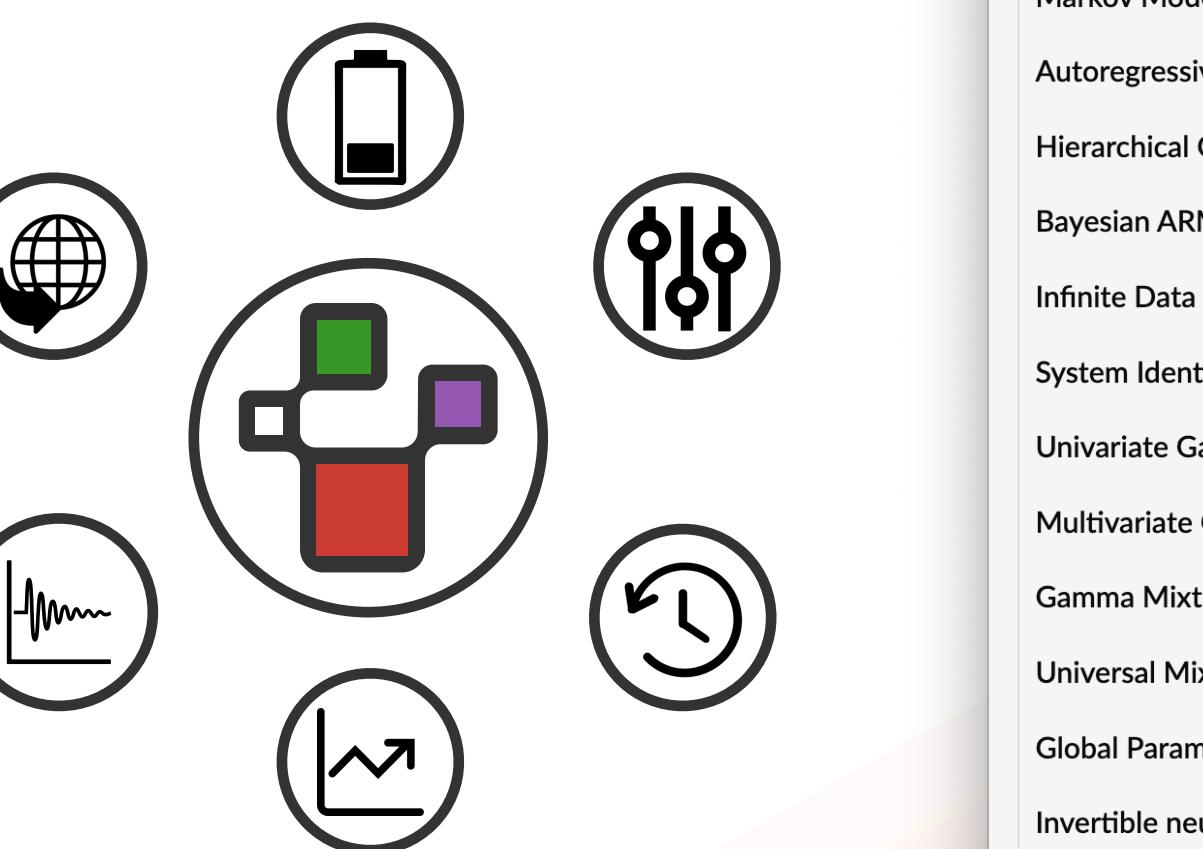
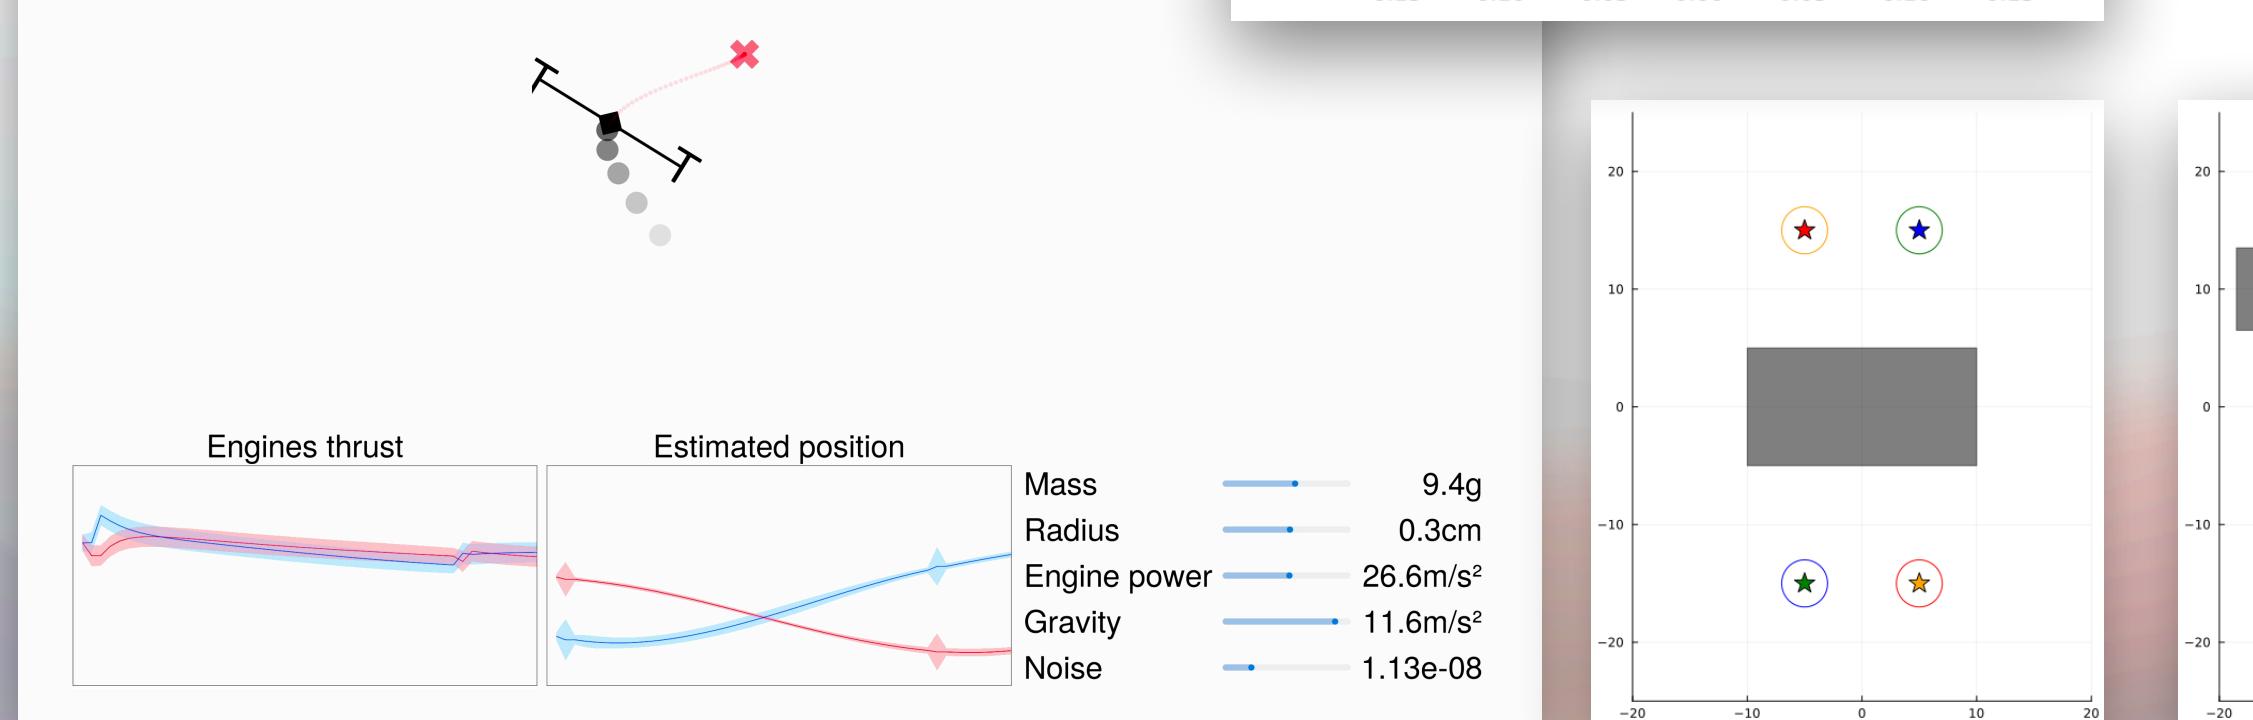
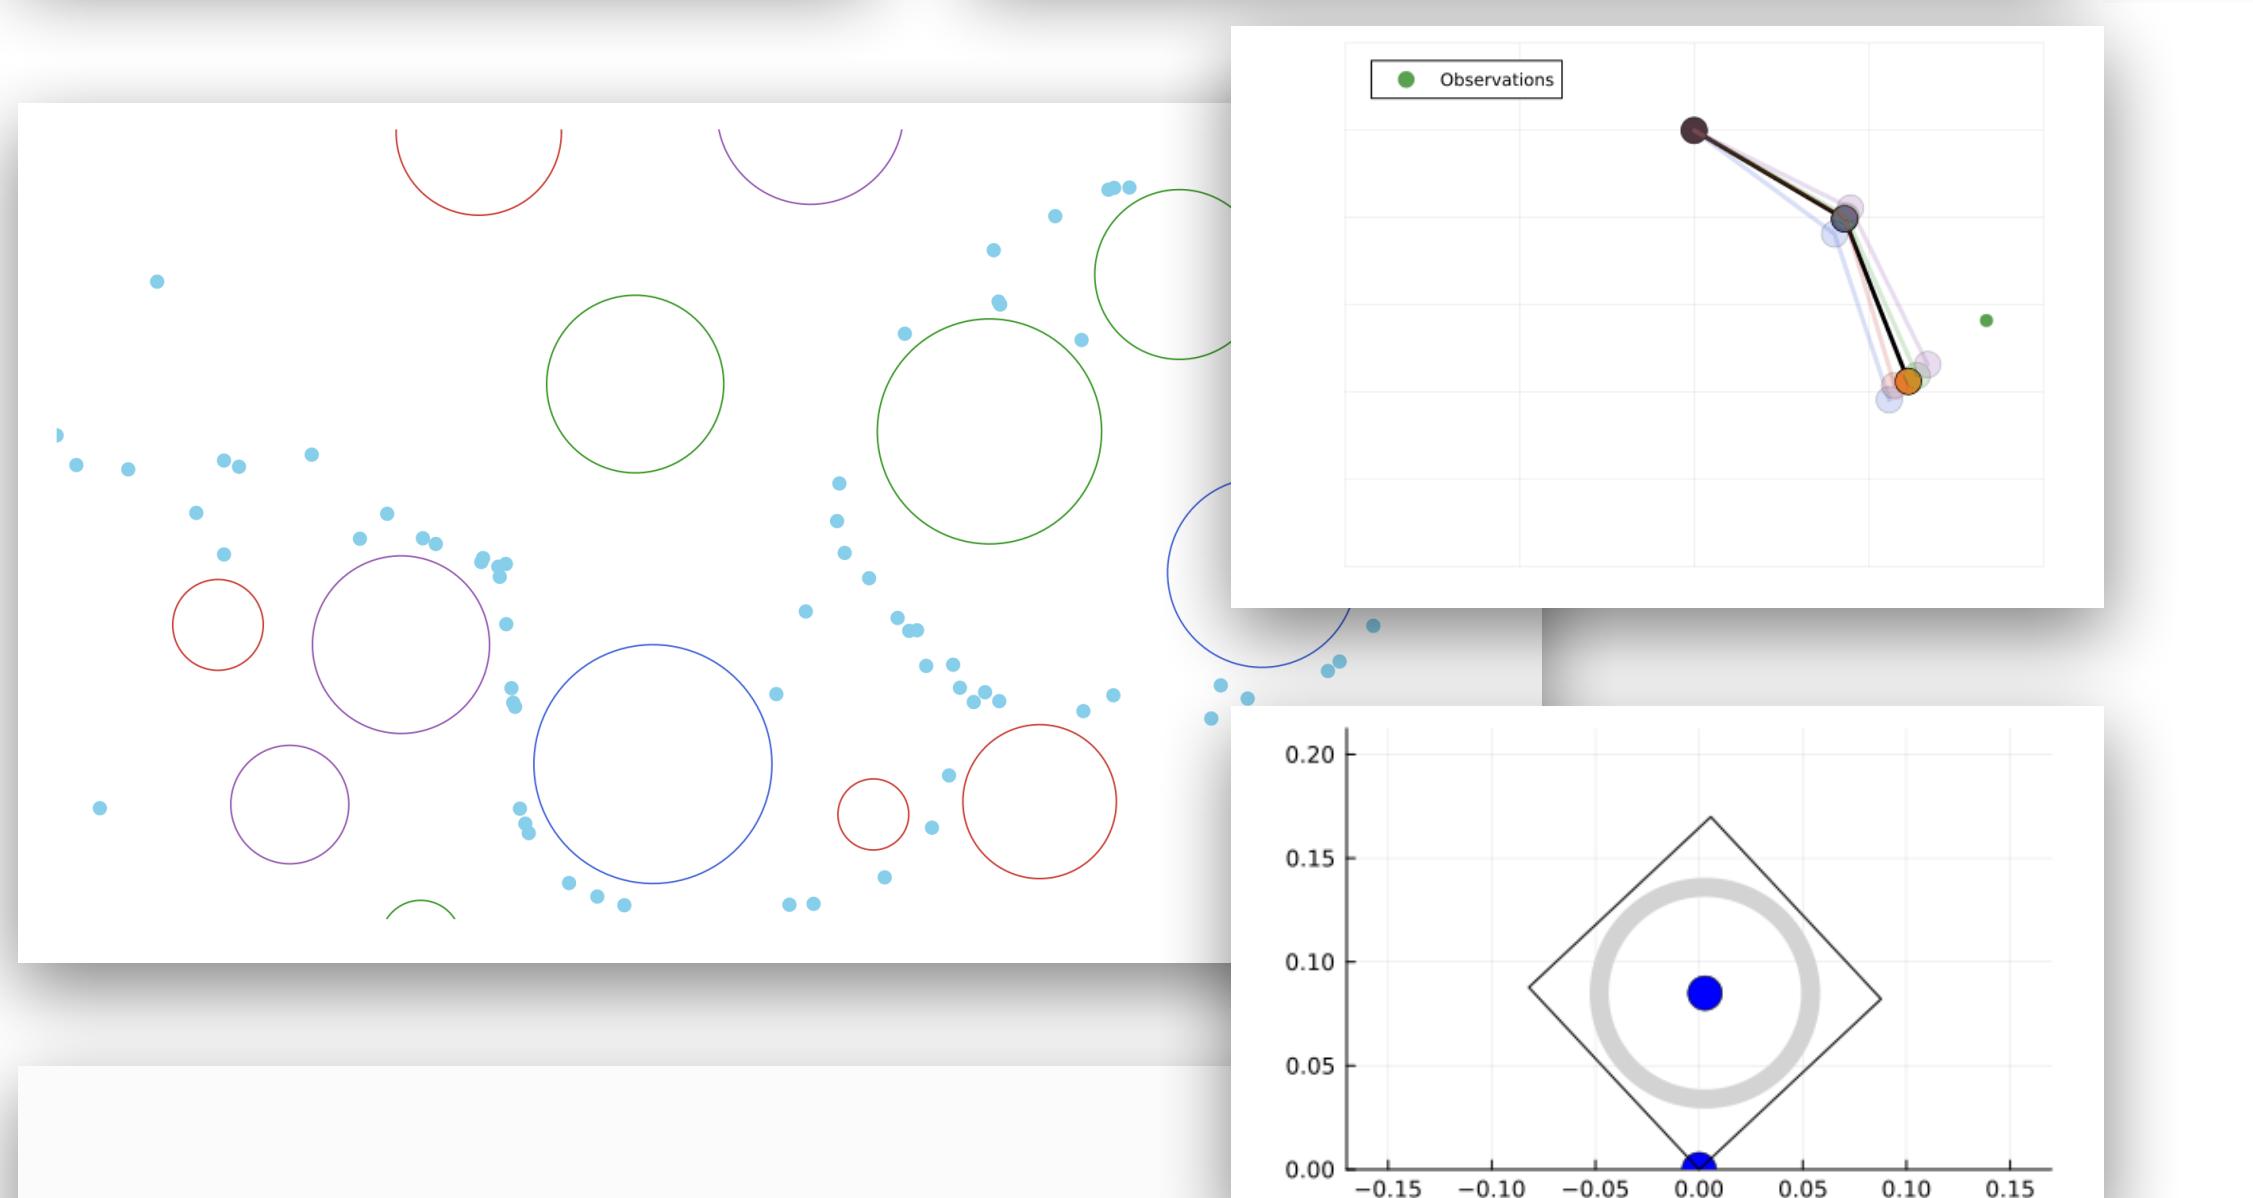
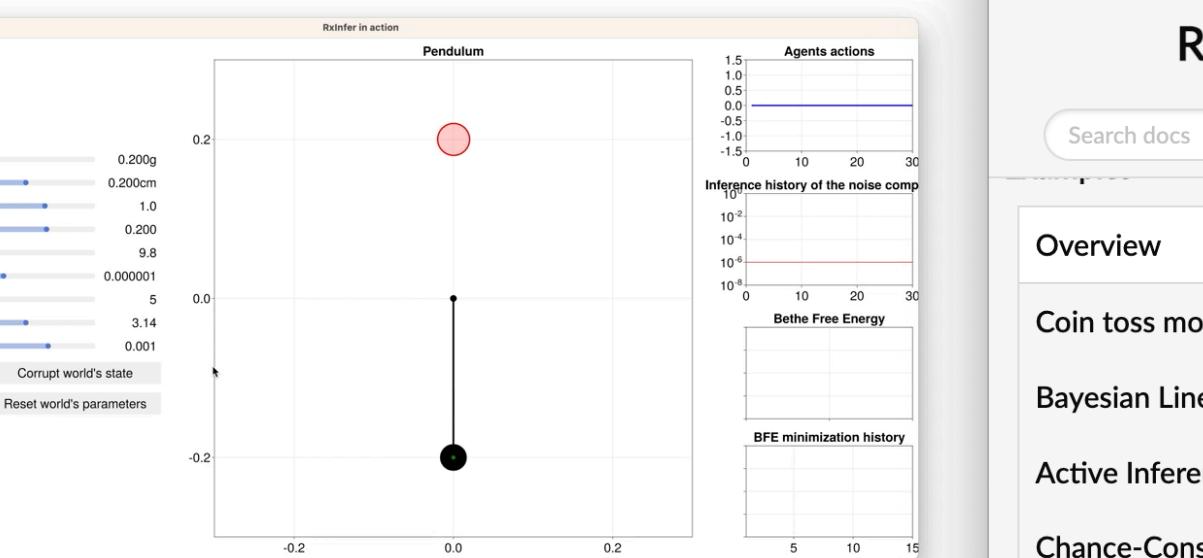
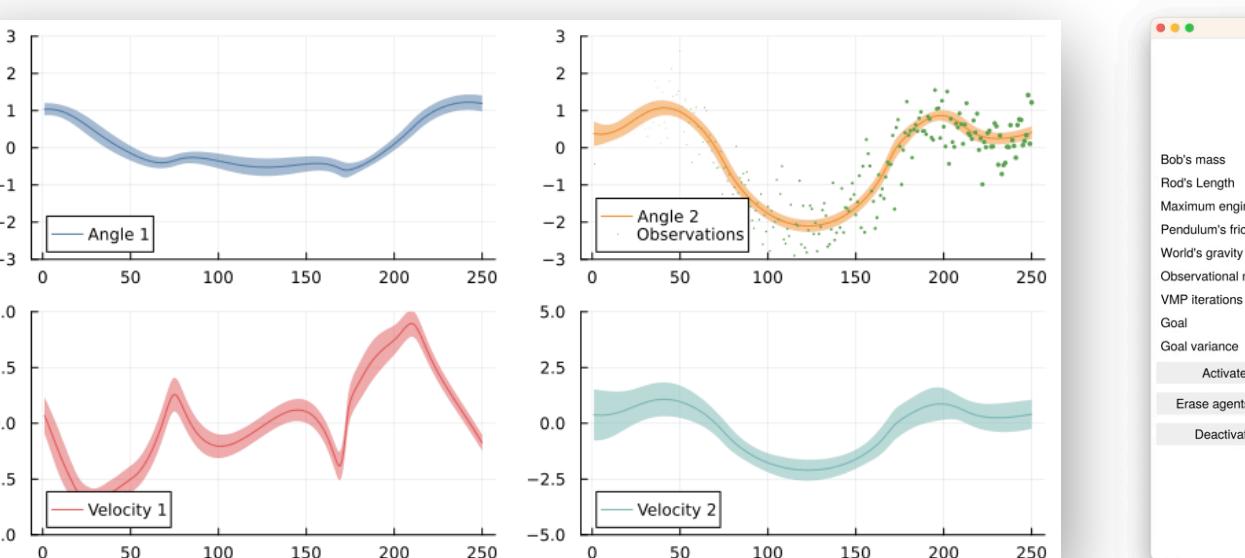
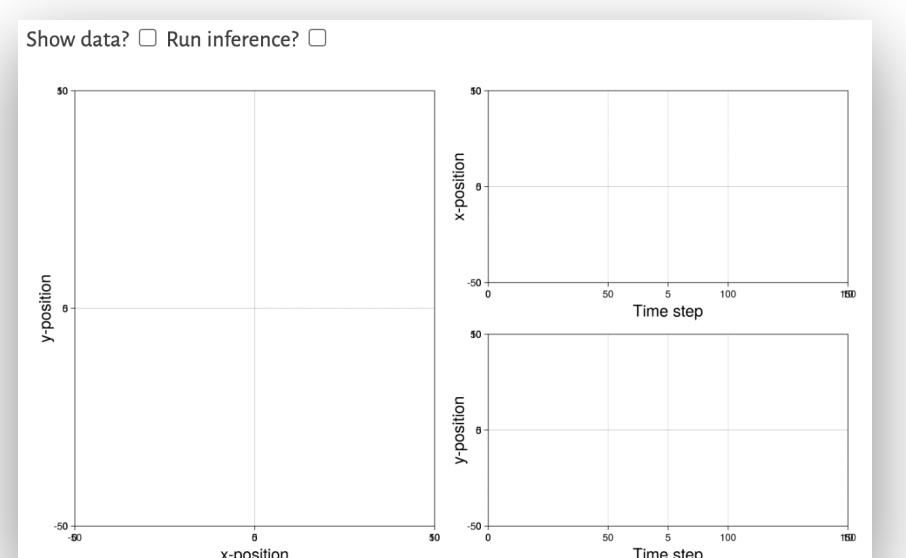
RxInfer.jl

Search docs

## Overview

[Coin toss model \(Beta-Bernoulli\)](#)[Bayesian Linear Regression](#)[Active Inference Mountain car](#)[Chance-Constrained Active Inference](#)[Assessing People's Skills](#)[Gaussian Linear Dynamical System](#)[Ensemble Learning of a Hidden Markov Model](#)[Autoregressive Model](#)[Hierarchical Gaussian Filter](#)[Bayesian ARMA model](#)[Infinite Data Stream](#)[System Identification Problem](#)[Univariate Gaussian Mixture Model](#)[Multivariate Gaussian Mixture Model](#)[Gamma Mixture Model](#)[Universal Mixtures](#)[Global Parameter Optimisation](#)[Invertible neural networks: a tutorial](#)[Conjugate-Computational Variational Message Passing \(CVI\)](#)[Solve GP regression by SDE](#)[Nonlinear Smoothing: Noisy Pendulum](#)[Nonlinear Smoothing: Rabbit Population](#)[Nonlinear Virus Spread](#)[Nonlinear Sensor Fusion](#)[Kalman filter with LSTM network driven dynamic](#)[Handling Missing Data](#)[Custom Nonlinear Node](#)[Probit Model \(EP\)](#)All examples have been pre-generated automatically from the [example](#)

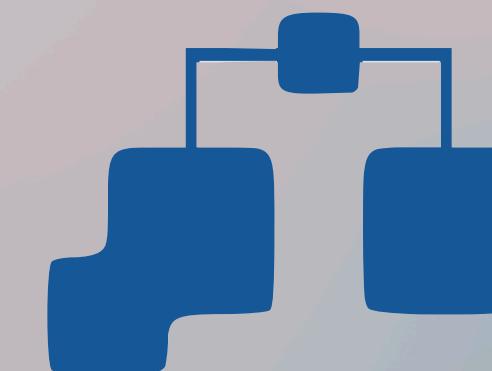
- [Coin toss model \(Beta-Bernoulli\)](#): An example of Bayesian inference observations.
- [Bayesian Linear Regression](#): An example of Bayesian linear regression
- [Active Inference Mountain car](#): This notebook covers RxInfer usage simple mountain car problem.
- [Chance-Constrained Active Inference](#): This notebook applies reactive the context of chance-constraints.
- [Assessing People's Skills](#): The demo is inspired by the example from Machine Learning book. We are going to perform an exact inference results of the test.
- [Gaussian Linear Dynamical System](#): An example of inference procedure with multivariate noisy observations using Belief Propagation (Sum Bayesian Filtering and Smoothing).
- [Ensemble Learning of a Hidden Markov Model](#): An example of structural Hidden Markov Model with unknown transition and observational r
- [Autoregressive Model](#): An example of variational Bayesian Inference Reference: Albert Podusenko, [Message Passing-Based Inference for Autoregressive Models](#)
- [Hierarchical Gaussian Filter](#): An example of online inference procedure univariate noisy observations using Variational Message Passing algorithm [Message Passing-based Inference in the Hierarchical Gaussian Filter](#)
- [Bayesian ARMA model](#): This notebook shows how Bayesian ARMA can be implemented in RxInfer.jl
- [Infinite Data Stream](#): This example shows RxInfer capabilities of running
- [System Identification Problem](#): This example attempts to identify an
- [Univariate Gaussian Mixture Model](#): This example implements variational Gaussian mixture model with mean-field assumption.
- [Multivariate Gaussian Mixture Model](#): This example implements variational Gaussian mixture model with mean-field assumption.
- [Gamma Mixture Model](#): This example implements one of the Gamma Mixture Models <https://biaslab.github.io/publication/mp-based-inference-in-gmm/>.
- [Universal Mixtures](#): Universal mixture modeling.
- [Global Parameter Optimisation](#): This example shows how to use RxInfer optimisation packages such as Optim.jl.
- [Invertible neural networks: a tutorial](#): An example of variational Bay networks. Reference: Bart van Erp, [Hybrid Inference with Invertible Neural Networks](#)
- [Conjugate-Computational Variational Message Passing \(CVI\)](#): This example shows how to use CVI for the non-conjugate message-passing based inference by exploiting the conjugate message-passing.
- [Solve GP regression by SDE](#): In this notebook, we solve a GP regression Differential Equation' (SDE). This method is well described in the discussion of the paper 'Variational Message Passing for spatio-temporal Gaussian process regression.' by Arno Solin and Jouni Hartikainen.
- [Nonlinear Smoothing: Noisy Pendulum](#): In this demo, we will look at state transitions: tracking a noisy single pendulum. We translate a discrete form to a probabilistic model.
- [Nonlinear Smoothing: Rabbit Population](#): In this demo, we will look at state transitions. We will start with a one-dimensional problem; the number seems overly simple, but it is a good way to demonstrate the basic principle of nonlinear smoothing.
- [Nonlinear Virus Spread](#): In this demo we consider a model for the spread of a virus in a population. We are interested in estimating the reproduction rate from infected individuals.
- [Nonlinear Sensor Fusion](#): Nonlinear object position identification using particle filters.
- [Kalman filter with LSTM network driven dynamic](#): In this demo, we implement a Kalman filter with an LSTM network driven dynamic in Nonlinear State-Space Model using the LSTM.
- [Handling Missing Data](#): This example shows how to extend the set of



Version v2.11.1

# Reactive Probabilistic Programming for Scalable Bayesian Inference

Dmitry Bagaev



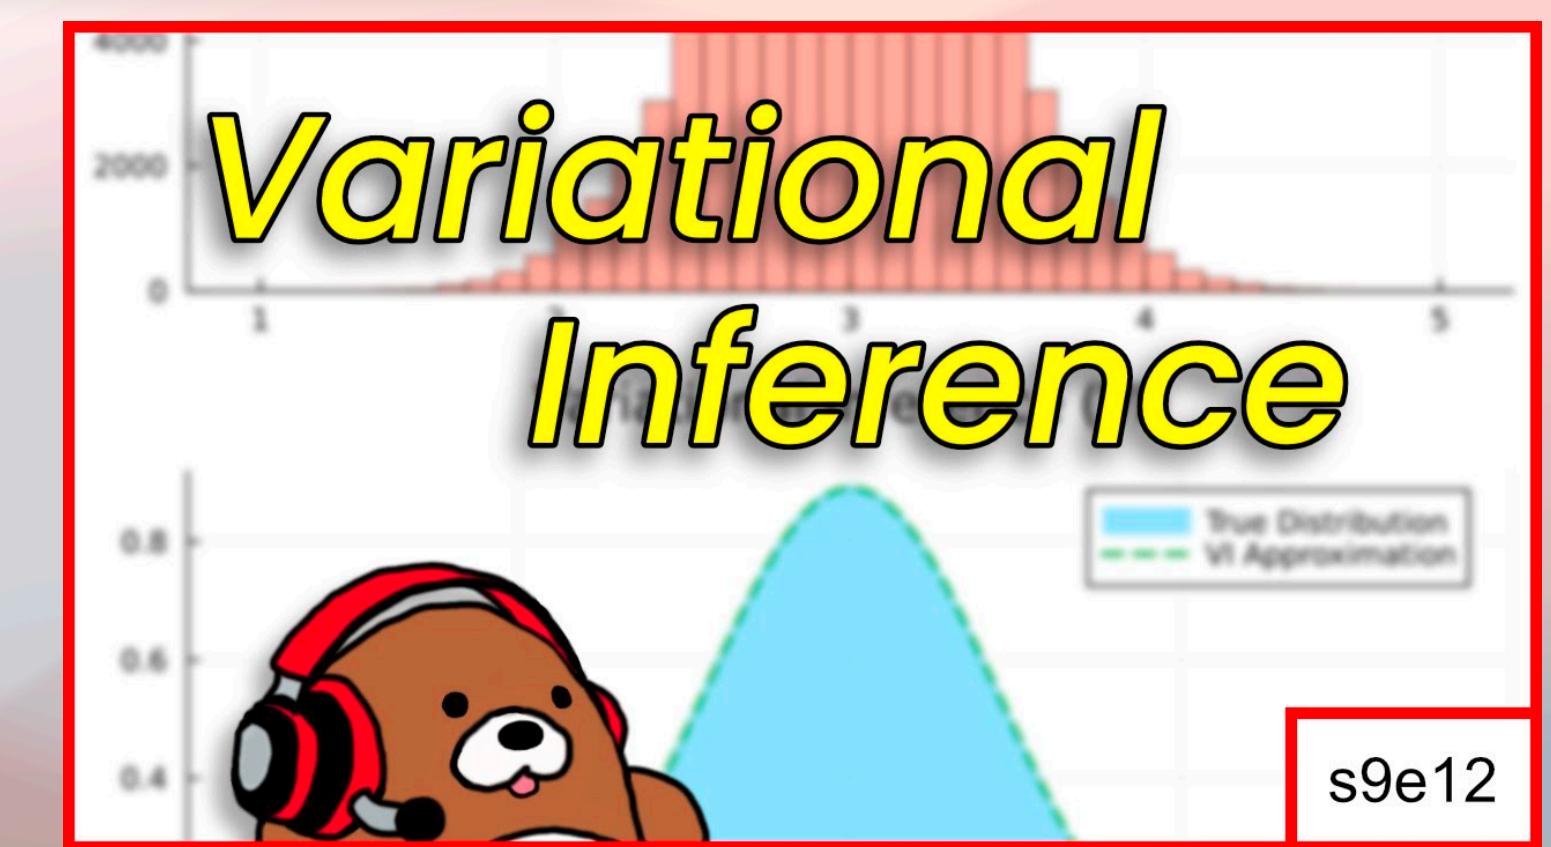
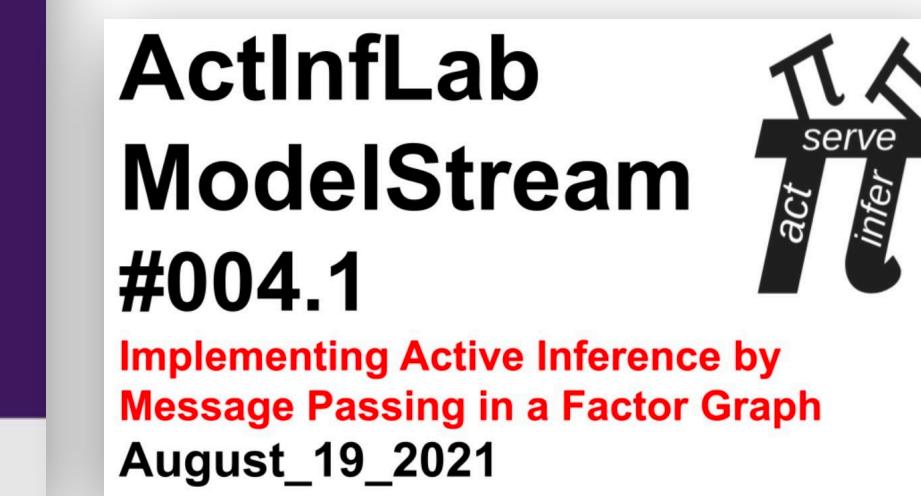
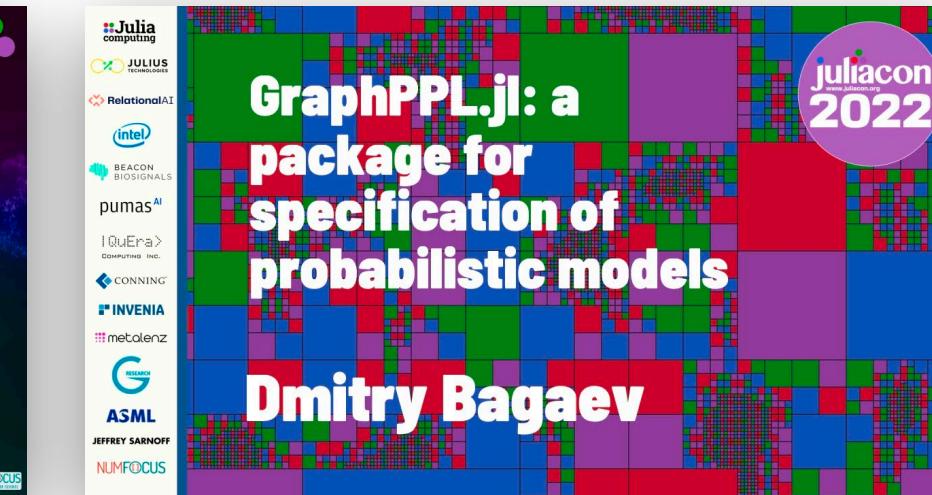
BIASlab



EINDHOVEN  
UNIVERSITY OF  
TECHNOLOGY

# Recognition beyond BIASlab

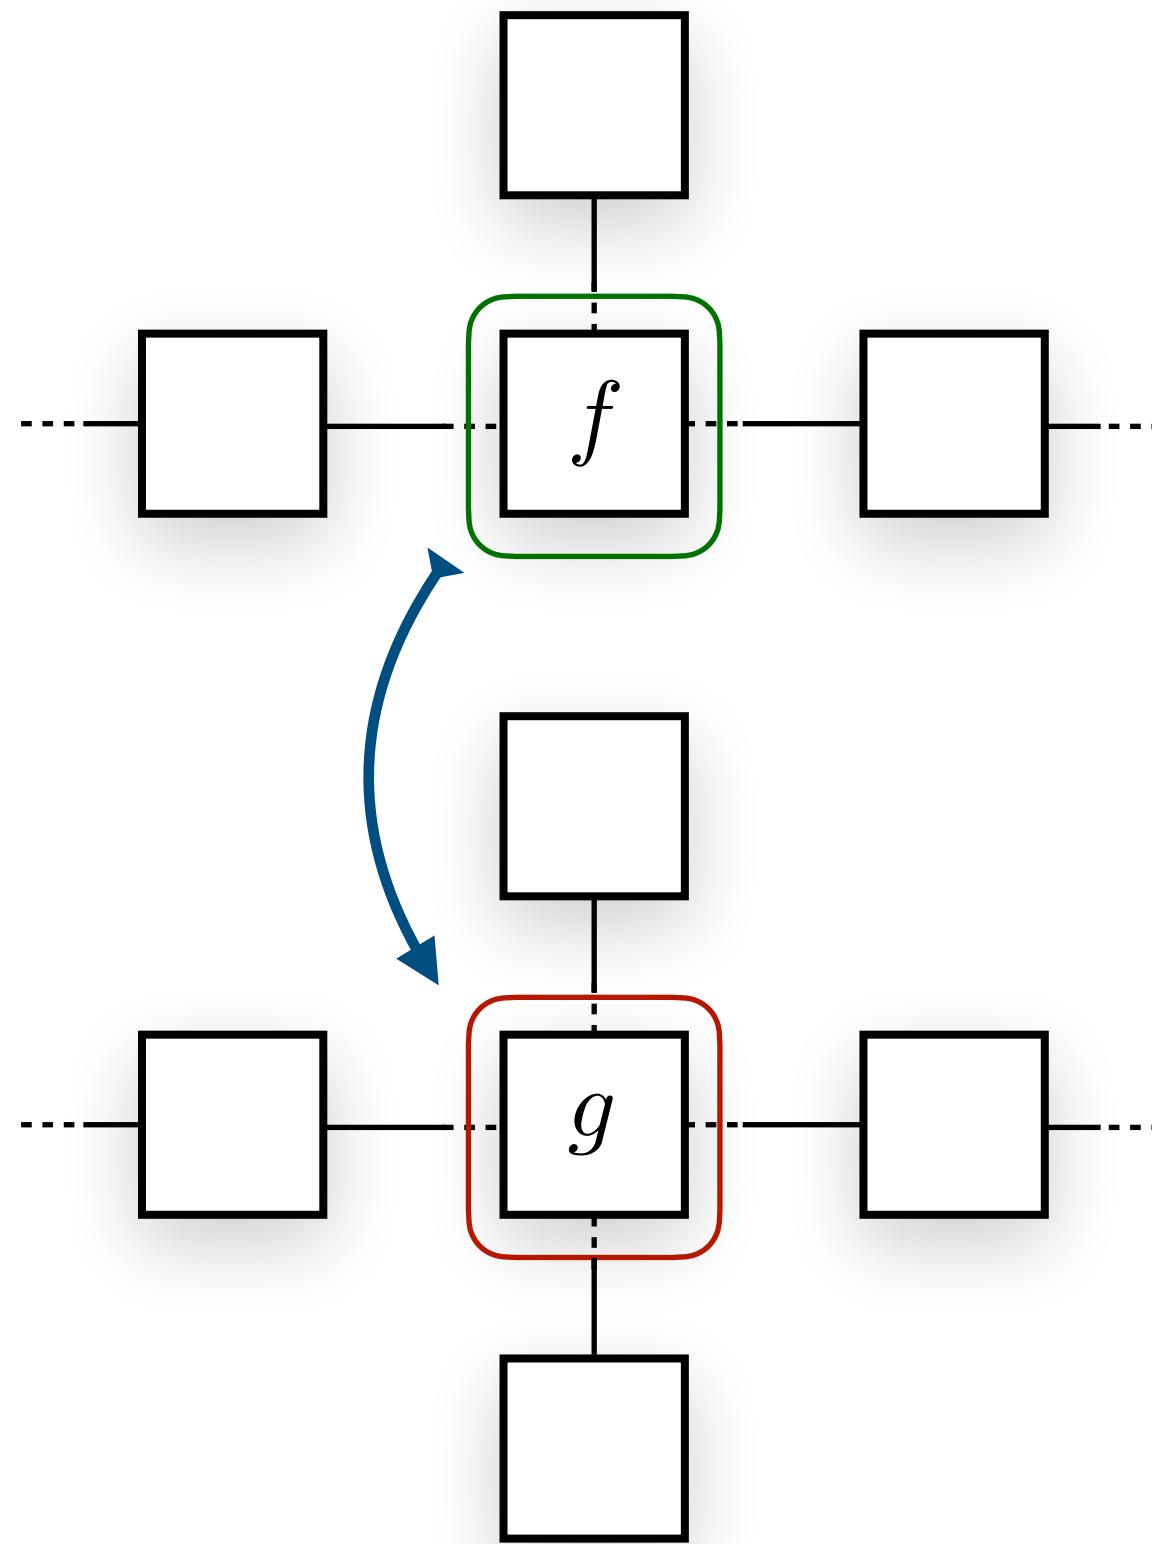
Scalable Bayesian Inference architecture for autonomous systems



# Future research ideas

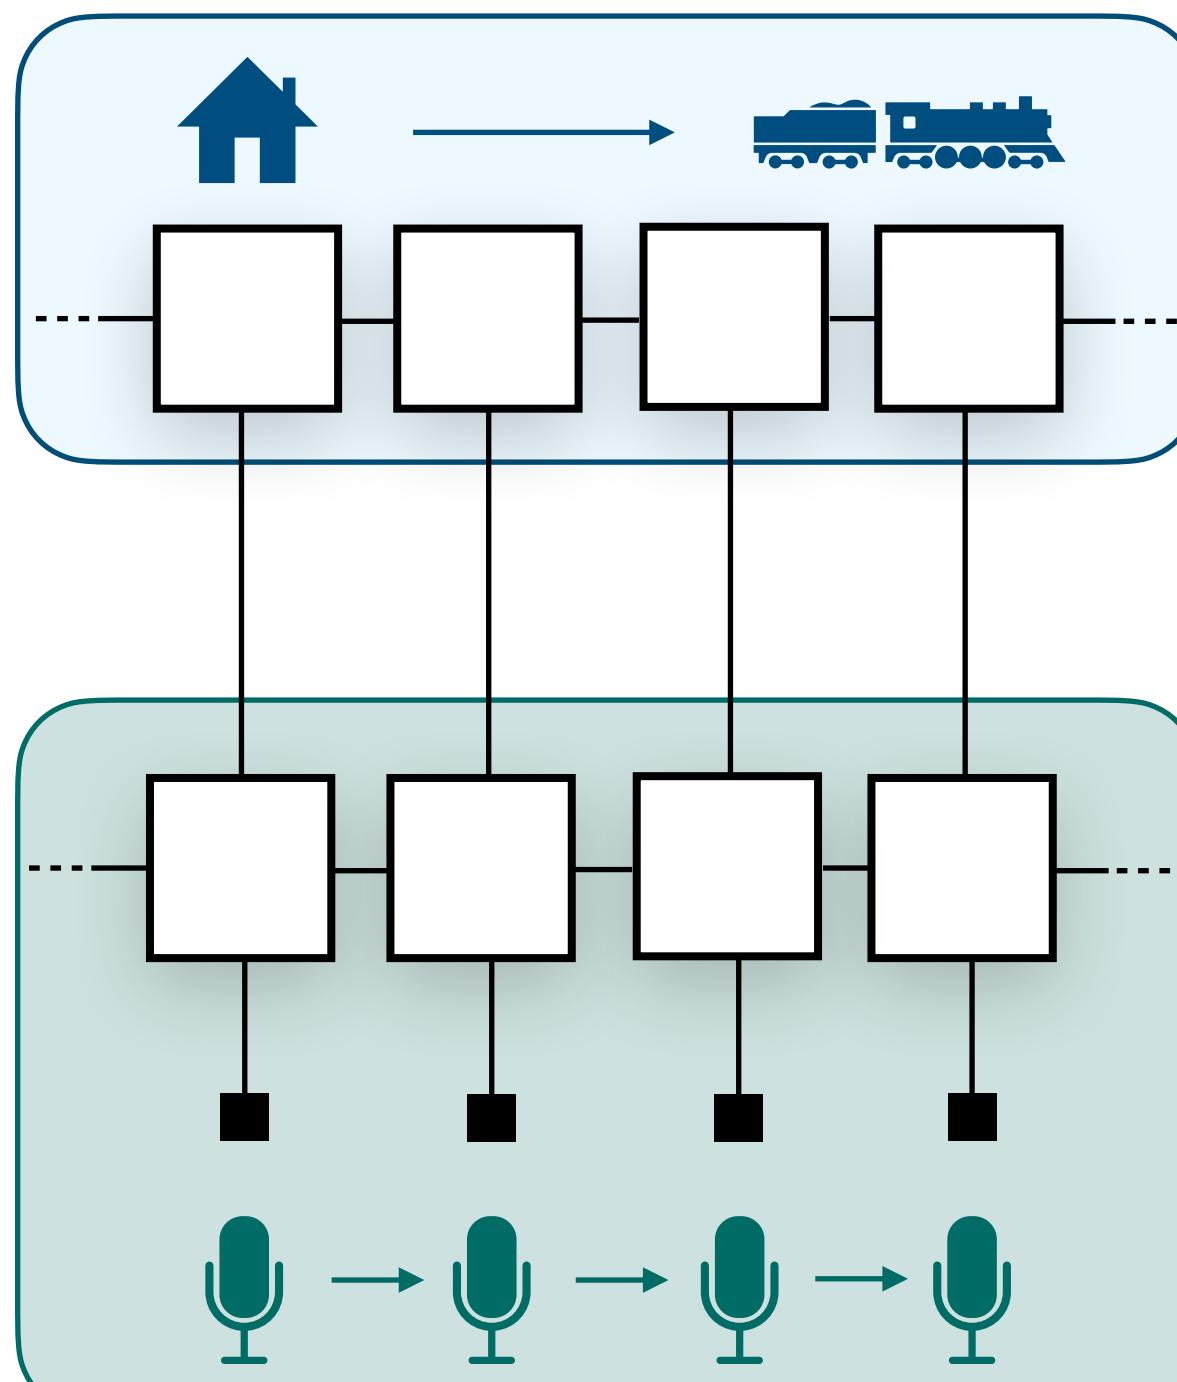
## Structure and constraints adaptation in real time

Infer structure & constraints automatically



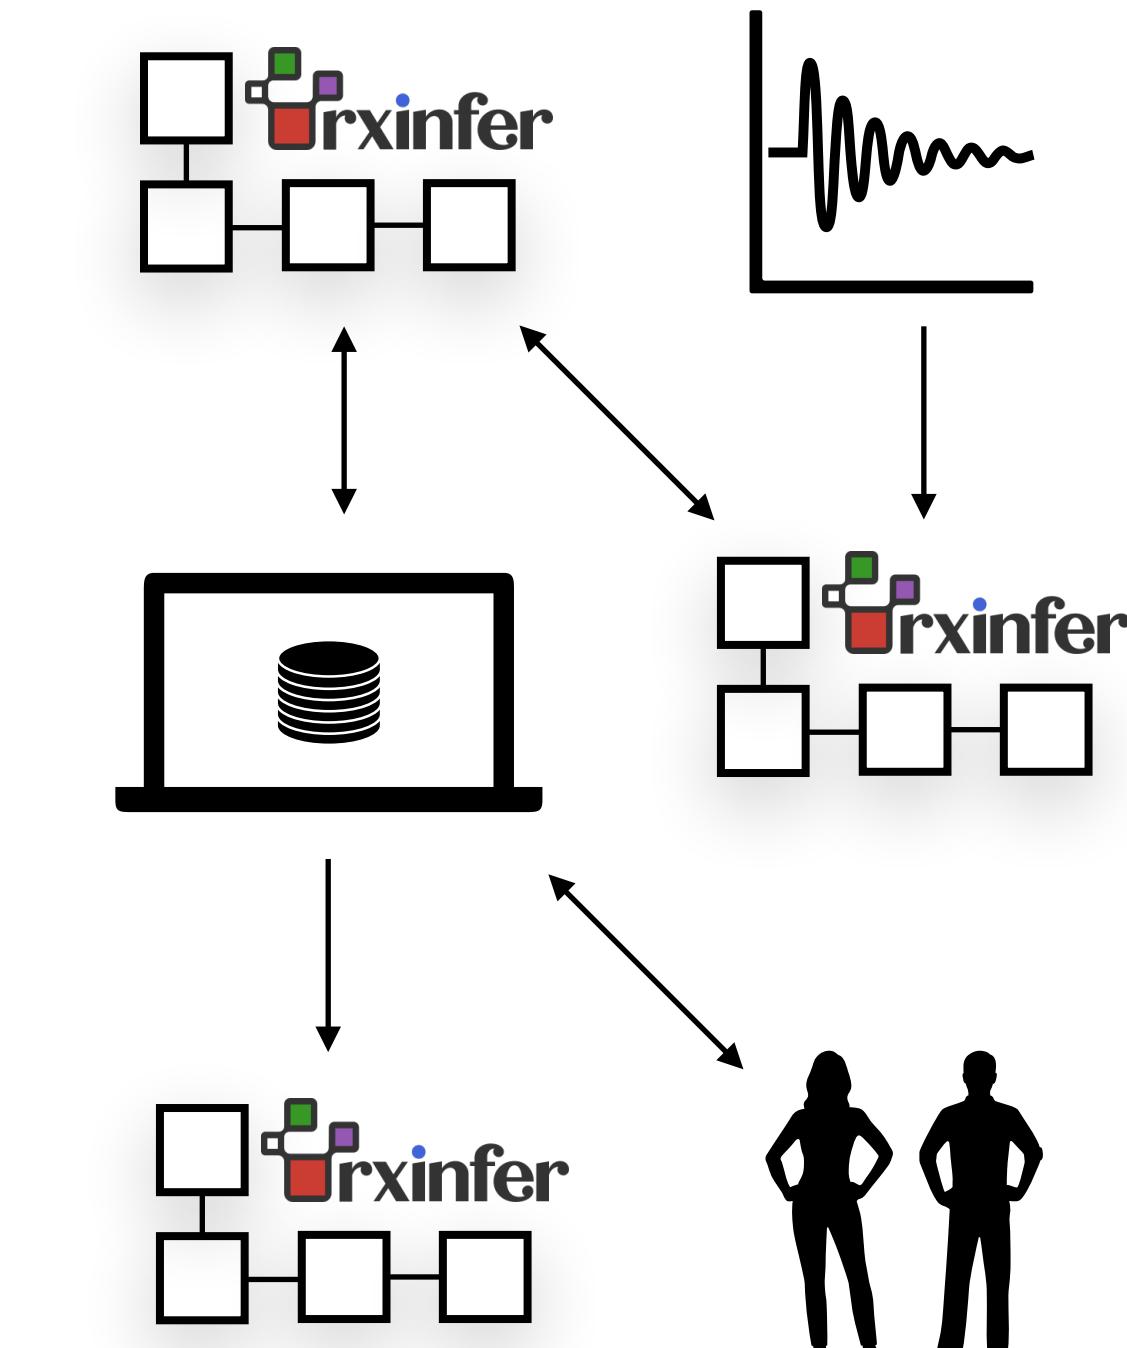
## Tradeoff accuracy and computational complexity

Some states change their values fast, while others change their values over long period



## Multi-agent interaction in changing environment

Reactivity is a natural way to support multi-agent interaction



## Many more

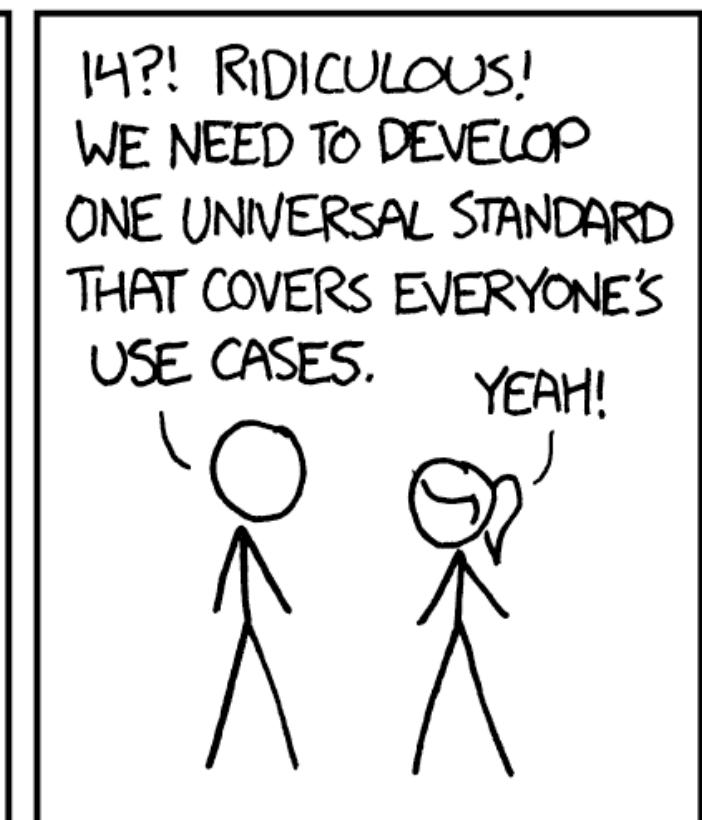
- Universality
- Dynamic constraints
- Documentation
- Usability
- EFE support
- Non-parametric
- Modular PPL
- Robustness
- Fixing



# How do we approach it?

And why don't we use another solution?

HOW STANDARDS PROLIFERATE:  
(SEE: A/C CHARGERS, CHARACTER ENCODINGS, INSTANT MESSAGING, ETC.)



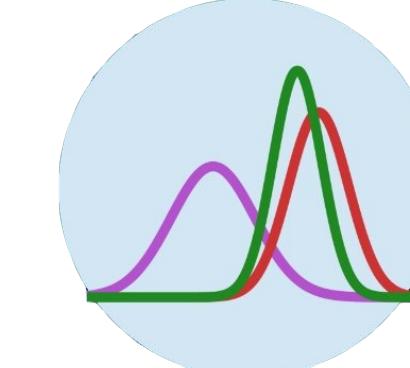
**Stan**

<https://mc-stan.org/>



**Pyro / NumPyro**

<https://pyro.ai/>



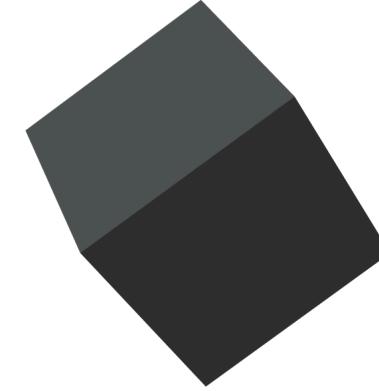
**Turing**

<https://turinglang.org/>



**PyMC3**

<https://www.pymc.io/>



**infer.NET**

<https://dotnet.github.io/infer/>



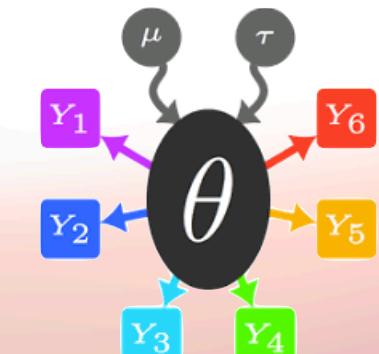
**BAT.jl**

<https://bat.mpp.mpg.de/>



**JASP**

<https://jasp-stats.org/>



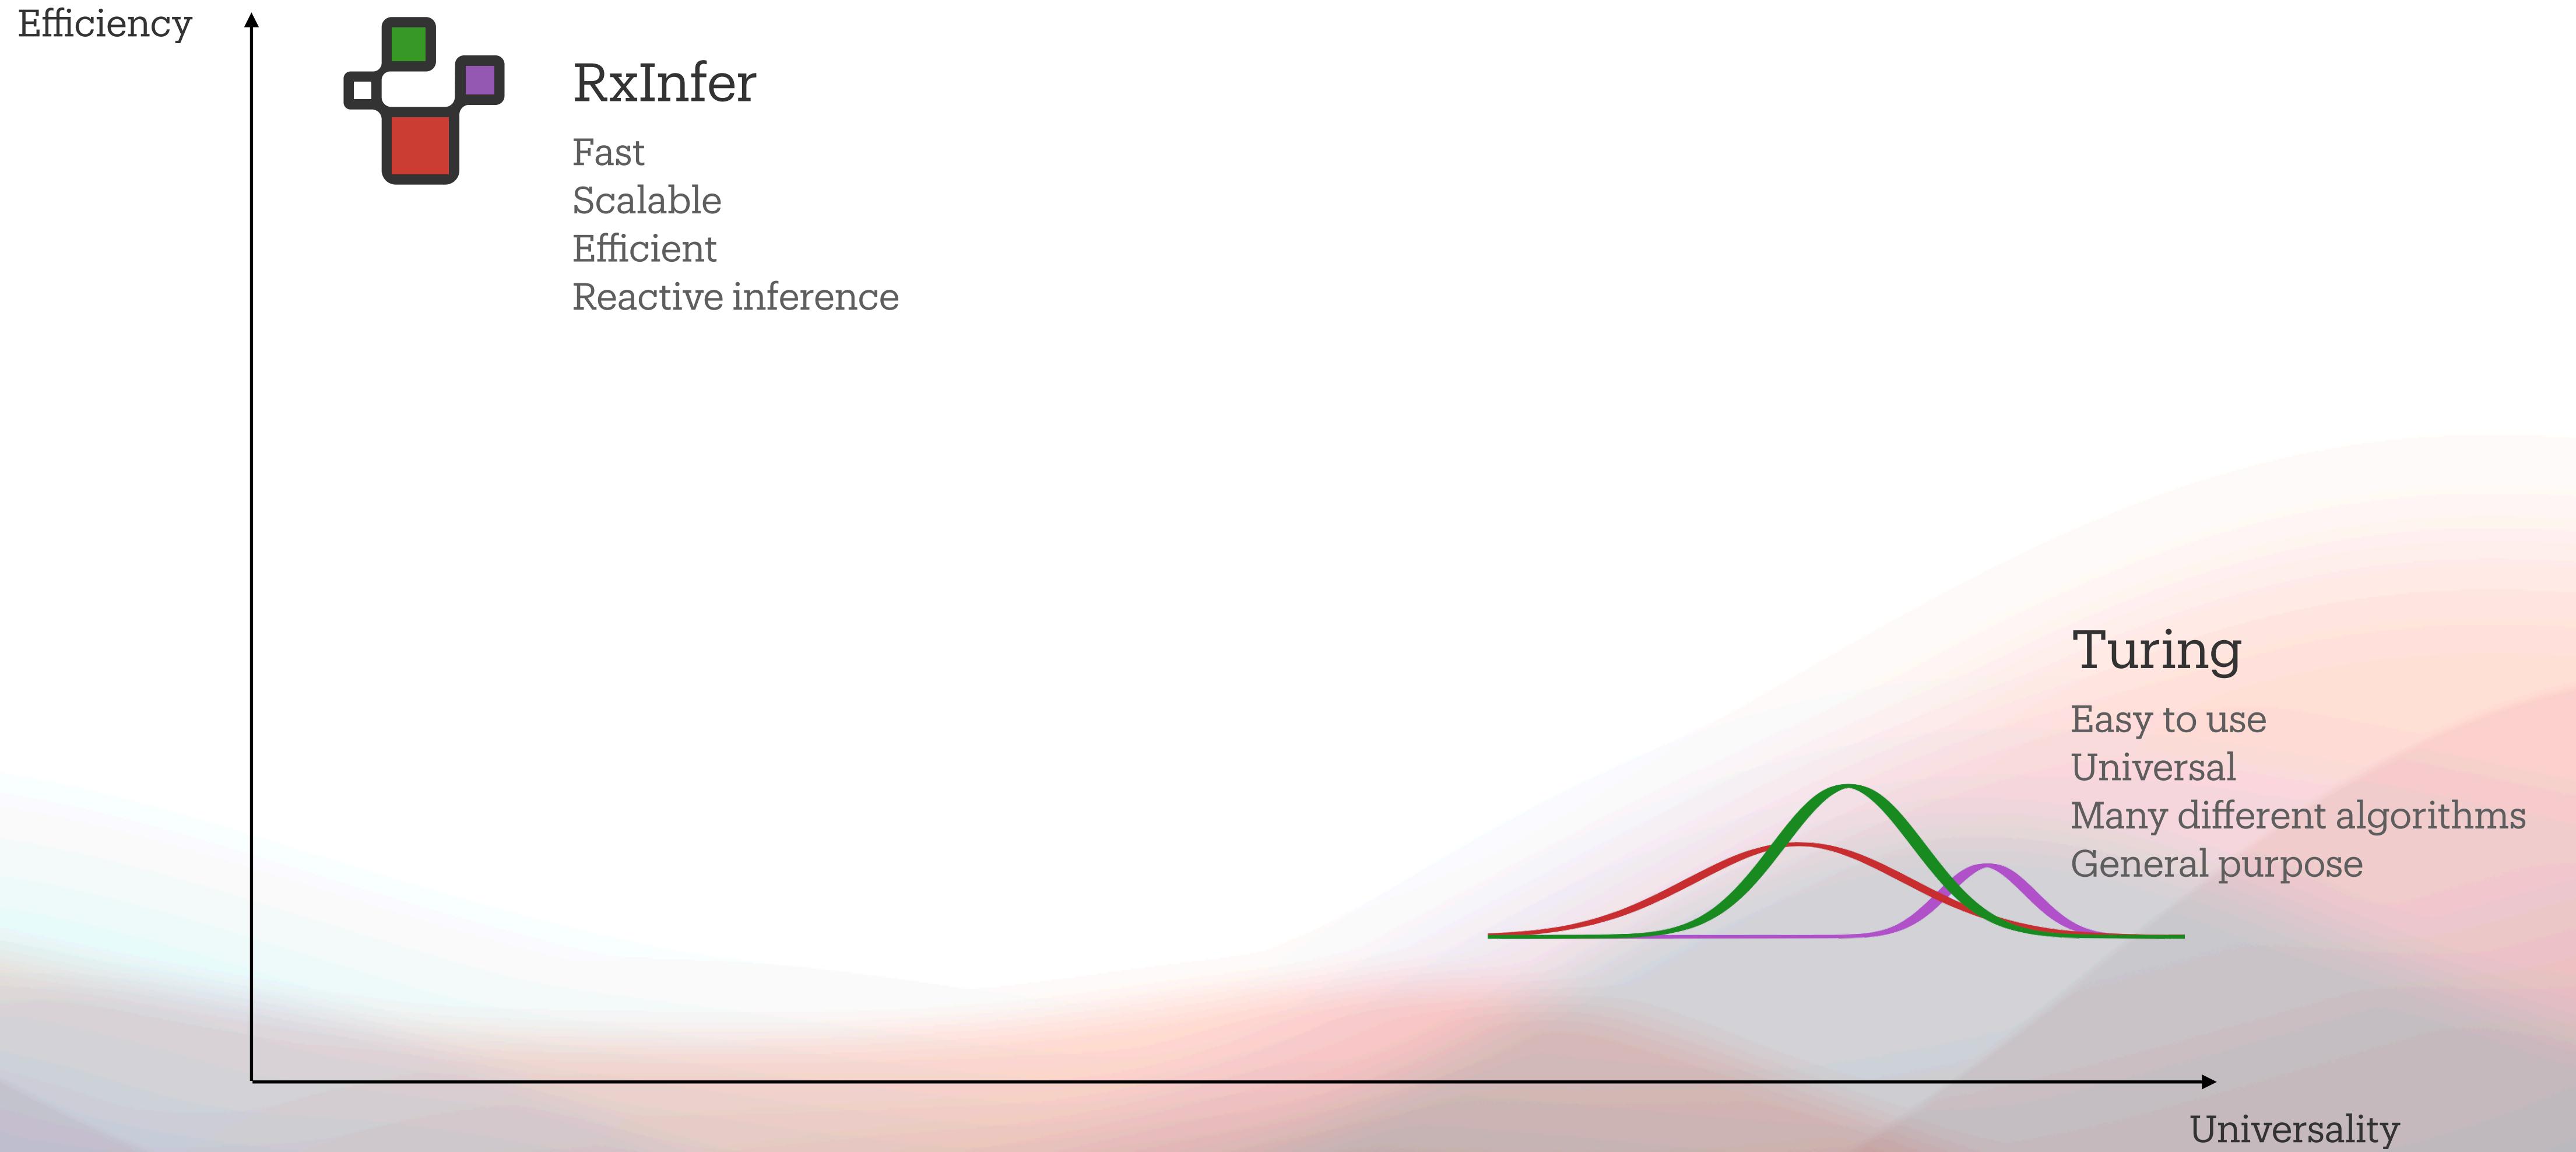
**BUGS**

<https://www.mrc-bsu.cam.ac.uk/software/bugs/>

**And others**

# How do we approach it?

And why don't we use another solution?



# User-friendliness

The toolbox should be easy to use

Generative model

Variational distribution

Compute posteriors

**How to specify?**

**How to specify?**

$$p(s, y)$$

$$q(s) \in \mathcal{Q}$$

$$q^*(s) = \arg \min_{q \in \mathcal{Q}} F [ q(s), p(s, y) ]$$

```
@constraints begin
    q(x, Q, P) = q(x)q(Q)q(P)

    q(Q) :: PointMass
    q(P) :: PointMass
    q(x) :: Gaussian
end
```

$$\mathcal{Q} = \left\{ q(x, Q, P) = q(x)q(Q)q(P), q(Q) = \delta(Q - \hat{Q}), q(P) = \delta(P - \hat{P}), q(X) = \mathcal{N}(x) \right\}$$

```
@model function rotate_ssm(n, x0, A, B)
    x = randomvar(n)
    y = datavar(Vector{Float64}, n)

    x_prior ~ MvNormal(μ = mean(x0), cov(x0))
    x_prev = x_prior

    Q ~ InverseWishart(3, diageye(2))
    P ~ InverseWishart(3, diageye(2))

    for i in 1:n
        x[i] ~ MvNormal(μ = A * x_prev, Σ = Q)
        y[i] ~ MvNormal(μ = B * x[i], Σ = P)
        x_prev = x[i]
    end
```

$$\begin{aligned} p(Q) &= \mathcal{IW}(Q|\nu_Q, K_Q) \\ p(P) &= \mathcal{IW}(P|\nu_P, K_P) \\ p(x_k|x_{k-1}) &= \mathcal{N}(x_k|Ax_{k-1}, Q) \\ p(y_k|x_k) &= \mathcal{N}(y_k|Bx_k, P) \end{aligned}$$

**How to minimise?**

```
result = inference(
    model = rotate_ssm(n, s0, A, B),
    constraints = constraints,
    data = (y = dataset, ),
    iterations = 10,
    free_energy = true
)
```

# Focus on customizability

The toolbox as a playground

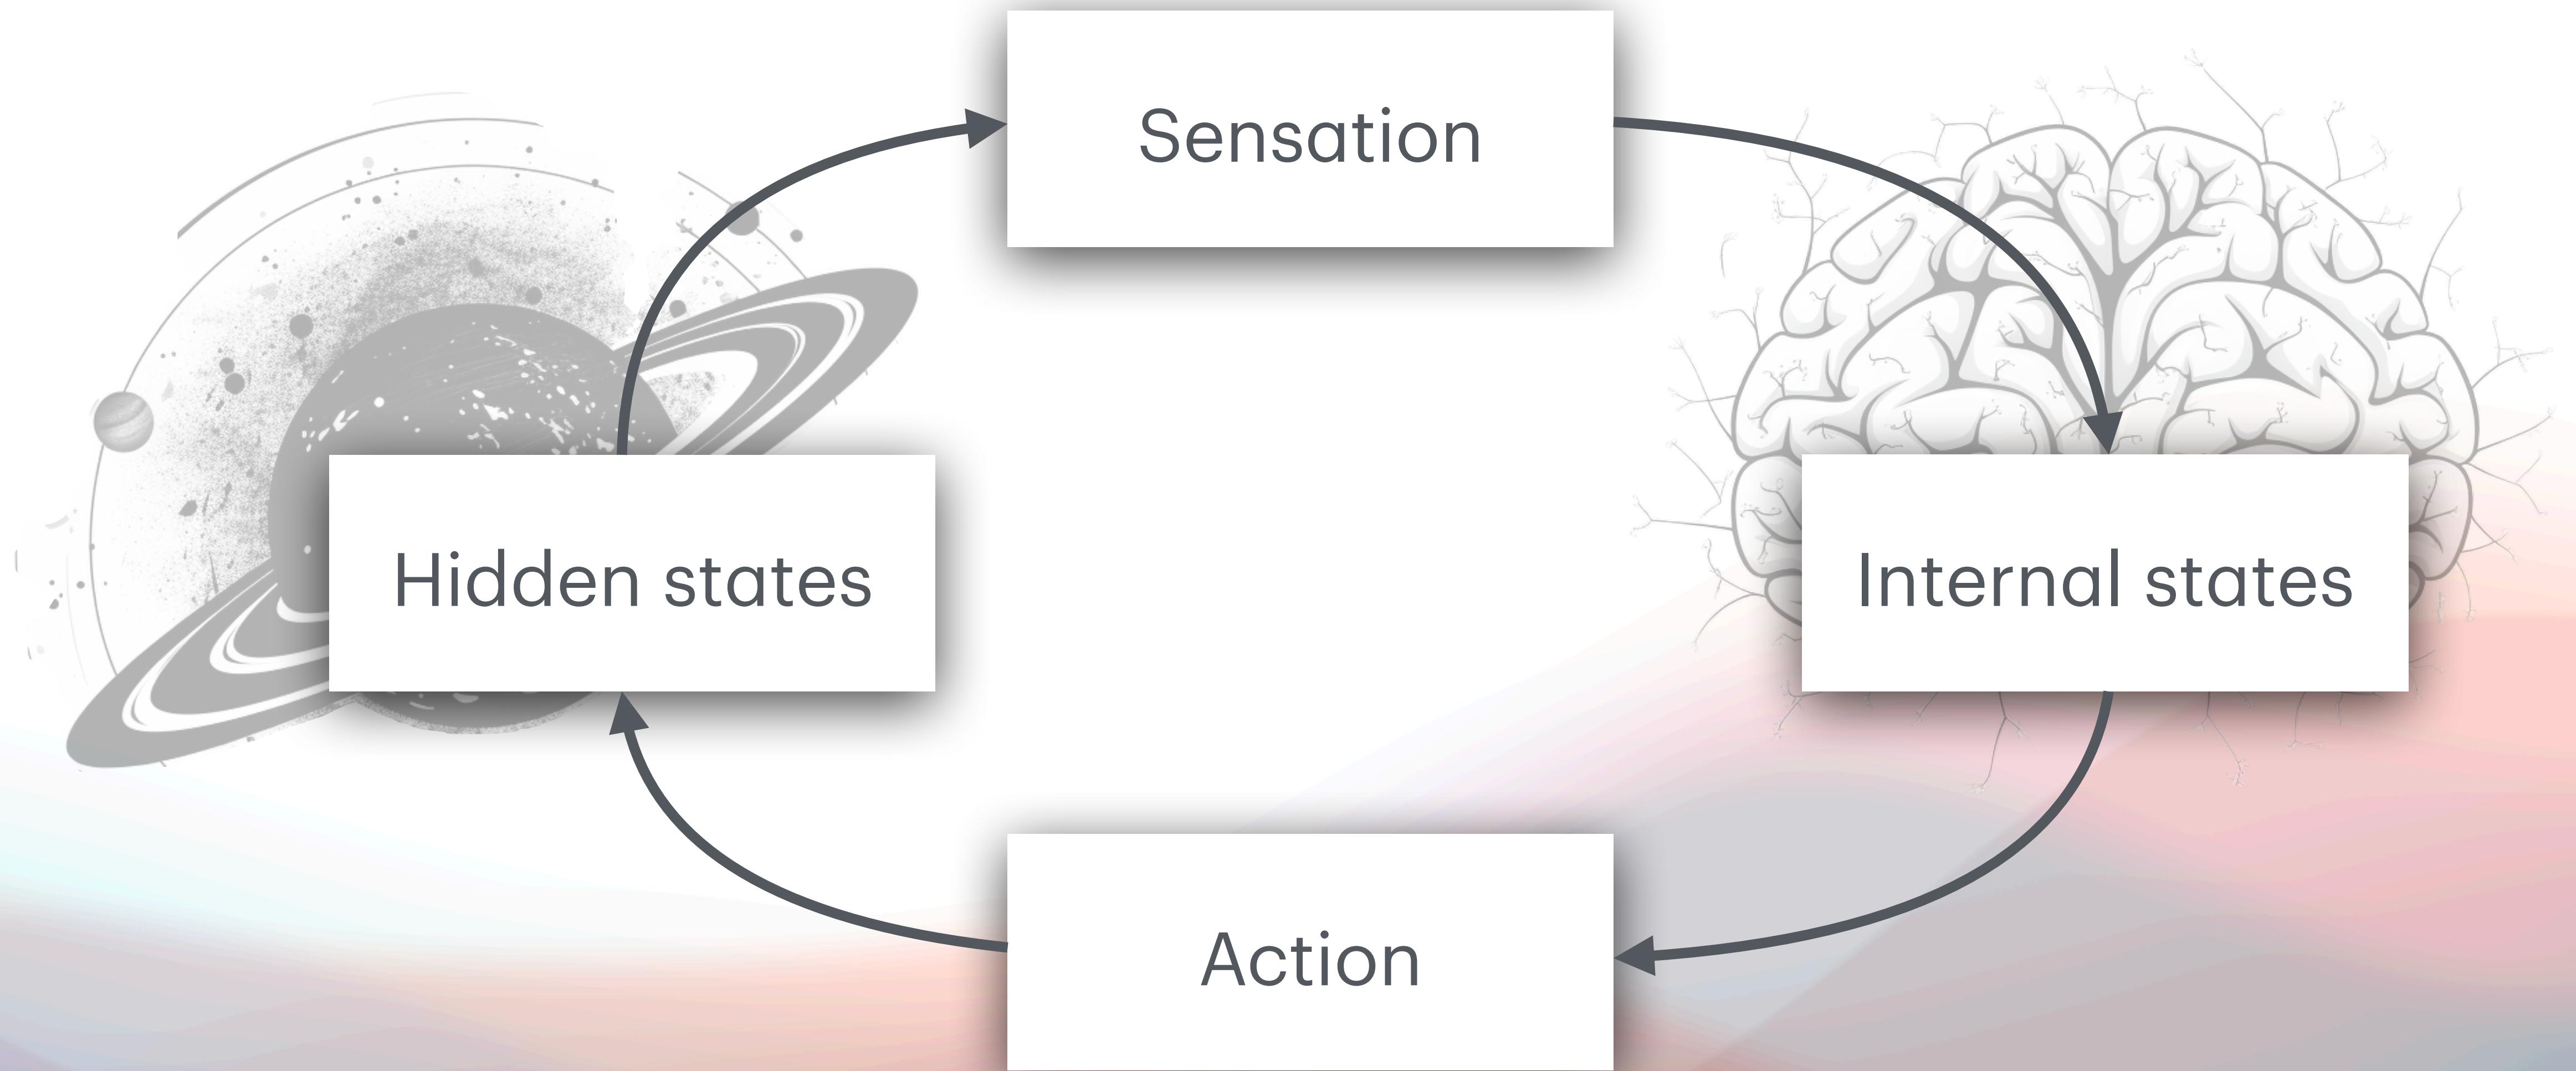
## Advanced message passing playground

- Easy to add new custom factor nodes
- Easy to add new message update rules
  - Belief propagation
  - Mean-field/Structured VI
  - Expectation maximisation
  - Expectation propagation
- Easy to redefine product of messages
- Easy to add new functional form constraints
- Easy to change approximation methods
  - Unscented transform
  - Linearization
  - Conjugate-Computation VI
- Custom plugins for message passing itself
- A large part of the toolbox is an “extension”
- Debugging messages



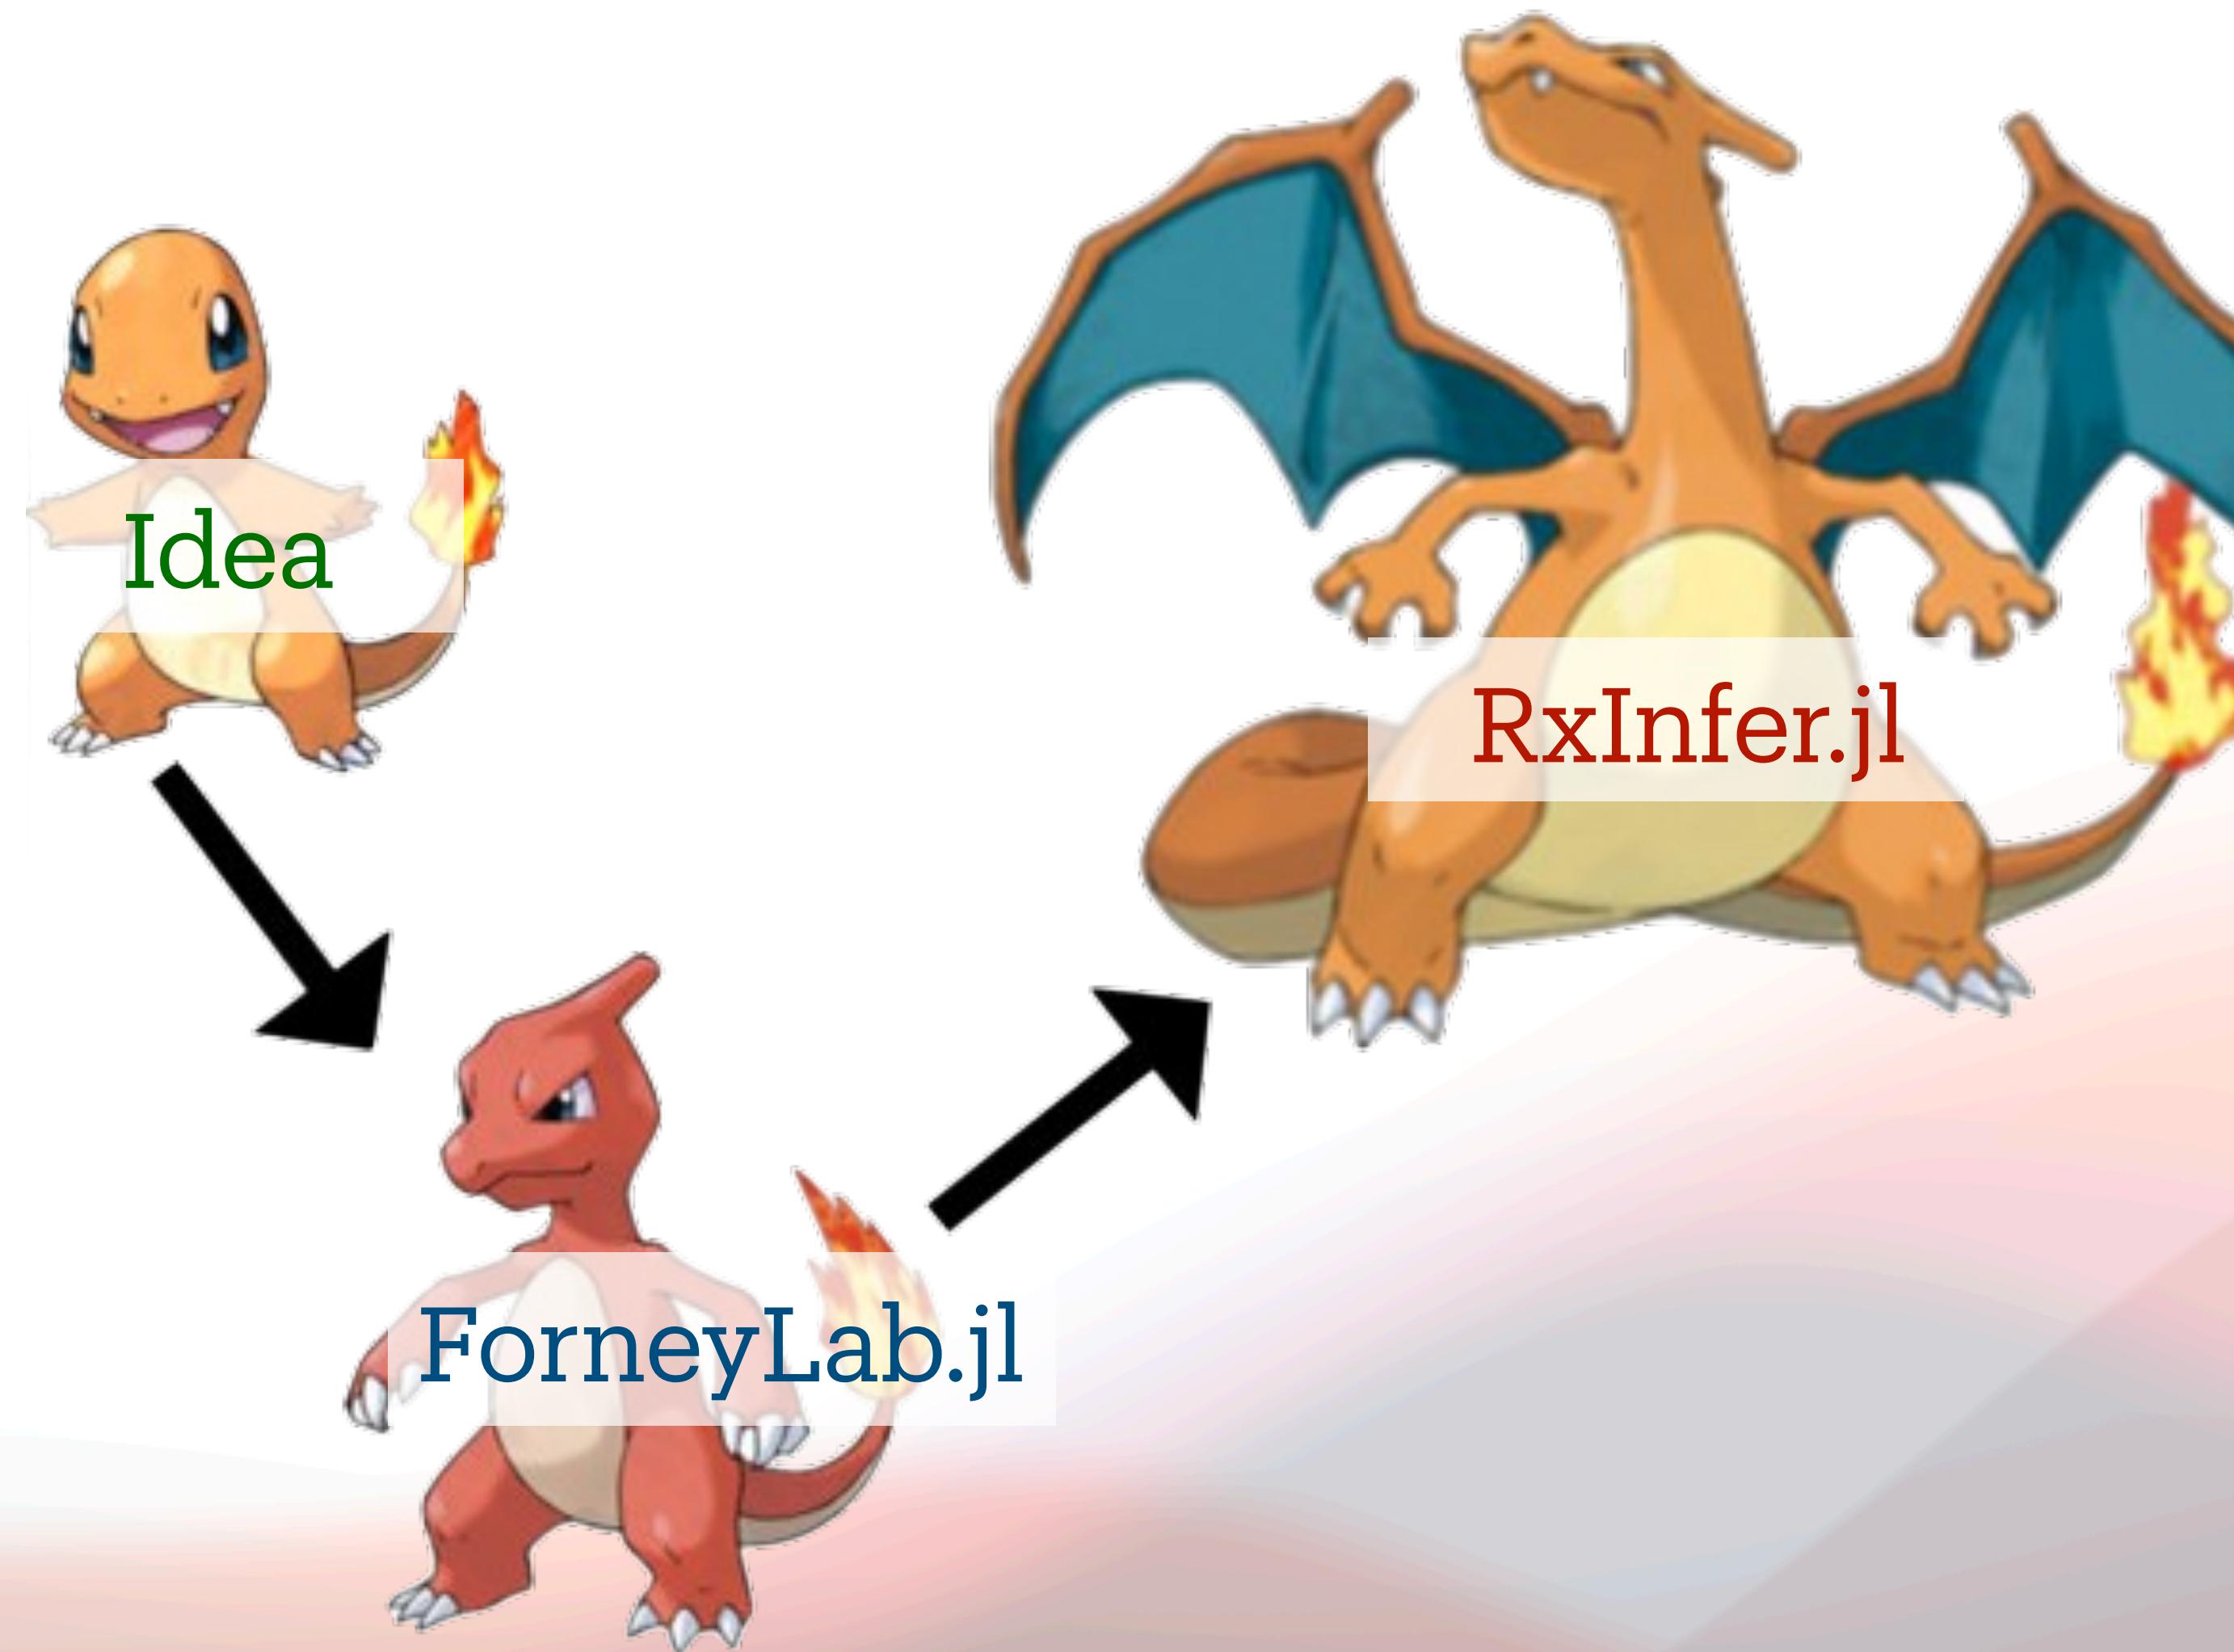
# How do we approach it?

Bayesian Brain hypothesis, Free Energy Principle and Active Inference



# Comparison to ForneyLab

Evolution of ideas and implementations



# Evolution of ideas and implementations

A lot of people contributed

## The team



bvdmitri



albertpod



bartvanerp



ismailsenoz

## The contributors



wmkouw



HoangMHNguyen



Chengfeng-Jia



Nimrais



wouterwln



Sepideh-Adamiat



ThijsvdLaar



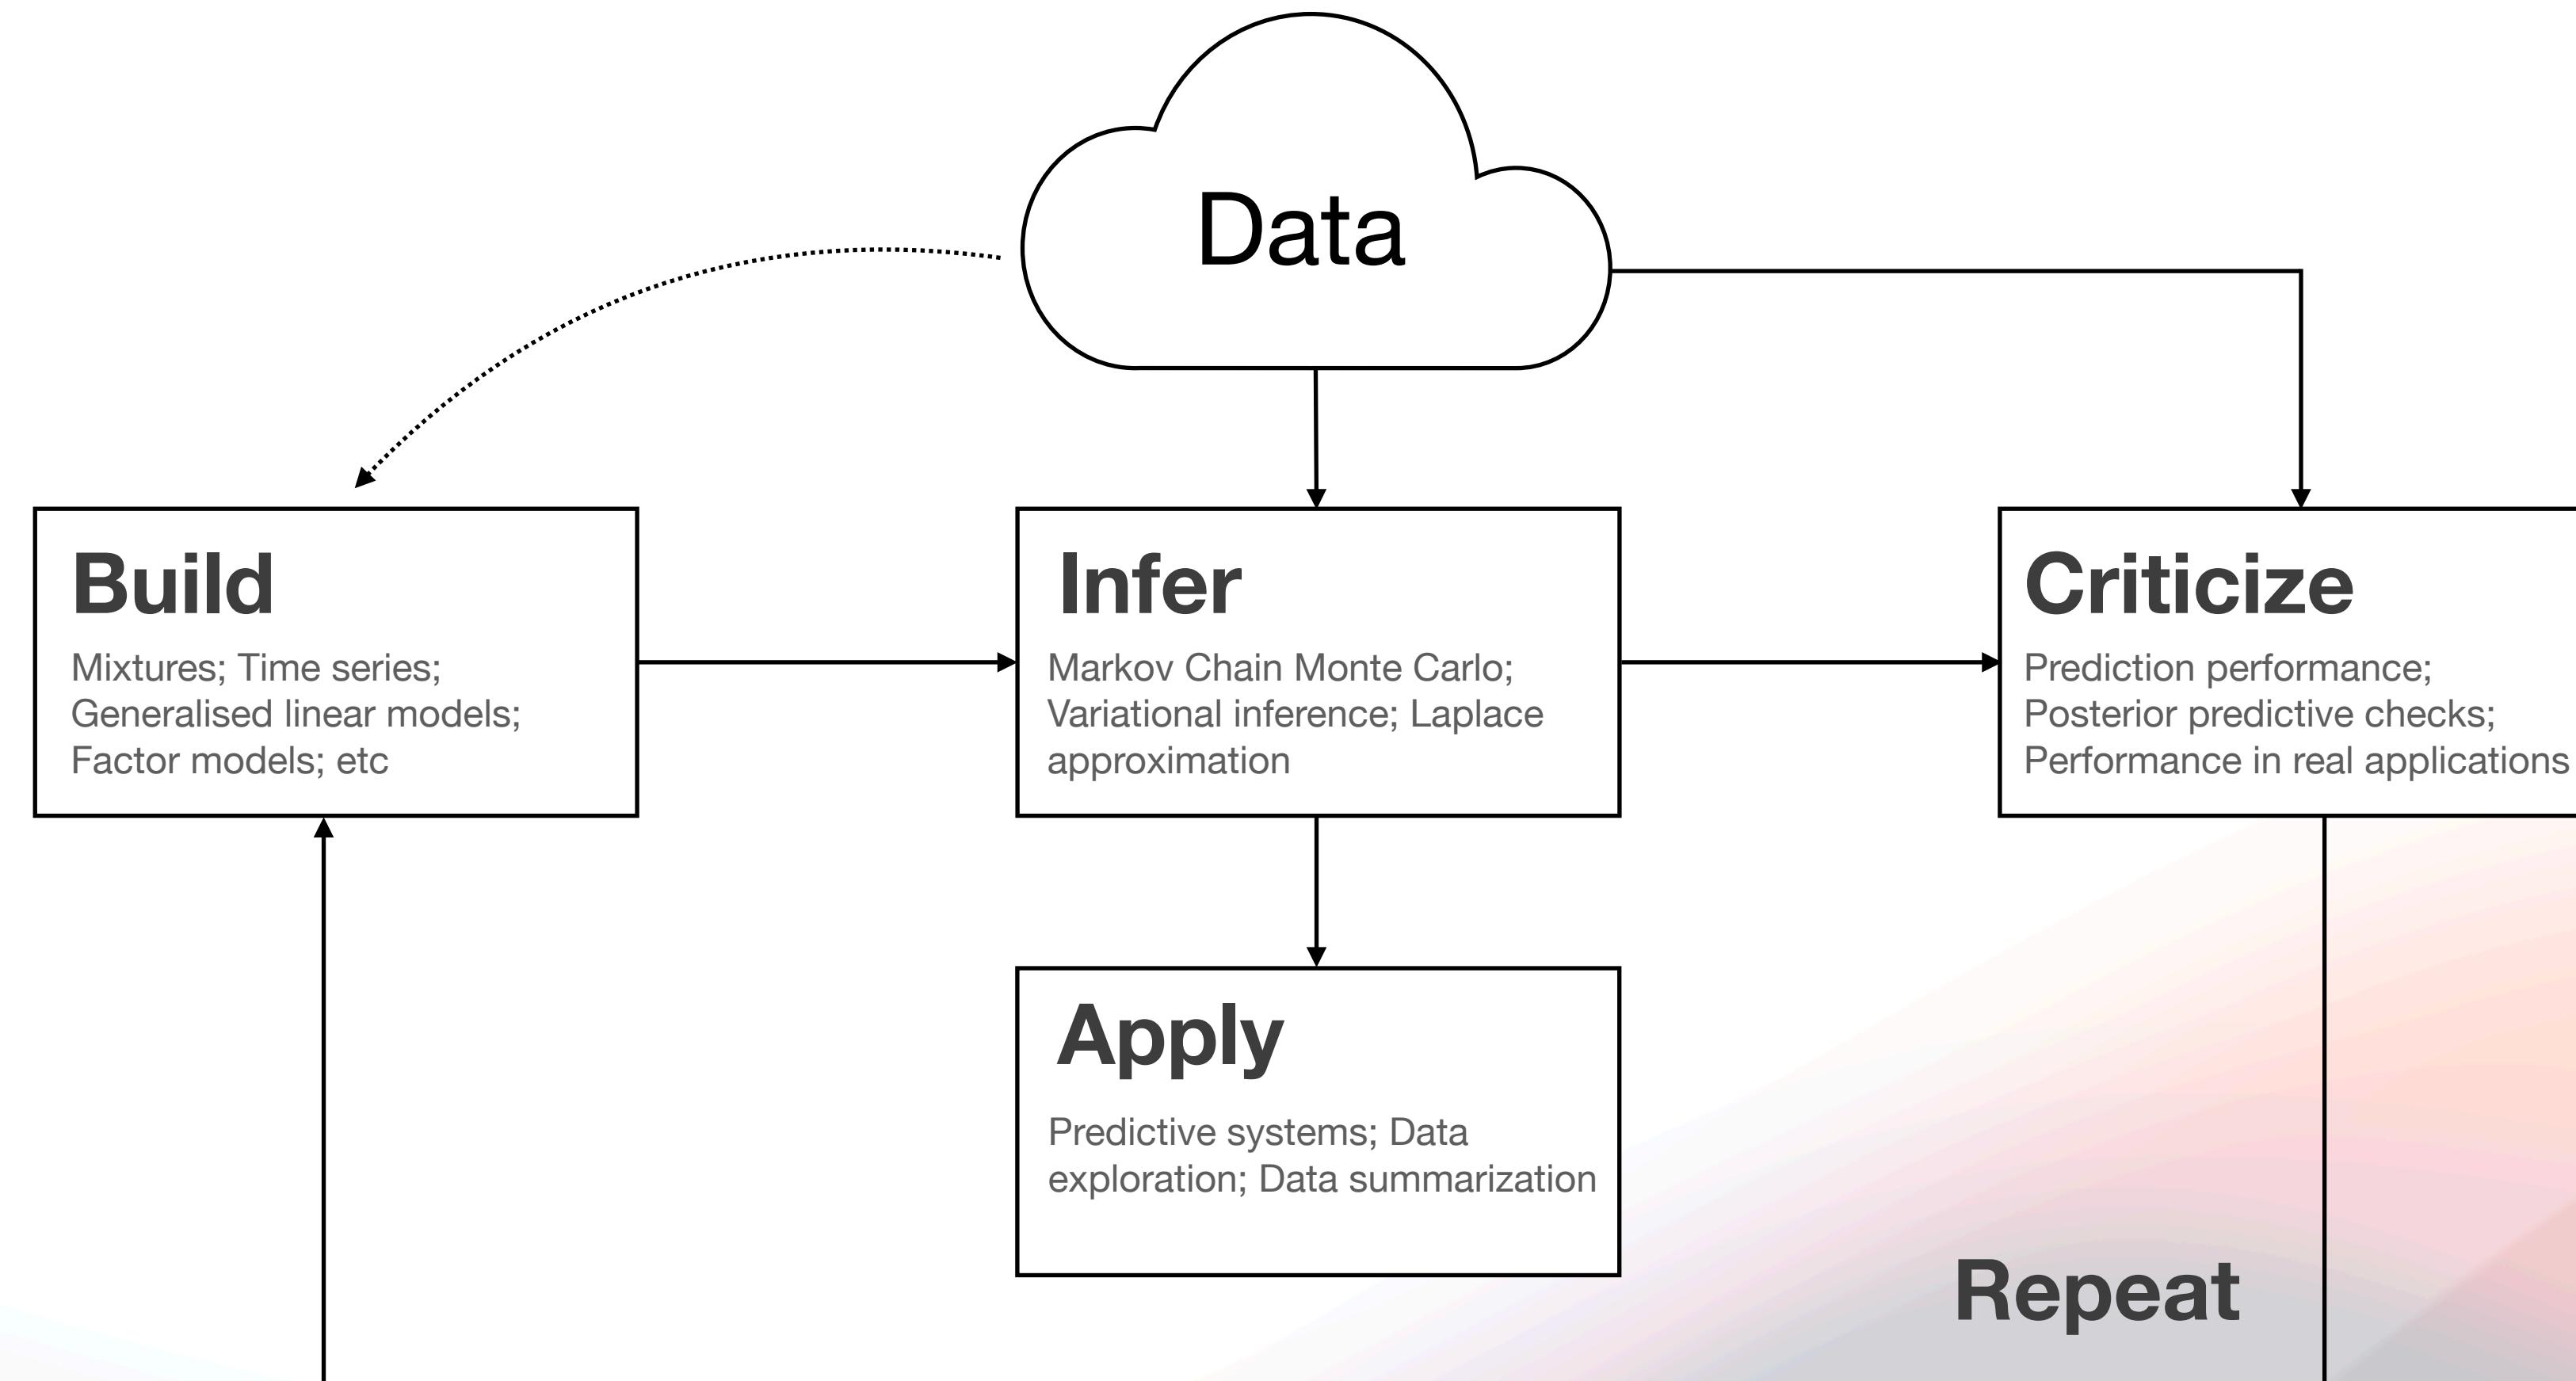
alstat



and others...

# Build, Infer, Apply, Criticize

until satisfactory or until your PhD contract runs out



# Efficient Bayesian inference on the edge

What can you do with RxInfer?



We do Real-Time Bayesian Inference Audio Processing On These Little  
Beauties With Low Power And Performance

# How do we specify prior knowledge?

Prior knowledge is an experience



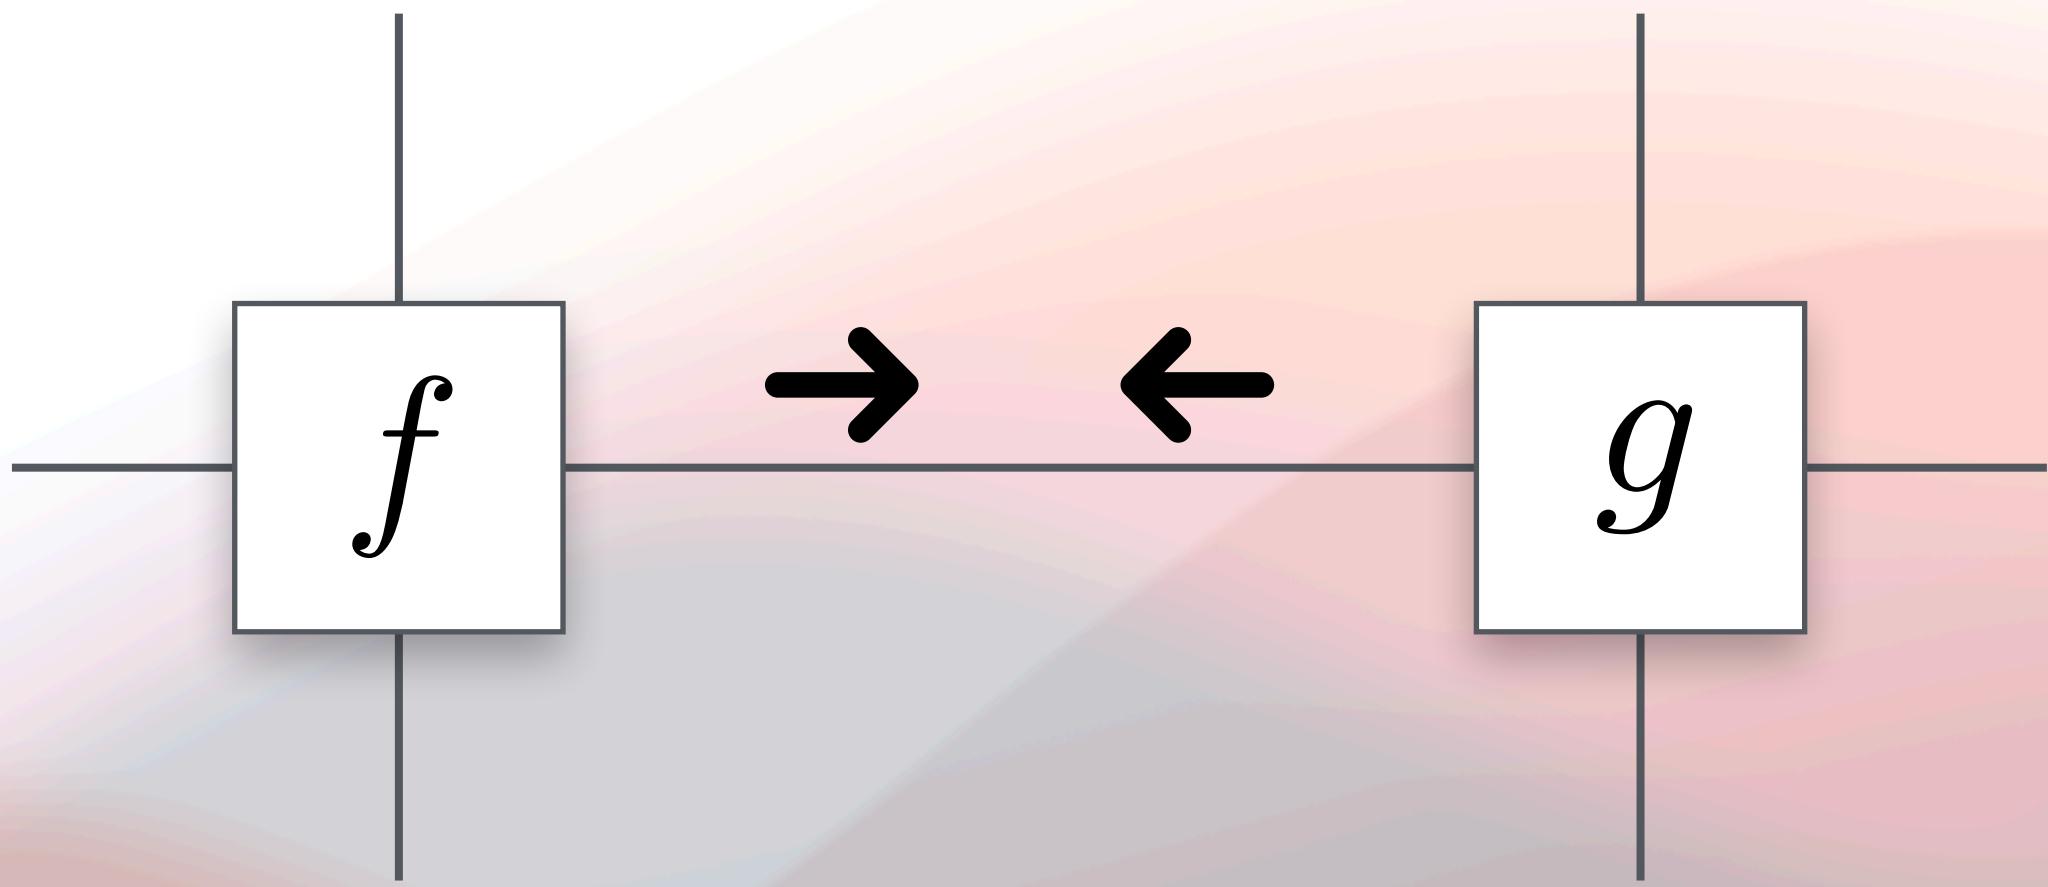
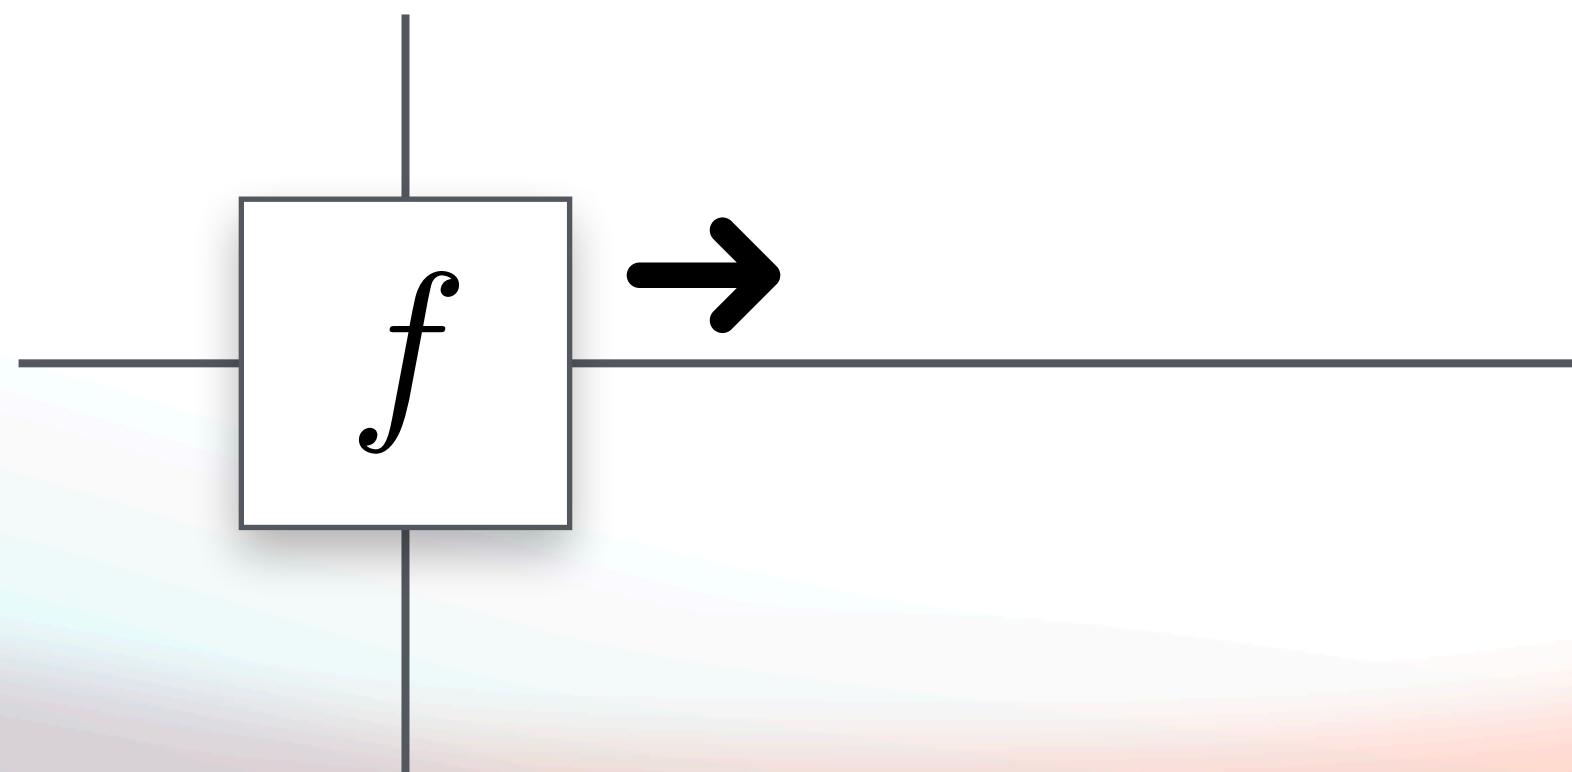
# Message passing in a nutshell

Variational Message Passing has two essential operations

$$f_a^n(\mathbf{s}_a^n) = \exp \left[ \int \left\{ \prod_{m \in \mathcal{C}(a), m \neq n} q_a^m(\mathbf{s}_a^m) \right\} \log f_a(\mathbf{s}_a) d\mathbf{s}_{a \setminus n} \right]$$

$$\mu_{ib}(s_i) = \int f_a^n(\mathbf{s}_a^n) \prod_{j \in \mathcal{C}(a,i), j \neq i} \mu_{ja}(s_j) ds_j$$

$$q_i(s_i) = \frac{\mu_{ia}(s_i)\mu_{ib}(s_i)}{\int \mu_{ia}(s_i)\mu_{ib}(s_i) ds_i}$$



Blank slide hello

Blank slide hello